

# Learned Query Optimizer: What is New and What is Next

Rong Zhu, Lianggui Weng, Bolin Ding, Jingren Zhou

Institute for Intelligent Computing  
Alibaba Group



# Speakers



Rong Zhu  
Alibaba Group



Lianggui Weng  
Alibaba Group



Bolin Ding  
Alibaba Group



Jingren Zhou  
Alibaba Group

# AI-Native Data Management

# Towards being AI Native: Tasks and Values

- Leveraging AI for traditional database operators and improving performance
  - Database tuning: AI models as tuners
  - Statistics estimation: AI models as histograms
  - Query optimization: AI models as optimizers
- Leveraging AI for usability, data analytics, and new applications
  - Text-to-SQL: an LLM-aided translation task
  - Data Manipulation: AI-aided data cleaning, data augmentation, and feature generation
  - Text-to-Insights: AI-aided analytical pipeline generation, copilot
- Data processing systems for AI
  - LLMs as users of the data system, and their evaluation results as feedbacks
  - Multi-modal data and processors (machine learning models)
  - Scalability and composability of data processors

# Learned Query Optimizer in AI-Native Data Management

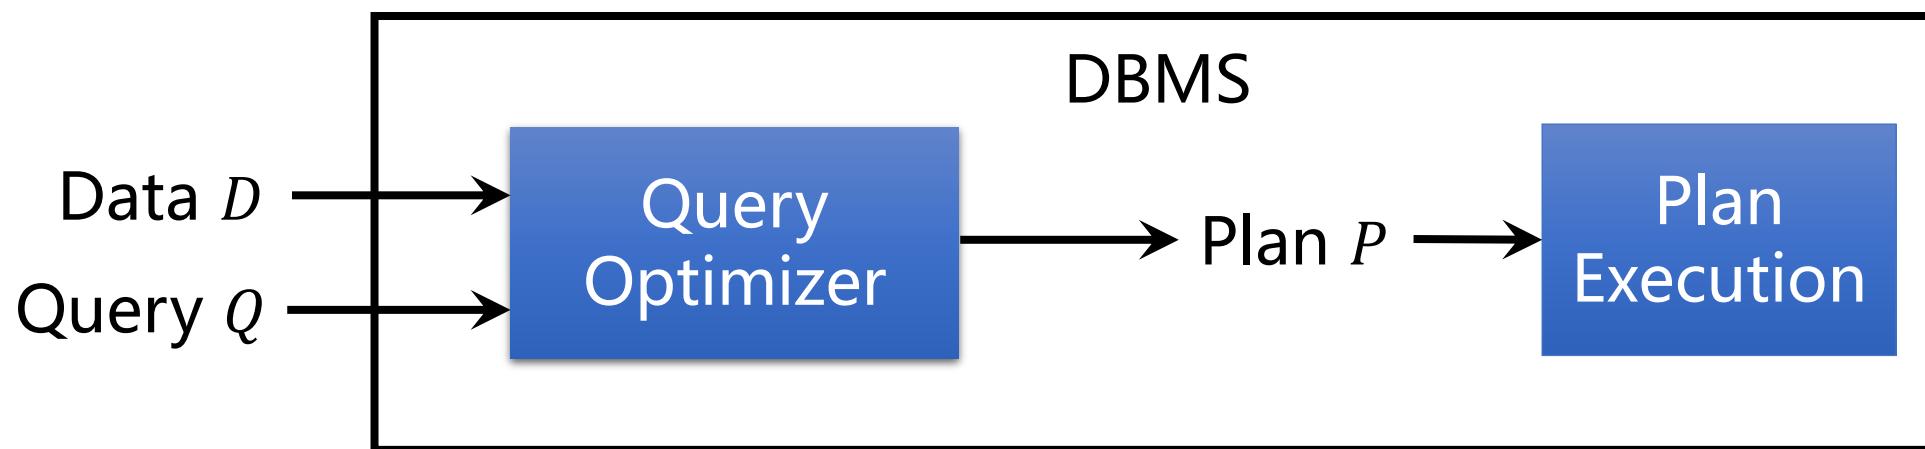
- ❑ Part 1: Preliminaries
- ❑ Part 2: Learn to Cost Plans
- ❑ Part 3: Learn to Explore Plan Space
- ❑ Part 4: Applications and Deployment
- ❑ Part 5: Summary and Future Work

# Part 1: Preliminaries

- Query Optimizer
- Learned Query Optimizer

# Query Optimizer: a Well-Defined Problem

- Generating best physical plan for input query
- Directly relates to the system performance
- Extensively studied in the literature

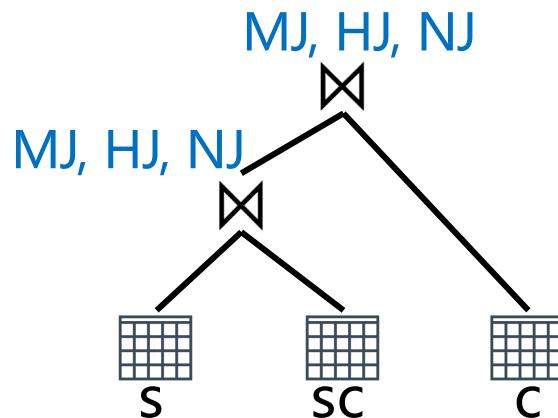


# Query Optimizer: a Difficult Task

Undefined part: plan quality

- Determined by data, query, system environment ...

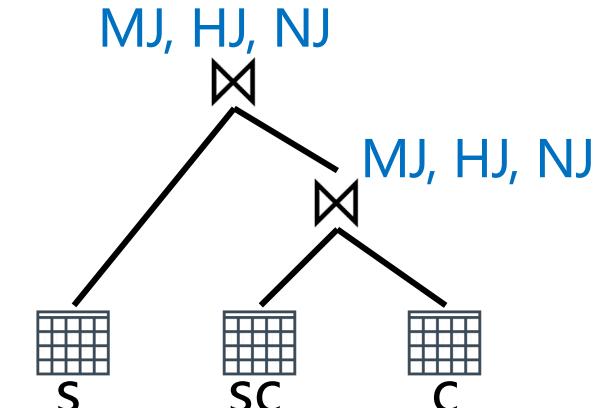
Large Search Space: join orders, access paths, ...



Seq Scan, Index Scan, ...

```
SELECT COUNT(*)  
FROM s, sc, c  
WHERE s.sid = sc.sid AND sc.cid = c.cid  
AND s.year < 2009 AND c.credit >= 3
```

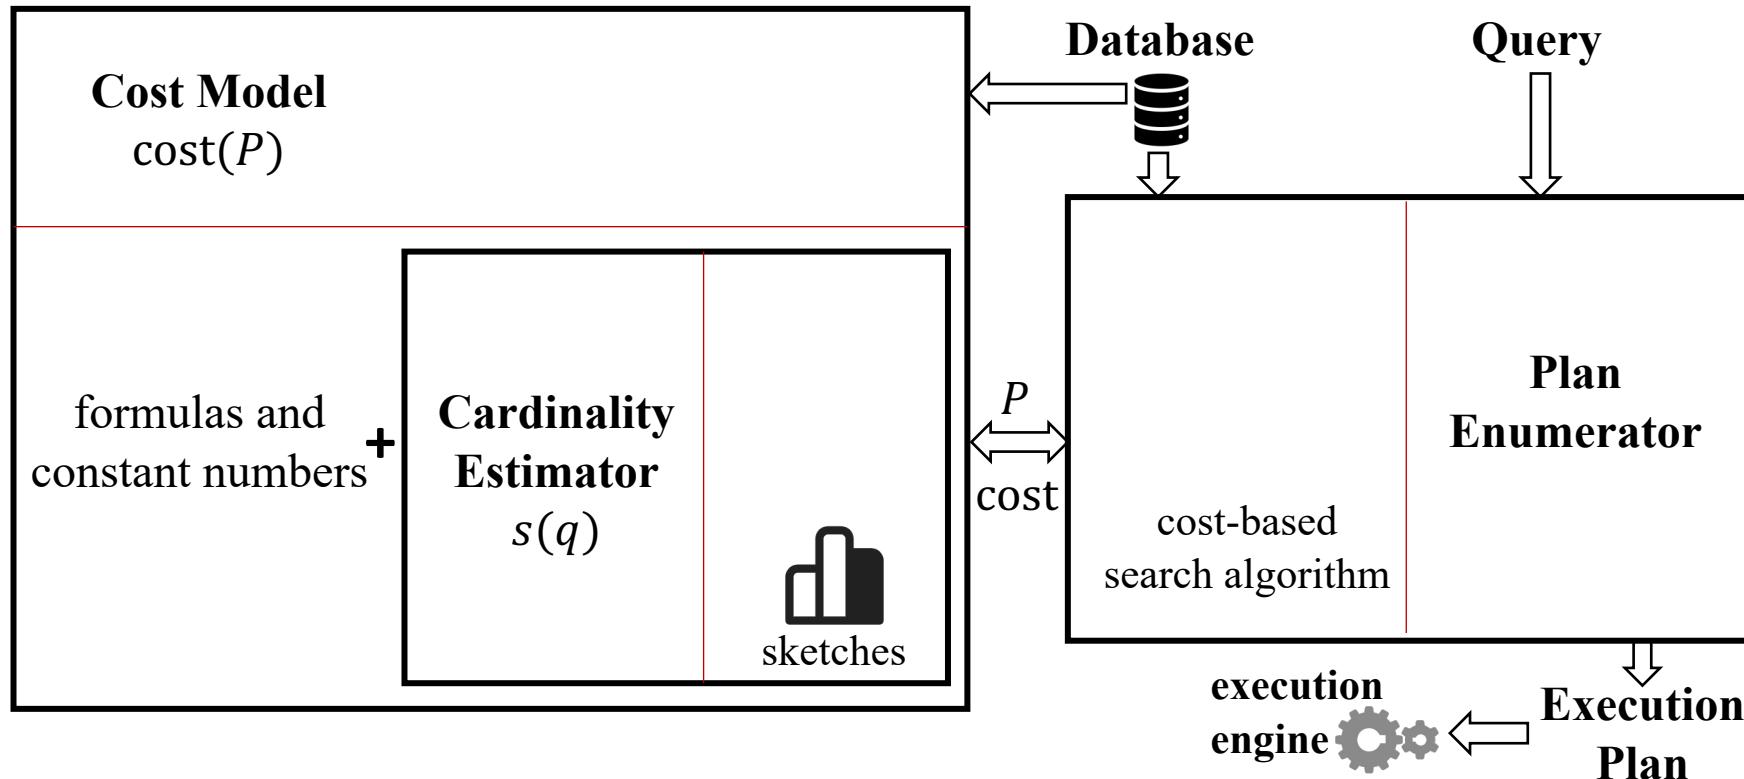
Which plan is better?



Seq Scan, Index Scan, ...

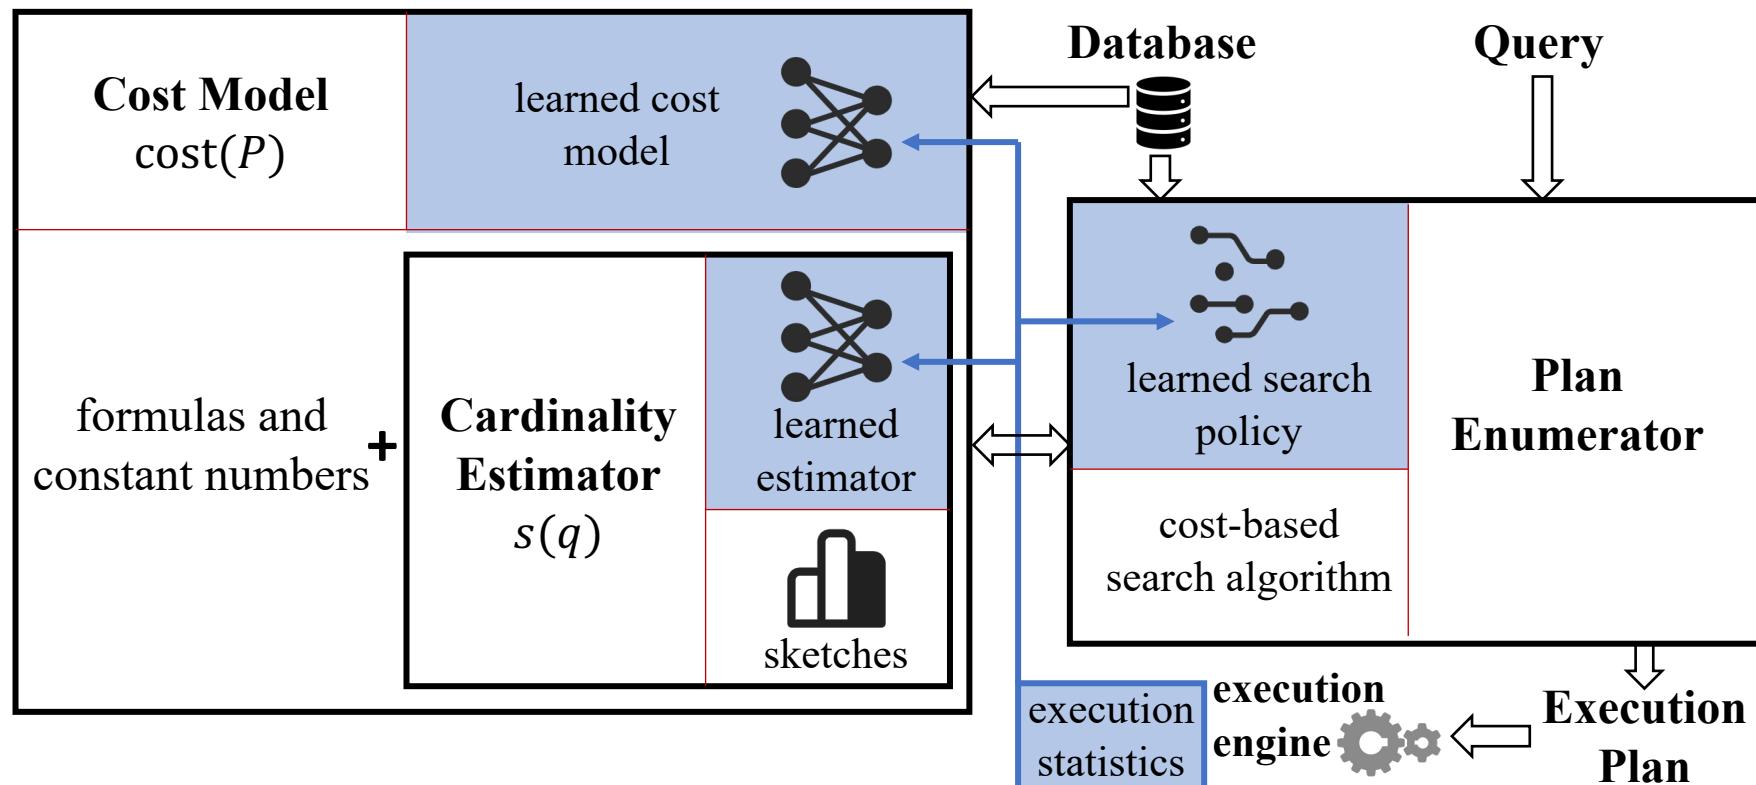
# Learned Query Optimizer: a Pioneer in AI4DB

- A typical cost-based query optimizer



# Learned Query Optimizer: a Pioneer in AI4DB

- Learned query optimizer: possibilities and opportunities



Learned Parts in Query Optimizer

# Learned Query Optimizer in AI-Native Data Management

- ❑ Part 1: Preliminaries
- ❑ Part 2: Learn to Cost Plans
- ❑ Part 3: Learn to Explore Plan Space
- ❑ Part 4: Applications and Deployment
- ❑ Part 5: Summary and Future Work

# Part 2: Learn to Cost Plans

- Cardinality & Cost
- Learn to Estimate Cardinality
- Learned Cost Model

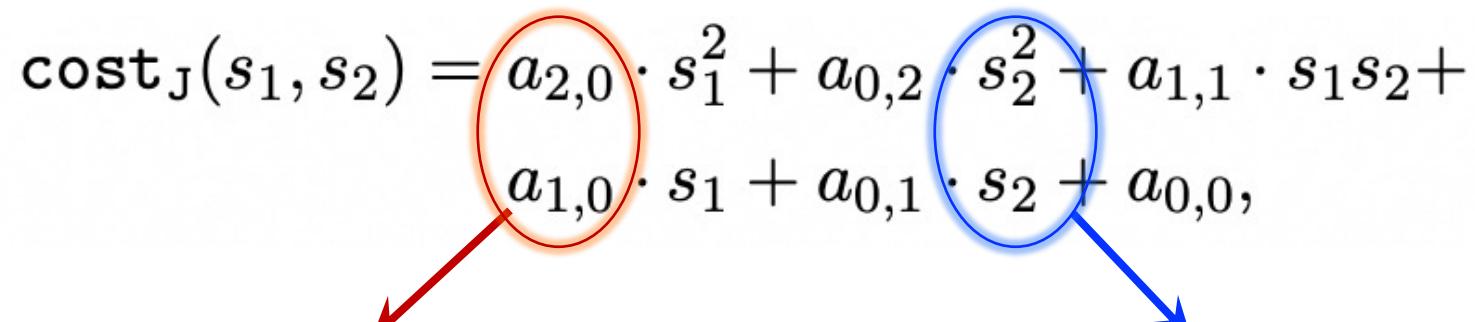
# Part 2: Learn to Cost Plans

- Cardinality & Cost
- Learn to Estimate Cardinality
- Learned Cost Model

# Cardinality & Cost

- Cardinality: size  $C(q)$  of the result of a query  $q$ 
  - Require to estimate before actual execution of queries
- Cost: an approximation to query execution time/resource usage
  - Cost model in forms of experience-driven formulas
  - Example: a join on two relations with size (cardinality)  $s_1$  and  $s_2$

$$\text{cost}_J(s_1, s_2) = a_{2,0} \cdot s_1^2 + a_{0,2} \cdot s_2^2 + a_{1,1} \cdot s_1 s_2 + a_{1,0} \cdot s_1 + a_{0,1} \cdot s_2 + a_{0,0},$$

  
Expert-tuning Factors      Estimated Cardinality

# Cardinality & Cost

- Cardinality to Cost: the error bound
  - Q-error of estimated cardinality  $\hat{s}$  to exact value  $s$

$$\|s/\hat{s}\|_Q = \max(s/\hat{s}, \hat{s}/s).$$

- Error bound on the cost of a join operator

**Proposition 2.1** (Moerkotte et al., 2009). Suppose the Q-errors of  $\hat{s}_1$  and  $\hat{s}_2$  for estimating  $s_1$  and  $s_2$  are bounded by  $\epsilon$ , we have the Q-error of  $\text{cost}_J(\hat{s}_1, \hat{s}_2)$  for estimating  $\text{cost}_J(s_1, s_2)$  bounded by  $\epsilon^2$ .

- Error bound on the plan  $\hat{P}$  with minimum cost and best plan  $P^*$

$$P^* = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; s) \quad \text{and} \quad \hat{P} = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; \hat{s}).$$

**Theorem 2.1** (Moerkotte et al., 2009). For a given query  $q$ , suppose the Q-error of cardinality estimation for intermediate result of any sub-query is bounded by  $\epsilon$ , we have  $\text{cost}(\hat{P}; s) \leq \epsilon^4 \text{cost}(P^*; s)$ .

# Cardinality & Cost

- Cardinality to Cost: the error bound
  - Q-error of estimated cardinality  $\hat{s}$  to exact value  $s$

$$\|s/\hat{s}\|_Q = \max(s/\hat{s}, \hat{s}/s).$$

worst-case

- Error bound on the cost of a join operator

**Proposition 2.1** (Moerkotte et al., 2009). Suppose the Q-errors of  $\hat{s}_1$  and  $\hat{s}_2$  for estimating  $s_1$  and  $s_2$  are bounded by  $\epsilon$ , we have the Q-error of  $\text{cost}_J(\hat{s}_1, \hat{s}_2)$  for estimating  $\text{cost}_J(s_1, s_2)$  bounded by  $\epsilon^2$ .

Inaccurate Cardinality

- Error bound on the plan  $\hat{P}$  with minimum cost and best plan  $P^*$

$$P^* = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; s) \quad \text{and} \quad \hat{P} = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; \hat{s}).$$



Inaccurate Cost



Sub-optimal Plan

**Theorem 2.1** (Moerkotte et al., 2009). For a given query  $q$ , suppose the Q-error of cardinality estimation for intermediate result of any sub-query is bounded by  $\epsilon$ , we have  $\text{cost}(\hat{P}; s) \leq \epsilon^4 \text{cost}(P^*; s)$ .

# Cardinality & Cost

- Cardinality to Cost: the error bound
    - Q-error of estimated cardinality  $\hat{s}$  to exact value  $s$ 
$$\|s/\hat{s}\|_Q = \max(s/\hat{s}, \hat{s}/s).$$
hopefully
    - Error bound on the cost of a join operator

**Proposition 2.1** (Moerkotte et al., 2009). Suppose the Q-errors of  $\hat{s}_1$  and  $\hat{s}_2$  for estimating  $s_1$  and  $s_2$  are bounded by  $\epsilon$ , we have the Q-error of  $\text{cost}_J(\hat{s}_1, \hat{s}_2)$  for estimating  $\text{cost}_J(s_1, s_2)$  bounded by  $\epsilon^2$ .
    - Error bound on the plan  $\hat{P}$  with minimum cost and best plan  $P^*$ 
$$P^* = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; s) \quad \text{and} \quad \hat{P} = \arg \min_{P \in \mathbb{P}(q)} \text{cost}(P; \hat{s}).$$

**Theorem 2.1** (Moerkotte et al., 2009). For a given query  $q$ , suppose the Q-error of cardinality estimation for intermediate result of any sub-query is bounded by  $\epsilon$ , we have  $\text{cost}(\hat{P}; s) \leq \epsilon^4 \text{cost}(P^*; s)$ .
- 
- Better Cardinality
- Better Cost
- Better Plan

# Part 2: Learn to Cost Plans

- Cardinality & Cost
- Learn to Estimate Cardinality
- Learned Cost Model

# Traditional Methods

- Most widely used in modern DBMS
- Histogram: attribute independent assumption  $\Pr(A) \approx \prod_i \Pr(A_i)$ 
  - Fast inference, low storage cost, high estimation error
  - Multi-dimensional histogram: high storage cost
- Sampling: execute the query on a smaller fraction of data
  - Estimation accuracy and inference speed trade-off
  - Kernel-density estimation, join sampling, ...

# Learned Methods: An Overview

- Data-driven methods:
  - Learned models approximating  $\Pr(A)$  of table  $T$  on attributes  $A$
  - $C(q) = \Pr(q) * |T|$  for  $q \subseteq T$
- Query-driven methods: for a workload (distribution)  $Q$ 
  - Learned models as a function  $C(q)$  for  $q \in Q$
- Hybrid methods:
  - Utilize both data and query information

# Data-Driven Methods

- Build models to capture data distribution  $\Pr(A)$
- Different methods to model  $\Pr(A)$ 
  - Auto-Regression model
  - Probabilistic graphical model (PGM)
  - ...

# Data-Driven Methods

- Auto-Regression models: chain rule of probability

$$\Pr(A) = \Pr(A_1) \cdot \Pr(A_2 | A_1) \cdot \Pr(A_3 | A_2, A_1) \cdots \Pr(A_n | A_{n-1}, A_{n-2}, \dots, A_1)$$

- Naru: single table
  - NeuroCard: multi-tables as a joined table
  - IAM: adding Gaussian mixture models
- 
- Histogram: attribute independent assumption  $\Pr(A) \approx \prod_i \Pr(A_i)$

(Naru) Deep Unsupervised Cardinality Estimation, VLDB 2019

(NeuroCard) NeuroCard: One Cardinality Estimator for All Tables, VLDB 2020

(IAM) Unsupervised Selectivity Estimation by Integrating Gaussian Mixture Models and an Autoregressive Model, EDBT 2022

# Data-Driven Methods

- Auto-Regression models: model training

Train a model to approximating  
 $\Pr(x_k|x_{k-1}, x_{k-2}, \dots, x_1)$  for  $x_i \in A_i$

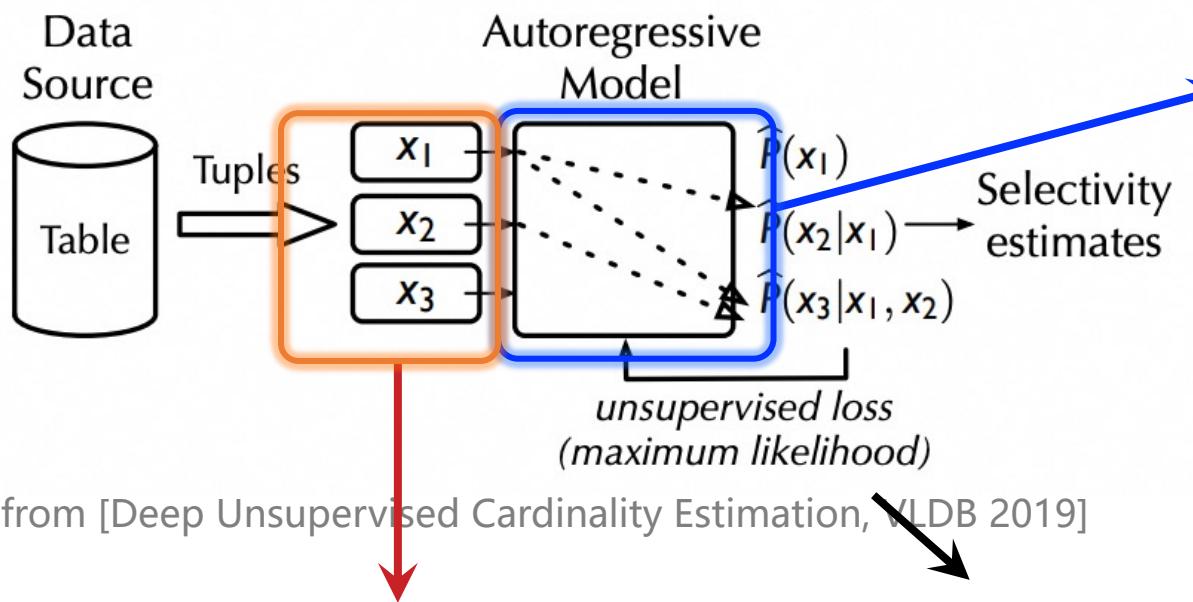


Figure from [Deep Unsupervised Cardinality Estimation, VLDB 2019]

Training data: draw tuple  $x$   
uniformly random from table

$$\text{Minimize loss } L = -\frac{1}{|R|} \sum_{x \in R} \log(\Pr(x))$$

# Data-Driven Methods

- Auto-Regression models: inference (probability computation)
  - Point query: compute the probability from model directly

$$\Pr(a, b, c) = \Pr(A_1 = a) \cdot \Pr(A_2 = b | A_1 = a) \cdot \Pr(A_3 = c | A_2 = b, A_1 = a)$$

- Range query: sum the probability of all points: low efficiency

$$\Pr(a, b, c \in U) = \sum_{a', b', c' \in U} \Pr(a') \cdot \Pr(b' | a') \cdot \Pr(c' | b', a')$$

- Speed up by sampling a number of points
  - Tradeoff between efficiency and accuracy

$$\Pr(a, b, c \in U) = \frac{|U|}{|S|} \cdot \sum_{a', b', c' \in S} \Pr(a') \cdot \Pr(b' | a') \cdot \Pr(c' | b', a')$$

# Data-Driven Methods

Auto-Regression models: chain rule of probability

$$\Pr(A) = \Pr(A_1) \cdot \Pr(A_2|A_1) \cdot \Pr(A_3|A_2, A_1) \cdots \Pr(A_n|A_{n-1}, A_{n-2}, \dots, A_1)$$

- Probabilistic graphic models (PGM)
  - Bayesian network: BayesNet and BayesCard
    - Both structural learning and model inference are NP-Hard problems

$$\Pr(A) = \prod_i \Pr(A_i | A_{pa(i)})$$

A	B	C	E	R
y	n	n	y	y
n	n	n	n	y
y	y	y	n	n
...				

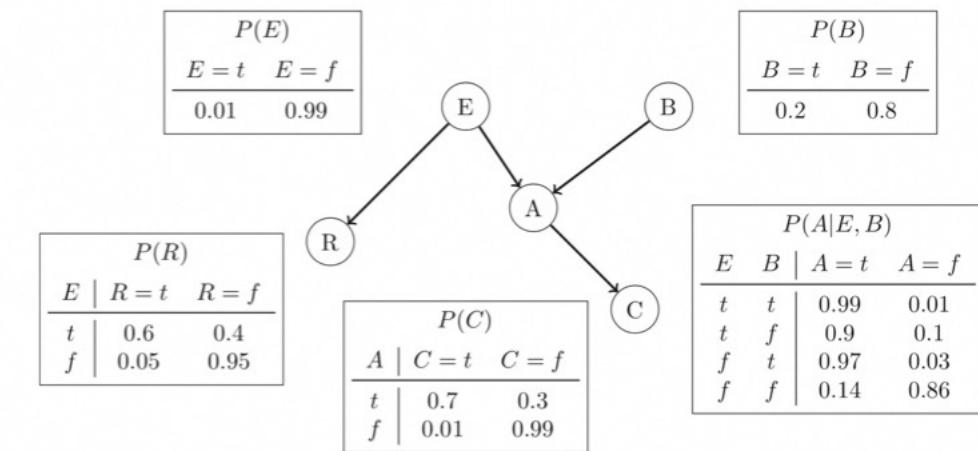


Figure from [A survey on Bayesian network structure learning from data, Progress in Artificial Intelligence, 2019]

(BayesNet) Lightweight Graphical Models for Selectivity Estimation Without Independence Assumptions, VLDB 2011

(BayesCard) BayesCard: A Unified Bayesian Framework for Cardinality Estimation, ArXiv, 2021

# Data-Driven Methods

- Different probabilistic graphic models (PGM)
  - Sum-Product network (SPN): DeepDB
    - Model training: split data or attribute
      - Split the data by rows to find local independent attributes
      - Split attributes (columns) and build histograms on local data
    - Model inference: bottom up, linear cost to the node size
    - Accuracy, speed, and model size: sensitive to attribute correlations

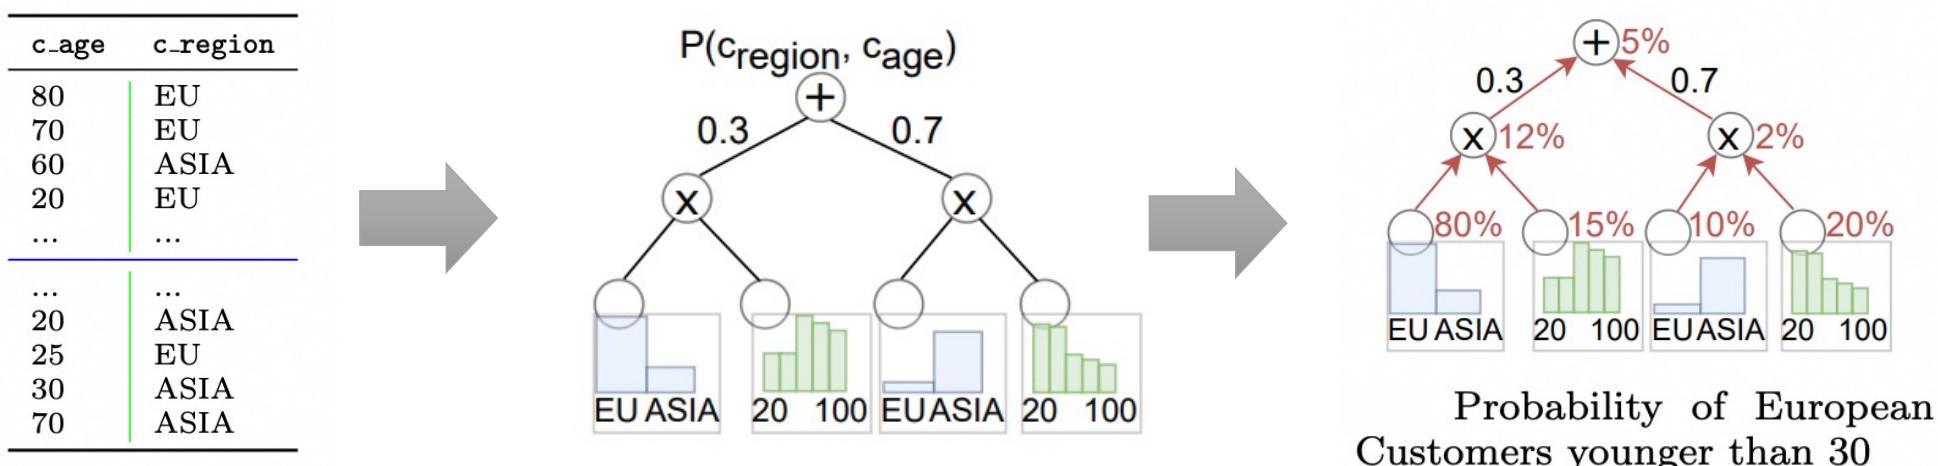
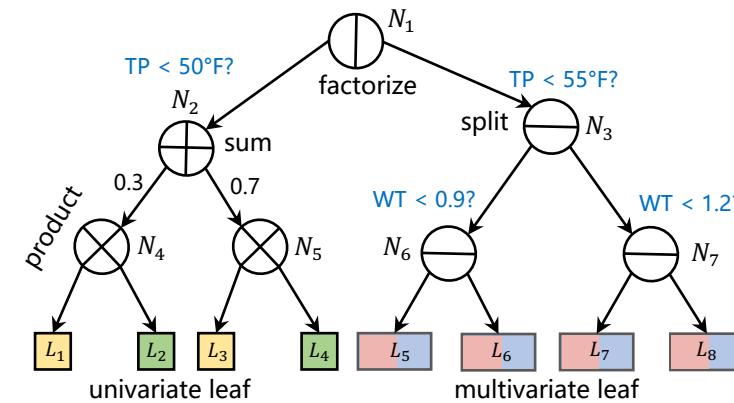


Figure from [DeepDB: Learn from Data, not from Queries!, VLDB 2020]

# Data-Driven Methods

- Different probabilistic graphic models (PGM)
  - FSPN: FLAT
  - Adaptively process highly and weakly correlated attributes
    - Factorize node: losslessly split highly and weakly correlated attributes
    - SPN: model weakly correlated attributes
  - Combining the advantages from BN and SPN, no assumptions on correlations
  - High accuracy and small model size, low Inference time

W = { WT, TP }		H = { WH, WF }	
weakly correlated		highly correlated	
WT (NTU)	TP (°F)	WH (m)	WF (m/s)
1.27	43.7	0.9	5.5
0.63	46.3	3.7	16.2
0.55	51.6	3.2	14.9
...	...	...	...
2.61	62.8	1.5	8.9
1.18	60.5	4.7	18.9
0.97	73.6	2.1	11.7



# Data-Driven Methods

- Other techniques
  - FactorJoin:  
histogram on join -> factor graph
  - FACE: normalizing flow
  - Iris: pretraining summarization models

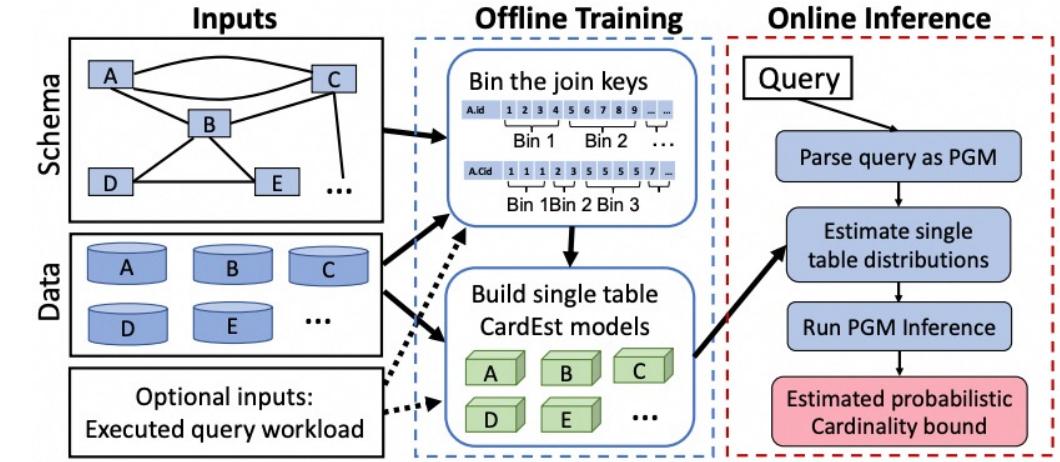


Figure from [FactorJoin: A New Cardinality Estimation Framework for Join Queries, SIGMOD 2023]

(FactorJoin) FactorJoin: A New Cardinality Estimation Framework for Join Queries, SIGMOD 2023

(FACE) FACE: A Normalizing Flow based Cardinality Estimator, VLDB 2021

(Iris) Pre-training Summarization Models of Structured Datasets for Cardinality Estimation, VLDB 2021

# Learned Methods: An Overview

- Data-driven methods:
  - Learned models approximating  $\Pr(A)$  of table  $T$  on attributes  $A$
  - $C(q) = \Pr(q) * |T|$  for  $q \subseteq T$
- Query-driven methods: for a workload (distribution)  $Q$ 
  - Learned models as a function  $C(q)$  for  $q \in Q$
- Hybrid methods:
  - Utilize both data and query information

# Query-Driven Methods

- Query featurization: from query to vectors
  - Tables
  - Join Conditions
  - Predicates: Columns, Values and Operators

```
SELECT COUNT(*) FROM title t, movie_companies mc WHERE t.id = mc.movie_id AND t.production_year > 2010 AND mc.company_id = 5
```

Table set {[0 1 0 1 ... 0], [0 0 1 0 ... 1]}      Join set {[0 0 1 0]}      Predicate set {[1 0 0 0 1 0 0 0.72], [0 0 0 1 0 0 1 0 0.14]}

table id                    samples                    join id                    column id            value                    operator id

Figure from [Learned Cardinalities: Estimating Correlated Joins with Deep Learning, CIDR 2019]

- Build supervised model mapping  $q$  to  $C(q)$ 
  - Statistical models
  - DNN-based models

# Query-Driven Methods

- Statistical models:
  - Linear model
  - Tree-based Ensembles
  - XGBoost
  - QuickSel: a mixture of models with overlaps

(Linear) A Black-Box Approach to Query Cardinality Estimation, CIDR 2007

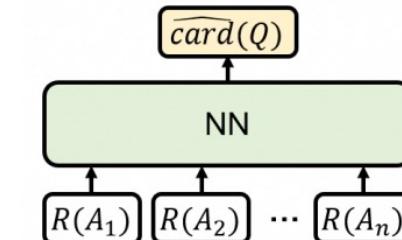
(Tree) Selectivity Estimation for Range Predicates using Lightweight Models, VLDB 2019

(XGBoost) Approximating Selectivity Functions using Low Overhead Regression Models, VLDB 2020

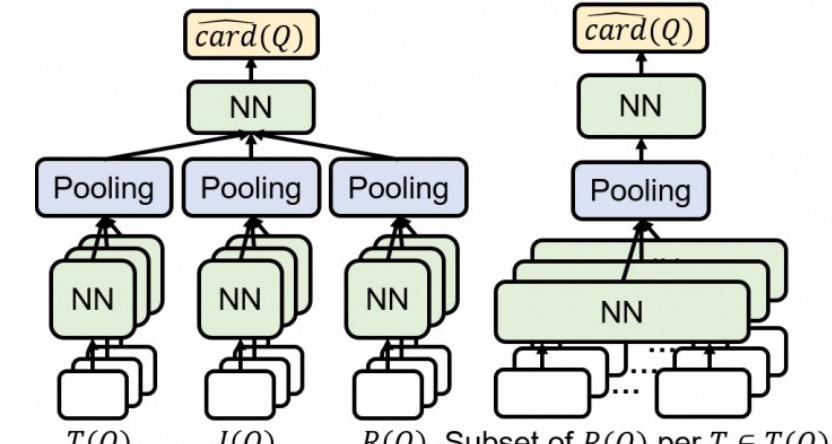
(QuickSel) QuickSel: Quick Selectivity Learning with Mixture Models, SIGMOD 2020

# Query-Driven Methods

- DNN-based models:
  - FCN: fully connected neural network
    - FCN+Pool: adding pooling layers
  - MSCN: multi-set convolutional network
    - CRN: learning containment rate between a pair of queries
    - Robust-MSCN: query masking to improve generalization ability



(a) FCN.



(b) MSCN.

(c) FCN+Pool.

Figure from [Learned Cardinality Estimation: An In-depth Study, SIGMOD 2022]

(FCN) Cardinality estimation using neural networks, CASON 2015

(FCN+Pool) Learned Cardinality Estimation: An In-depth Study, SIGMOD 2022

(MSCN) Learned Cardinalities: Estimating Correlated Joins with Deep Learning, CIDR 2019

(CRN) Improved Cardinality Estimation by Learning Queries Containment Rates. EDBT 2020

(Robust-MSCN) Robust Query Driven Cardinality Estimation under Changing Workloads, VLDB 2023

# Query-Driven Methods

- Advanced techniques
  - Fauce: an ensemble of deep models
    - Learn cardinality and uncertainty together
  - GL+: DNNs + segmentation techniques
    - Resolve the data hungry problem
  - NNGP: Bayesian deep learning models
  - LPCE: re-optimization framework

(Fauce) Fauce: Fast and Accurate Deep Ensembles with Uncertainty for Cardinality Estimation, VLDB 2021

(GL+) Learned Cardinality Estimation for Similarity Queries, SIGMOD 2021

(NNGP) Lightweight and Accurate Cardinality Estimation by Neural Network Gaussian Process, SIGMOD 2022

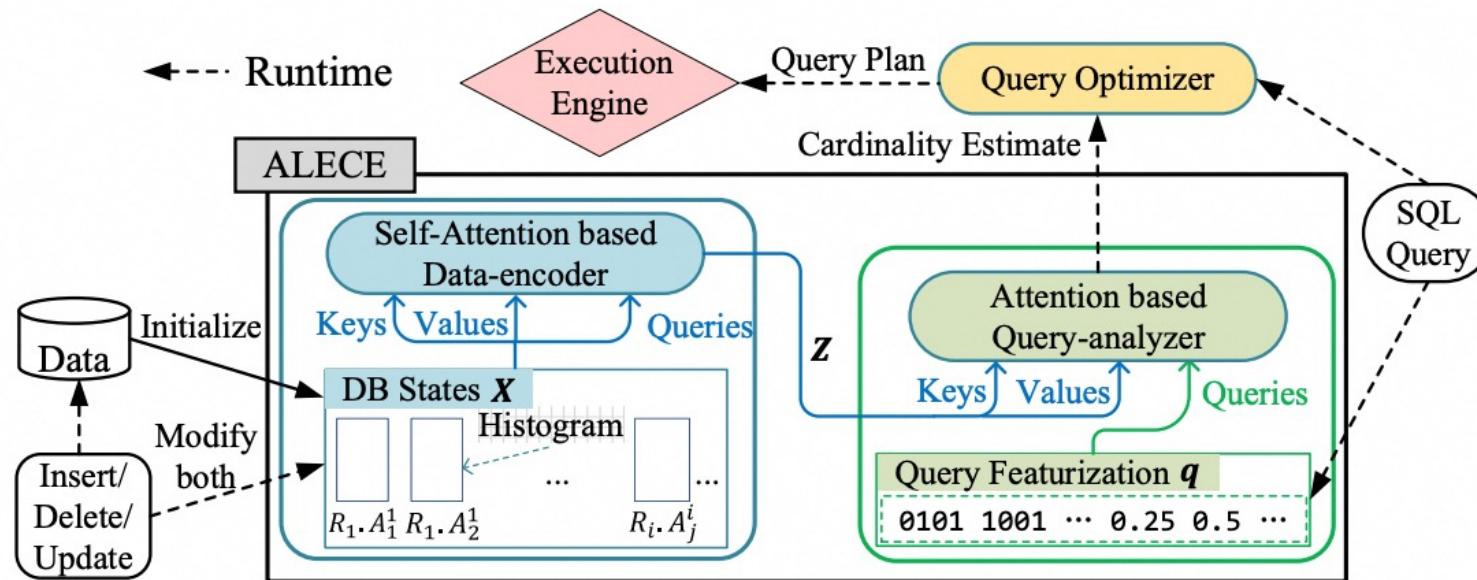
(LPCE) Speeding Up End-to-end Query Execution via Learning-based Progressive Cardinality Estimation, SIGMOD 2023

# Learned Methods: An Overview

- Data-driven methods:
  - Learned models approximating  $\Pr(A)$  of table  $T$  on attributes  $A$
  - $C(q) = \Pr(q) * |T|$  for  $q \subseteq T$
- Query-driven methods: for a workload (distribution)  $Q$ 
  - Learned models as a function  $C(q)$  for  $q \in Q$
- Hybrid methods:
  - Utilize both data and query information

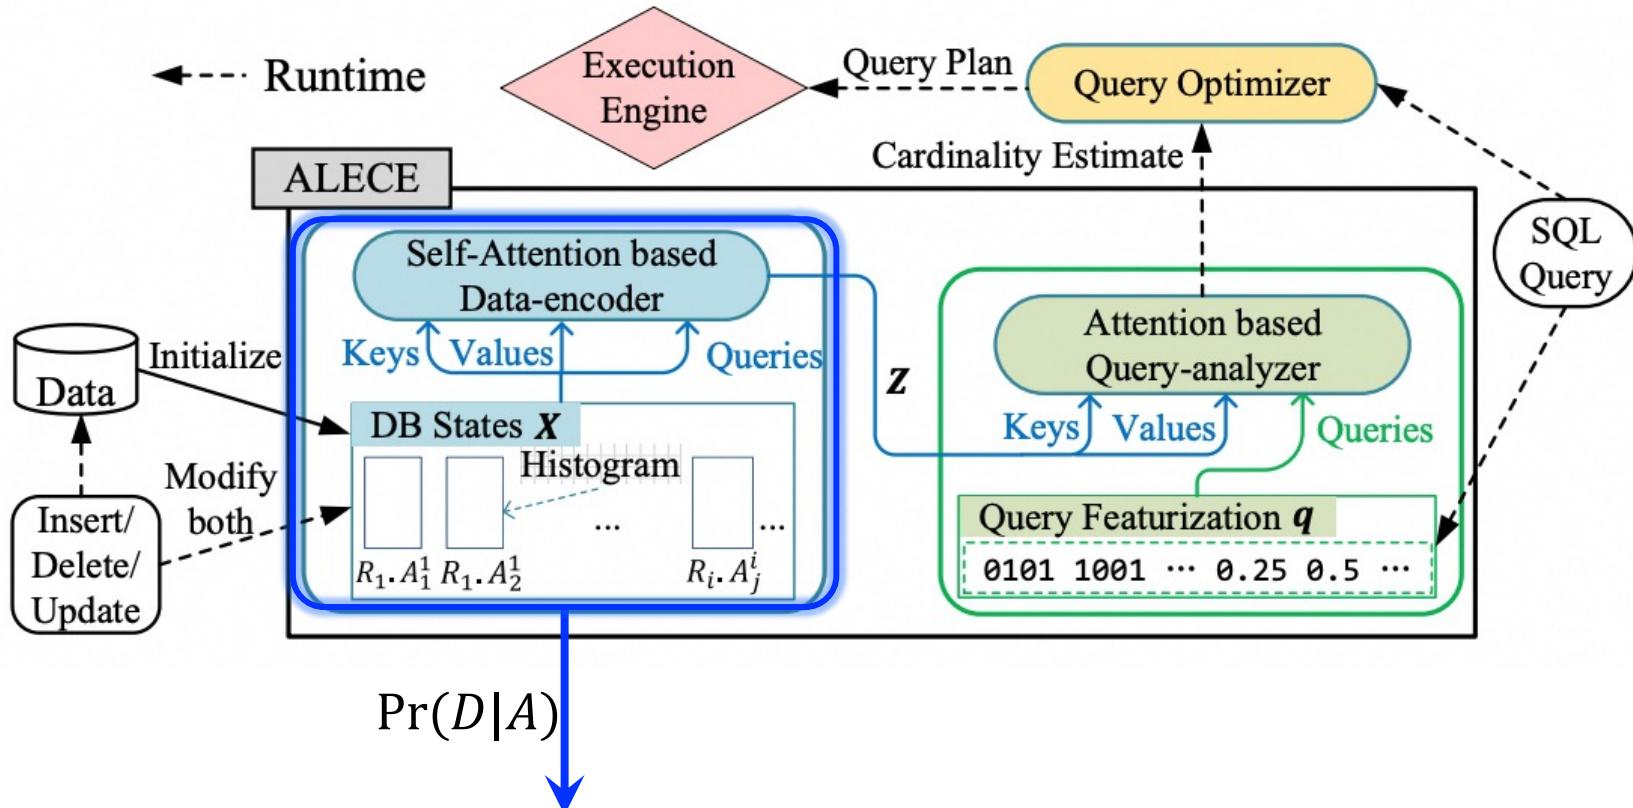
# Hybrid Methods

- ALECE: discover the hidden relationships between queries and underlying data using attention + transformer models
  - Data features (e.g., histogram of each attribute) + Query features
  - Process static and dynamic data in the same way



# Hybrid Methods

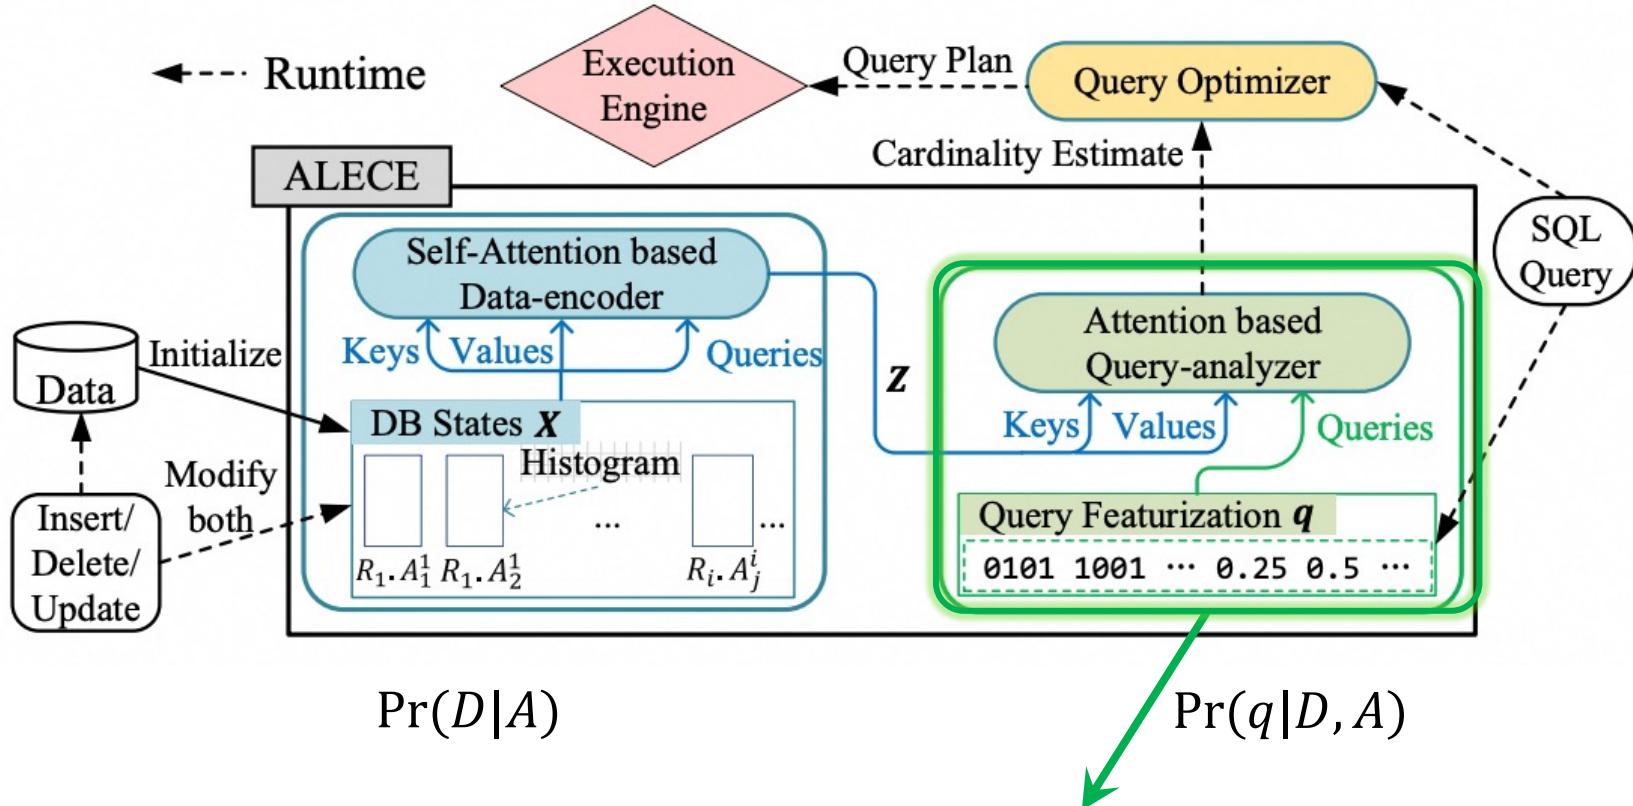
- ALECE: utilize both query and data information



**Data Encoder:** an attention module learn the information of high-order data distributions using low-level histograms on each attribute

# Hybrid Methods

- ALECE: utilize both query and data information



Query Analyzer: an attention module learn the joint information between data distribution and featuralized queries

# Benchmark Evaluations

- Cardinality estimation on single table
  - Analyze whether these methods are suitable for deploying
- End-to-end evaluation of query optimization/processing
  - Guidance to select methods in various practical scenarios
- New benchmarks and comprehensive evaluations
  - New dataset STATS: real-world data with complicated distributions
  - End-to-End evaluation on actual DBMS (PostgreSQL)

(Single Table) Are We Ready For Learned Cardinality Estimation?, VLDB 2021

(Space Exploration) Learned Cardinality Estimation: A Design Space Exploration and A Comparative Evaluation, VLDB 2021

(STATS) Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation, VLDB 2021

# Benchmark Evaluations

- Learned methods could help, but not always

Category	Method	Workload					
		JOB-LIGHT			STATS-CEB		
		End-to-End Time	Exec. + Plan Time	Improvement	End-to-End Time	Exec. + Plan Time	Improvement
Baseline	PostgreSQL	3.67h	3.67h + 3s	0.0%	11.34h	11.34h + 25s	0.0%
	<b>TrueCard</b>	<b>3.15h</b>	<b>3.15h + 3s</b>	<b>14.2%</b>	<b>5.69h</b>	<b>5.69h + 25s</b>	<b>49.8%</b>
Traditional	MultiHist	3.92h	3.92h + 30s	-6.8%	14.55h	14.53h + 79s	-28.3%
	UniSample	4.87h	4.84h + 96s	-32.6%	> 25h	--	--
	WJSample	4.15h	4.15h + 23s	-13.1%	19.86h	19.85h + 45s	-75.0%
	PessEst	3.38h	3.38h + 11s	7.9%	6.10h	6.10h + 43s	46.2%
Query-driven	MSCN	3.50h	3.50h + 12s	4.6%	8.13h	8.11h + 46s	28.3%
	LW-XGB	4.31h	4.31h + 8s	-17.4%	> 25h	--	--
	LW-NN	3.63h	3.63h + 9s	1.1%	11.33h	11.33h + 34s	0.0%
	UAE-Q	3.65h	3.55h+356s	-1.9%	11.21h	11.03h+645s	1.1%
Data-driven	NeuroCard <sup>E</sup>	3.41h	3.29h + 423s	6.8%	12.05h	11.85h + 709s	-6.2%
	BayesCard	3.18h	3.18h + 10s	13.3%	7.16h	7.15h + 35s	36.9%
	DeepDB	3.29h	3.28h + 33s	10.3%	6.51h	6.46h + 168s	42.6%
	FLAT	3.21h	3.21h + 15s	12.9%	5.92h	5.80h + 437s	47.8%
Query + Data	UAE	3.71h	3.60h + 412s	-2.7%	11.65h	11.46h + 710s	-0.02%

# Benchmark Evaluations: Dynamic Settings

Data	Model	Insert-heavy												Update-heavy											
		E2E Time(S)	Q-error				P-error				Size (MB)	Building Time(Min)	Latency (ms)	E2E Time(S)	Q-error				P-error						
			50%	90%	95%	99%	50%	90%	95%	99%					50%	90%	95%	99%	50%	90%	95%	99%			
STATS	PG	7,790	190	$1.4 \cdot 10^5$	$1.1 \cdot 10^6$	$1.8 \cdot 10^7$	2.60	25.44	41.23	243	-	-	-	4,337	524	43,096	$3.7 \cdot 10^5$	$1.1 \cdot 10^7$	1.78	19.52	29.01	178			
	Uni-Samp	6,646	1.35	12.47	$>10^{10}$	$>10^{10}$	2.71	16.77	21.65	201	3.34	0.02	239	6,002	1.27	$9.3 \cdot 10^9$	$>10^{10}$	$>10^{10}$	1.10	12.86	33.19	153.11			
	NeuroCard	>30,000	17.37	1,388	7,402	$3.0 \cdot 10^5$	2.38	39.18	824	84,415	90.76	23.69	32.42	22,193	18.49	1,252	6,784	$3.1 \cdot 10^5$	2.10	31.75	347	11,730			
	FLAT	>30,000	12.77	1,979	12,897	$8.6 \cdot 10^5$	2.22	70.08	2,202	8,738	210.33	53.77	13.67	24,319	16.00	2,376	17,811	$9.4 \cdot 10^5$	2.20	29.76	724	$1.4 \cdot 10^5$			
	FactorJoin	>30,000	22.62	2,593	31,936	$1.6 \cdot 10^6$	2.35	66.12	876	8,384	1.64	0.42	1.33	>30,000	23.76	2,939	36,198	$2.4 \cdot 10^6$	2.75	115	912	4,409			
	MLP	>30,000	$2.4 \cdot 10^6$	$>10^{10}$	$>10^{10}$	$>10^{10}$	7.35	491	2,294	6,887	8.63	3.11	3.17	>30,000	$2.3 \cdot 10^6$	$>10^{10}$	$>10^{10}$	$>10^{10}$	6.66	502	3,025	11,584			
	MSCN	27,758	20.09	2,870	17,037	$2.6 \cdot 10^5$	2.46	87.23	736	8,738	1.61	11.87	1.02	25,766	18.89	2,382	16,947	$1.8 \cdot 10^6$	2.56	64.64	246	46,721			
	NNGP	12,883	9.88	827	4,652	$2.6 \cdot 10^5$	1.41	5.31	27.33	3,588	32.44	0.89	32.66	11,985	8.15	567	3,639	$2.5 \cdot 10^5$	1.52	5.62	13.81	2,753			
	ALECE	2,901	1.29	5.07	11.54	62.52	1.07	1.34	1.59	2.32	22.31	9.25	8.25	2,402	1.42	5.06	10.40	44.69	1.08	1.45	2.13	2.86			
	Optimal	2,871	1	1	1	1	1	1	1	1	-	-	-	2,349	1	1	1	1	1	1	1	1			
Job-light	PG	6,128	1.59	7.14	20.85	223	1.10	1.32	1.48	2.89	-	-	-	4,747	1.67	6.20	14.09	70.13	1.10	1.39	1.64	3.06			
	Uni-Samp	6,401	1.23	2.59	6.53	$>10^{10}$	1.08	1.85	2.79	146	23.95	0.52	163	4,491	1.33	3.61	7.50	$>10^{10}$	1.12	1.71	2.26	3.67			
	DeepDB	26,636	11.31	663	6,328	$5.4 \cdot 10^5$	1.66	13.71	26.01	49.57	11.52	11.93	11.92	23,190	10.64	628	4,883	$8.8 \cdot 10^5$	1.50	11.81	27.71	135			
	NeuroCard	27,744	14.73	978	6,766	$4.7 \cdot 10^5$	1.81	12.40	37.25	144	42.93	9.11	15.90	16,553	13.49	1,000	6,489	$1.5 \cdot 10^6$	1.60	15.10	25.94	50.68			
	FLAT	23,876	11.22	626	5,786	$4.4 \cdot 10^5$	1.66	13.98	27.26	89.67	11.10	17.83	3.96	22,531	10.13	490	4,221	$1.1 \cdot 10^6$	1.42	12.87	34.54	195			
	FactorJoin	16,665	25.14	3,456	27,938	$2.5 \cdot 10^6$	1.52	5.45	10.23	47.09	4.66	26.22	1.33	13,071	27.31	6,134	36,155	$2.2 \cdot 10^6$	1.42	4.83	9.43	22.75			
	MLP	5,610	4.30	23.14	59.37	2,411	1.13	1.70	2.10	3.12	6.19	2.68	2.28	5,351	5.59	62.02	138	1,441	1.17	2.02	2.88	4.09			
	MSCN	26,665	24.42	1,809	9,217	$1.4 \cdot 10^5$	2.07	20.98	28.58	106	1.56	33.37	0.77	29,118	26.77	1,448	7,980	$1.5 \cdot 10^5$	2.59	19.87	61.18	257			
	NNGP	15,430	$4.5 \cdot 10^8$	$>10^{10}$	$>10^{10}$	$>10^{10}$	3.57	8.39	11.97	24.47	32.44	0.83	29.86	4,959	4.99	59.86	190	50,036	1.20	2.05	2.99	4.96			
	ALECE	5,168	1.14	2.08	3.85	14.82	1.07	1.24	1.32	1.77	22.25	7.73	7.28	3,839	1.09	1.93	3.63	16.28	1.07	1.23	1.29	1.57			
	Optimal	5,063	1	1	1	1	1	1	1	1	-	-	-	3,792	1	1	1	1	1	1	1	1			
TPCH	PG	3,230	1.35	4.88	7.90	16.16	1.44	1.57	1.59	1.62	-	-	-	2,385	1.33	5.01	6.87	14.64	1.41	1.56	1.59	1.62			
	Uni-Samp	>30,000	1.22	2.94	$>10^{10}$	$>10^{10}$	1.81	2.22	45.60	514	23.29	0.17	24.52	>30,000	1.30	3.41	$>10^{10}$	$>10^{10}$	1.75	19.35	50.34	211			
	NeuroCard	10,616	43.43	7,071	27,735	$3.5 \cdot 10^5$	3.75	10.91	22.48	106	704.98	34.38	34.77	6,562	45.73	5,896	20,931	$1.2 \cdot 10^5$	3.16	11.04	16.49	245			
	FLAT	11,428	46.04	12,236	70,383	$8.4 \cdot 10^7$	3.56	10.91	25.71	1,182	168.65	41.42	10.06	7,864	42.32	7,328	26,470	$3.7 \cdot 10^5$	3.56	15.82	21.70	425			
	MLP	7,261	51,475	$1.4 \cdot 10^7$	$3.5 \cdot 10^7$	$4.7 \cdot 10^9$	4.71	14.11	22.74	36.09	8.63	3.57	3.62	8,342	79,317	$3.7 \cdot 10^7$	$9.1 \cdot 10^7$	$>10^{10}$	5.12	17.63	32.74	63.39			
	MSCN	7,718	36.36	6,709	26,744	$4.4 \cdot 10^5$	3.58	15.11	25.71	387	1.58	23.43	0.79	8,353	41.66	5,457	21,433	$1.7 \cdot 10^5$	3.80	14.86	25.69	472			
	NNGP	18,618	$1.3 \cdot 10^9$	$>10^{10}$	$>10^{10}$	$>10^{10}$	5.56	47.40	59.18	99.55	32.44	0.92	41.76	3,128	25.88	432	986	4,968	1.54	7.98	10.90	20.66			
	ALECE	3,233	1.15	2.29	3.37	8.86	1.43	1.57	1.60	1.62	22.34	11.62	10.35	2,380	1.25	2.11	2.95	10.87	1.41	1.56	1.59	1.62			
	Optimal	3,227	1	1	1	1	1	1	1	1	-	-	-	2,378	1	1	1	1	1	1	1	1			

Experimental data from [ALECE: An Attention-based Learned Cardinality Estimator for SPJ Queries on Dynamic Workloads, VLDB 2023]

# Part 2: Learn to Cost Plans

- Cardinality & Cost
- Learn to Estimate Cardinality
- Learned Cost Model

# Learned Cost Model

- Cost: an approximation to execution time/resource usage
  - Traditional cost model: expert-tuned formulas on cardinality

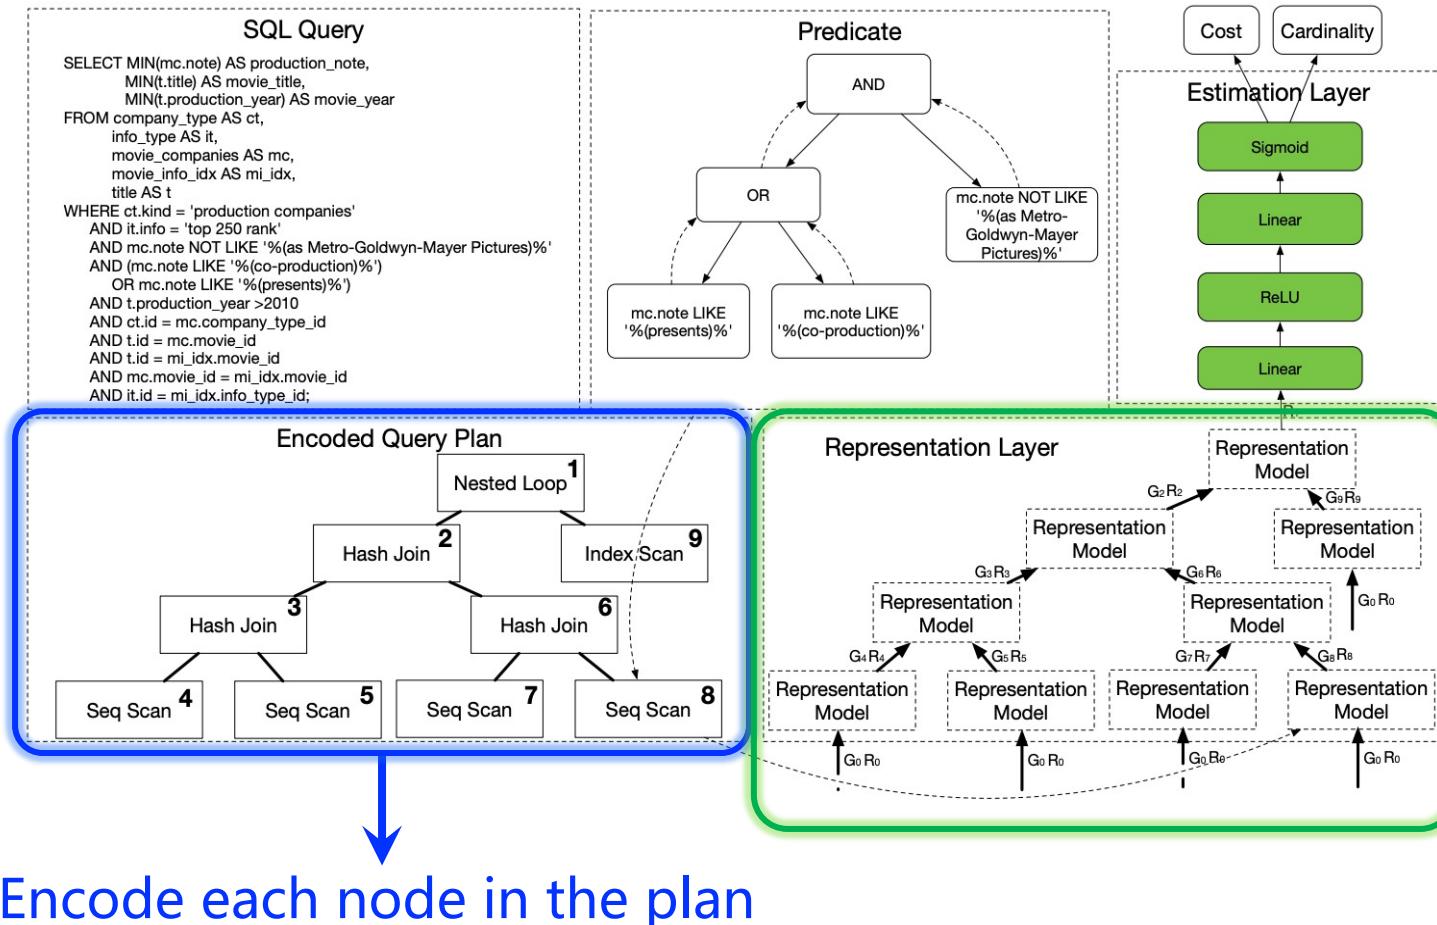
$$\text{cost}_J(s_1, s_2) = a_{2,0} \cdot s_1^2 + a_{0,2} \cdot s_2^2 + a_{1,1} \cdot s_1 s_2 + \\ a_{1,0} \cdot s_1 + a_{0,1} \cdot s_2 + a_{0,0},$$

- Learned methods
  - Costing a single plan:  
use previous plans to train a cost model
  - Costing plans for concurrent queries:  
consider the correlations between plans and predict the cost

# Cost Estimation for Single Plans

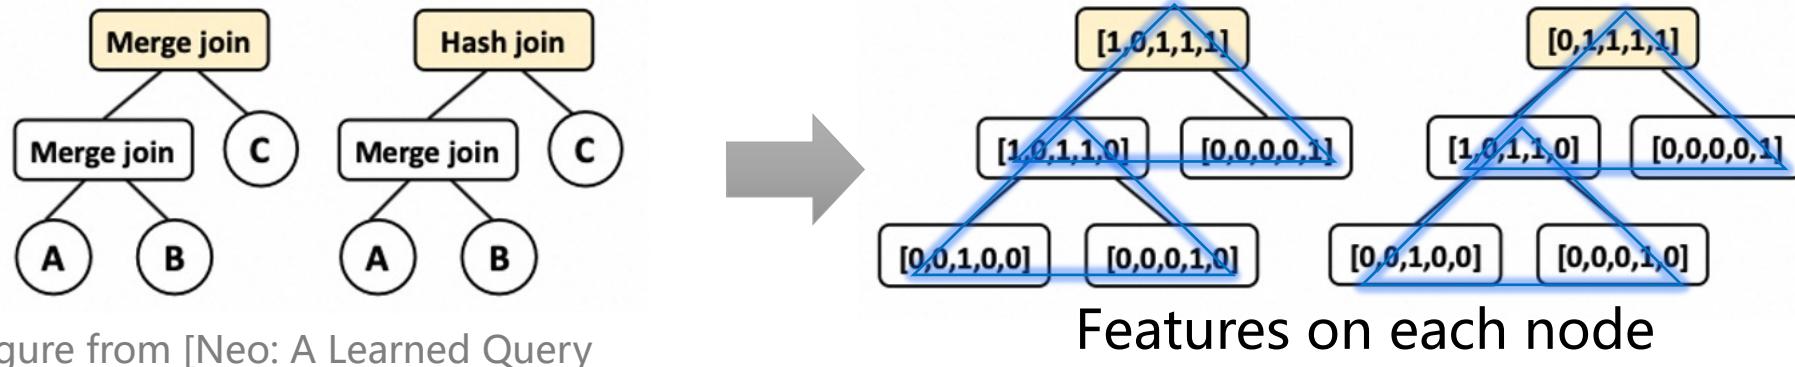
A typical LSTM model  
 $(h_t, m_t) = \text{LSTM}(x_p, h_{t-1}, m_{t-1})$

- Tree LSTM: plan as a sequence



# Cost Estimation for Single Plans

- Tree convolutional network (TCN): plan as a tree
    - Triangle sliding window on each node and its two children

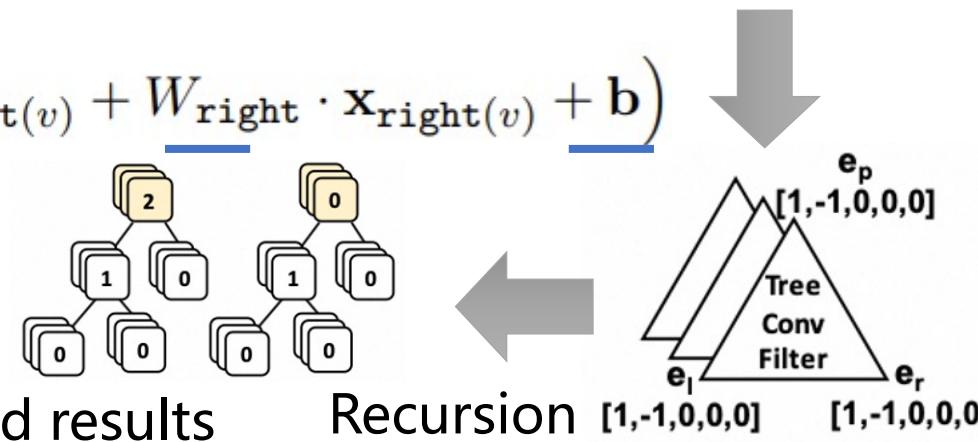


## Figure from [Neo: A Learned Query Optimizer, VLDB 2019]

$$\mathbf{x}'_v = \sigma \left( W_{\text{op}} \cdot \mathbf{x}_v + W_{\text{left}} \cdot \mathbf{x}_{\text{left}(v)} + W_{\text{right}} \cdot \mathbf{x}_{\text{right}(v)} + \mathbf{b} \right)$$

Additional layers on final output to predict the cost

Tree Convolution Filter:  
**shared model**  
**parameters** on each  
sliding window



# Cost Estimation for Single Plans

- Other ways to capture the tree structure information
  - Saturn: traverse and compress plans as vectors
  - QueryFormer: a transformer to learn embeddings of plans, applicable to multiple downstream tasks

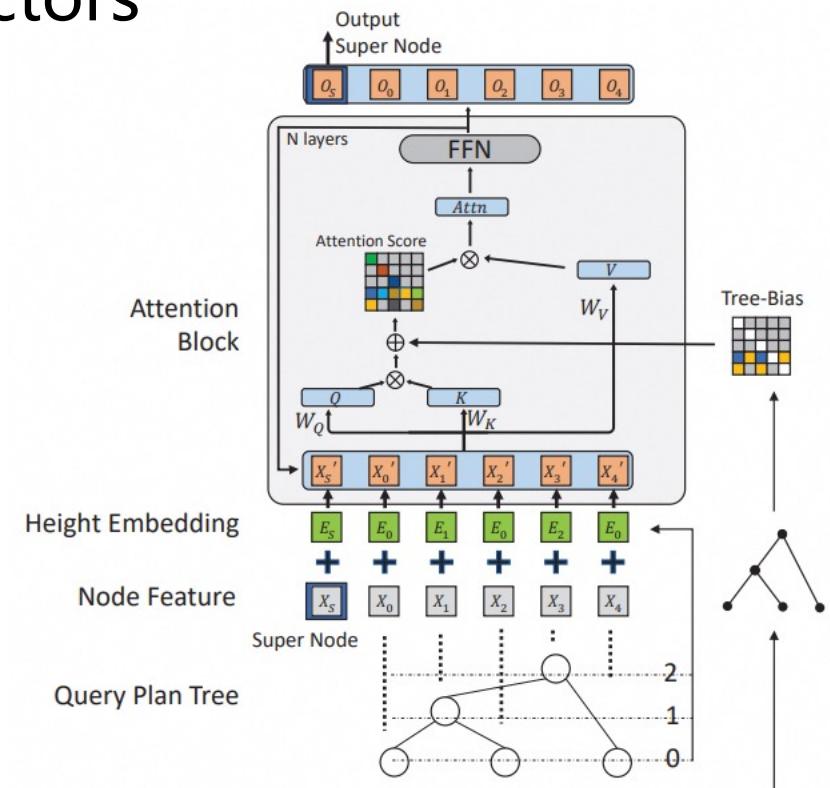


Figure from [QueryFormer: A Tree Transformer Model for Query Plan Representation, VLDB 2022]

# Cost Estimation for Plans of Concurrent Queries

- Challenges: queries have complex correlations
- Solutions: modeling the correlations using graphs
  - GPredictor: using graph neural networks for cost prediction

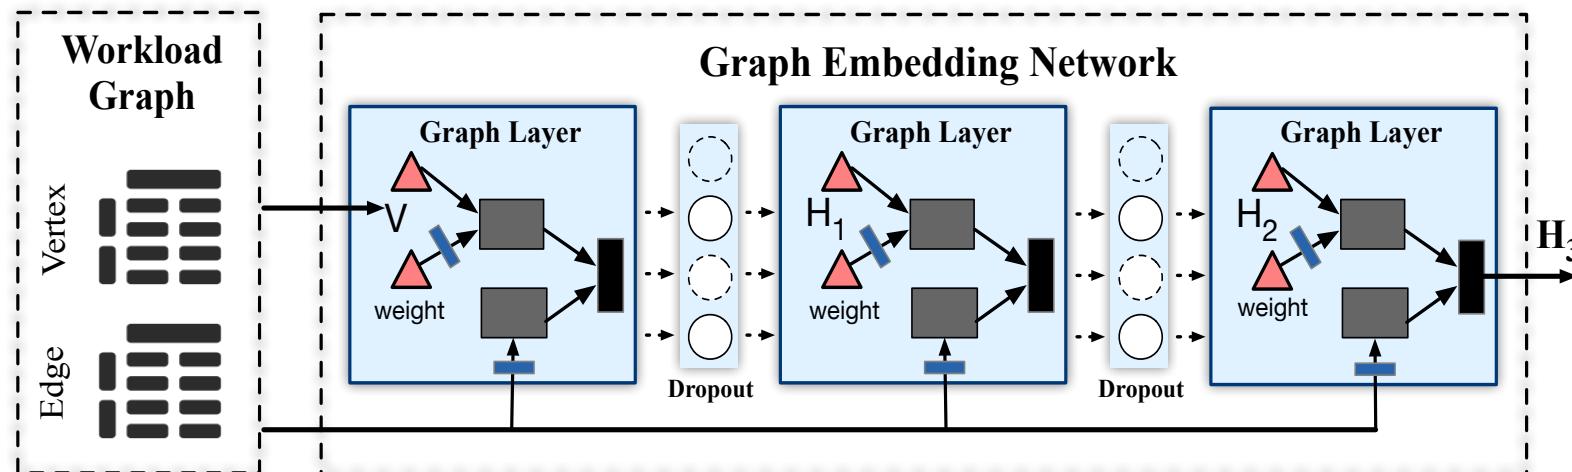


Figure from [Performance Prediction for Concurrent Queries using Graph Embedding, VLDB 2020]

# Cost Estimation for Plans of Concurrent Queries

- Prestroid: apply tree convolutional networks for concurrent queries
- RAAL: end-to-end deep cost model with attention mechanism applied in cloud environment

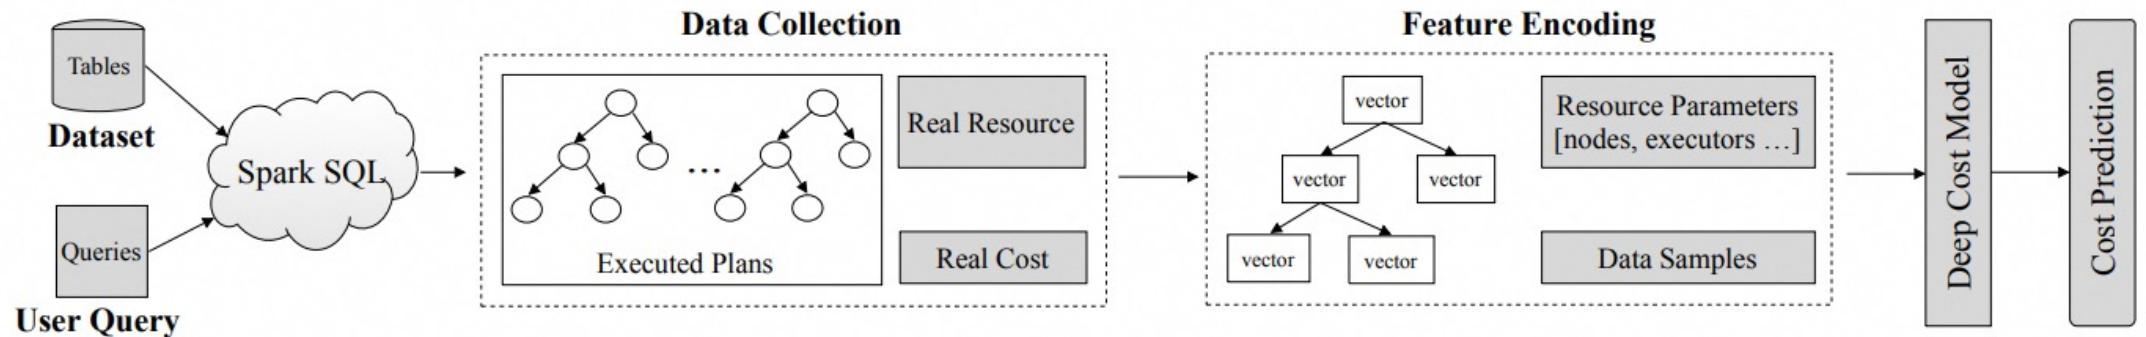


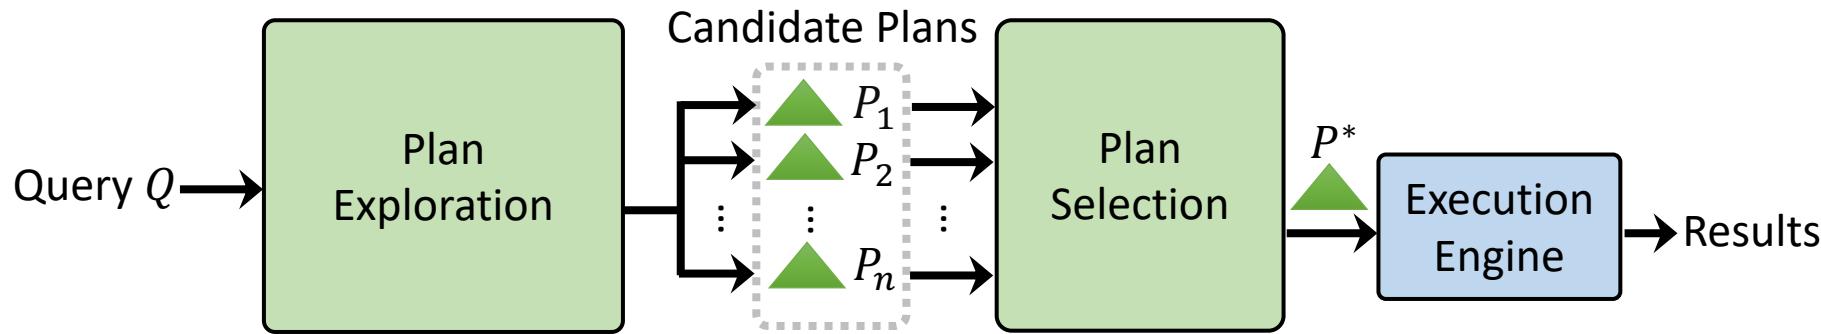
Figure from [A Resource-Aware Deep Cost Model for Big Data Query Processing, ICDE 2022]

# Part 3: Learn to Explore Plan Space

- Search by Value Networks
- Search by Steering Existing Query Optimizers
- Search and Rank Plans
- Quality Control of Learned Query Optimizers

# Plan Search in Query Optimizer

- The general framework



- Traditional plan exploration methods
  - Exact: Dynamic programming
  - Heuristic: Genetic or greedy search ...

# Part 3: Learn to Explore Plan Space

- Search by Value Networks
- Search by Steering Existing Query Optimizers
- Search and Rank Plans
- Quality Control of Learned Query Optimizers

# Search by Value Networks

- Value network: Predicate the best possible plan  $P$  containing partial plan  $P'$  of query  $Q$ 
  - Tree convolutional network or other models

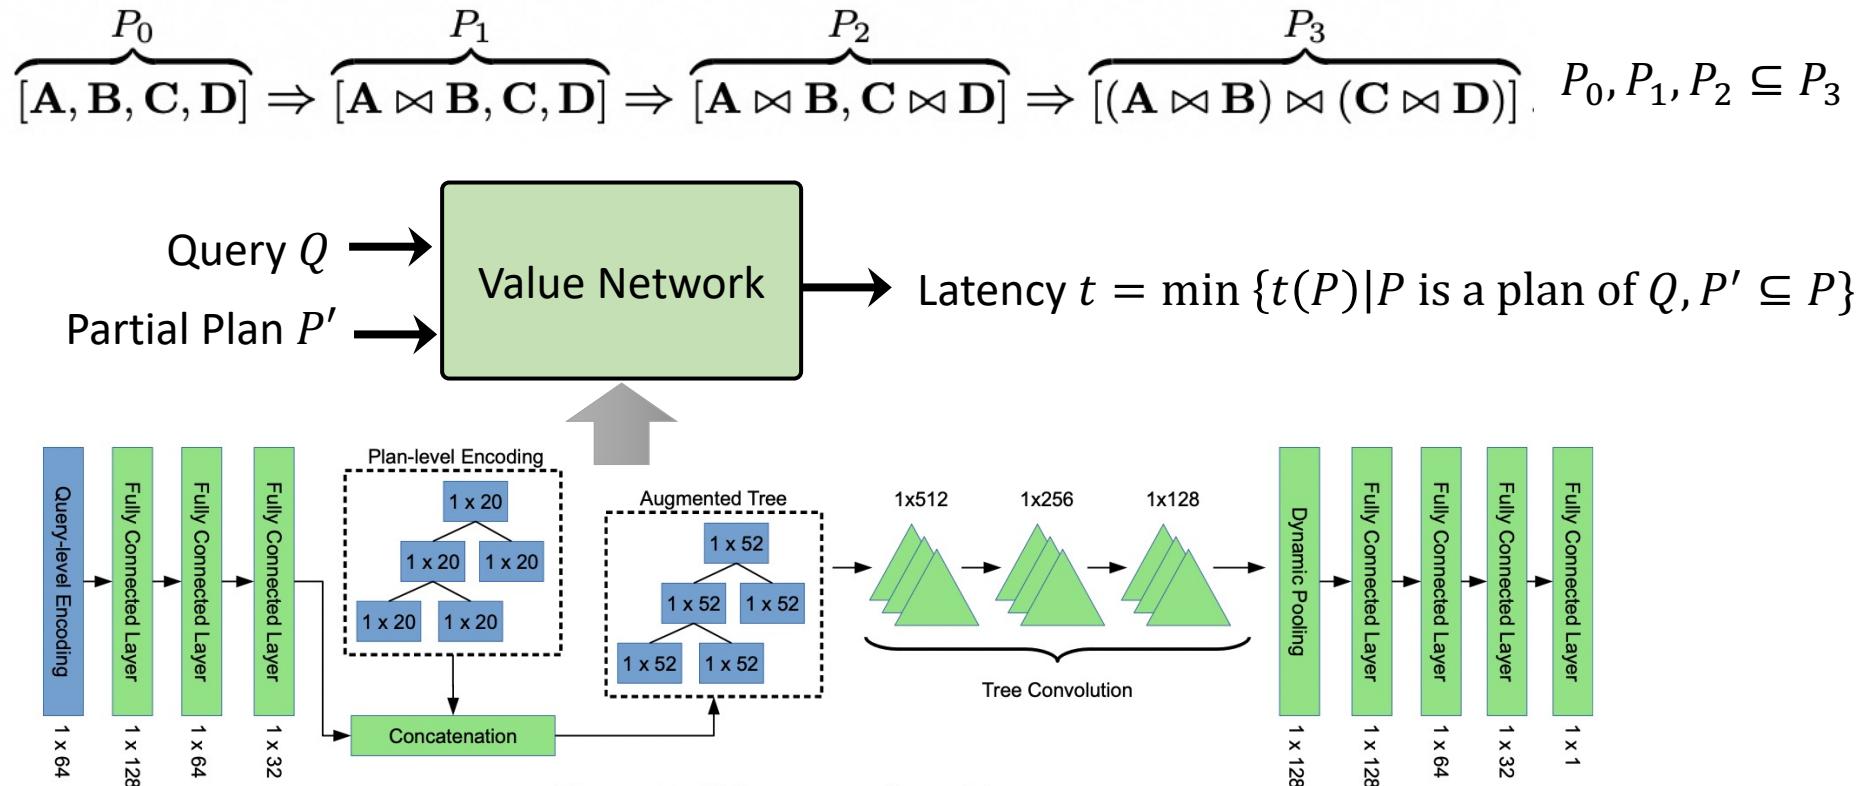
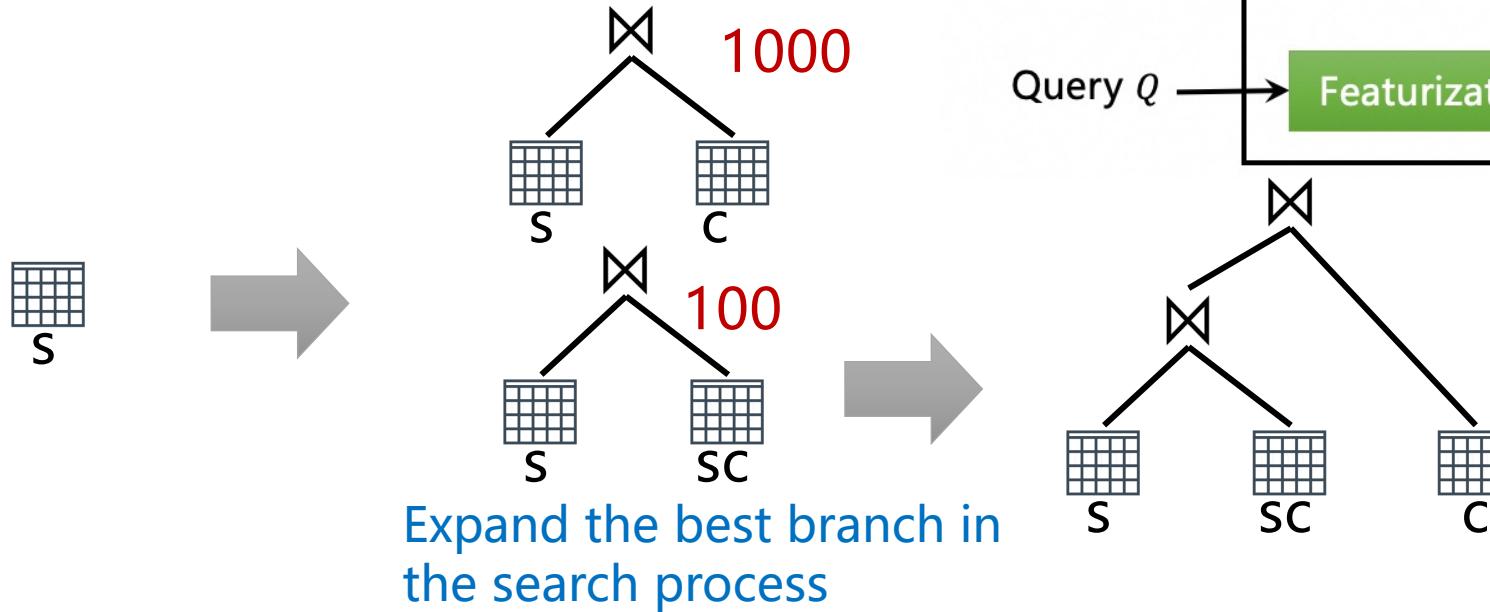
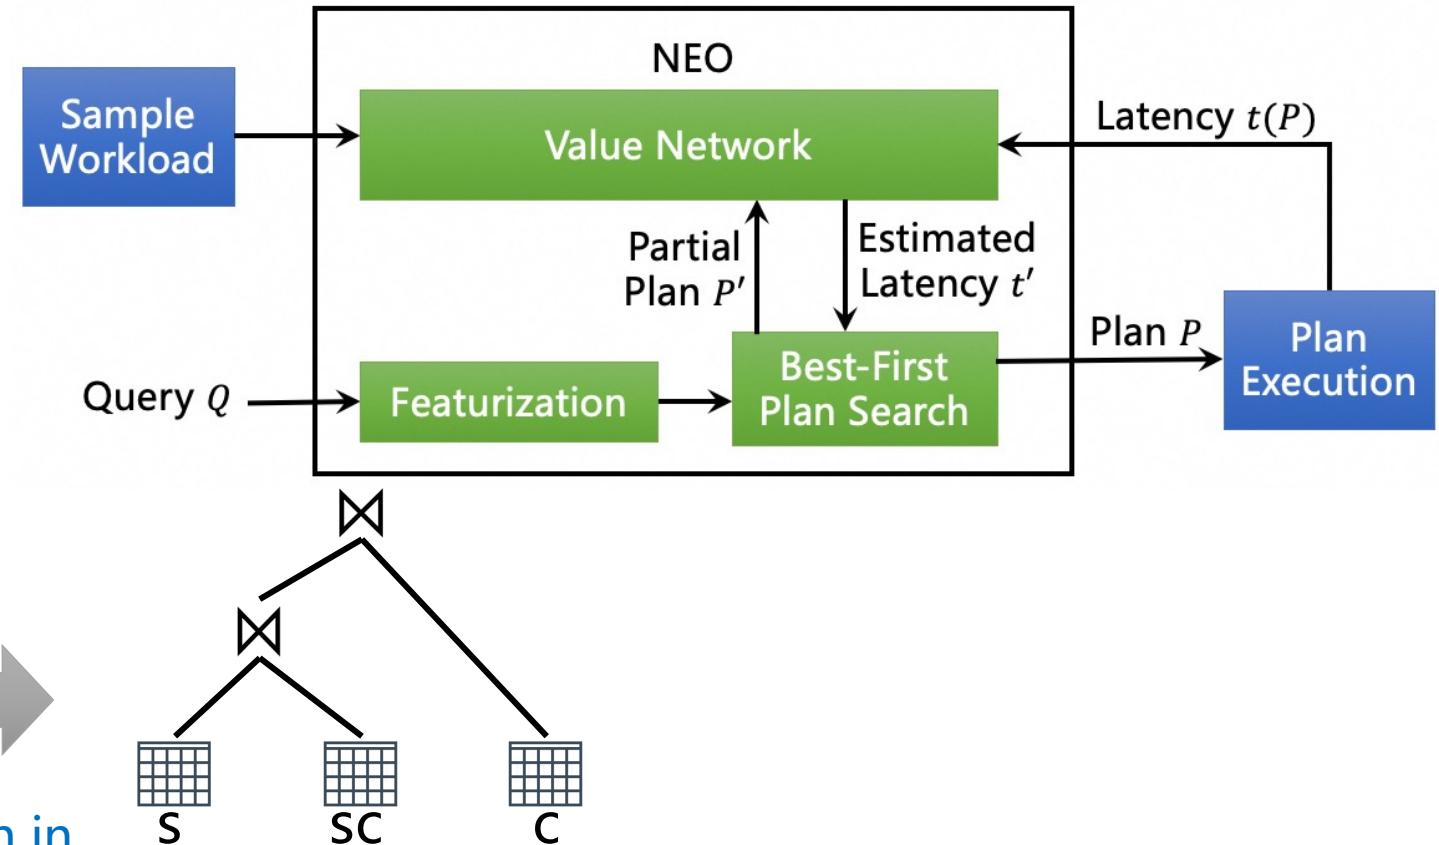


Figure from [Neo: A Learned Query Optimizer, VLDB 2019]

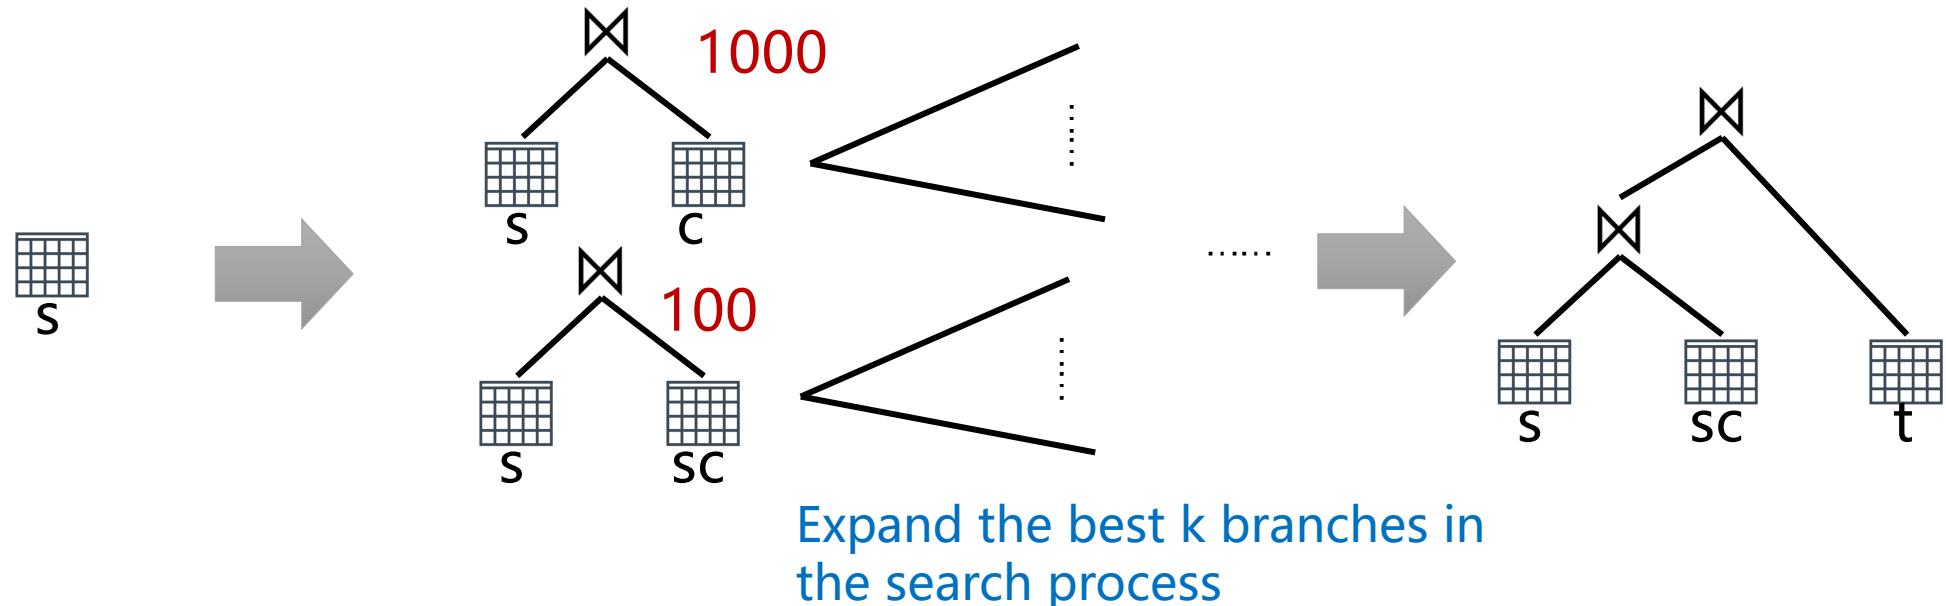
# Search by Value Networks

- Search plan space from scratch with value networks
  - Neo: best-first search



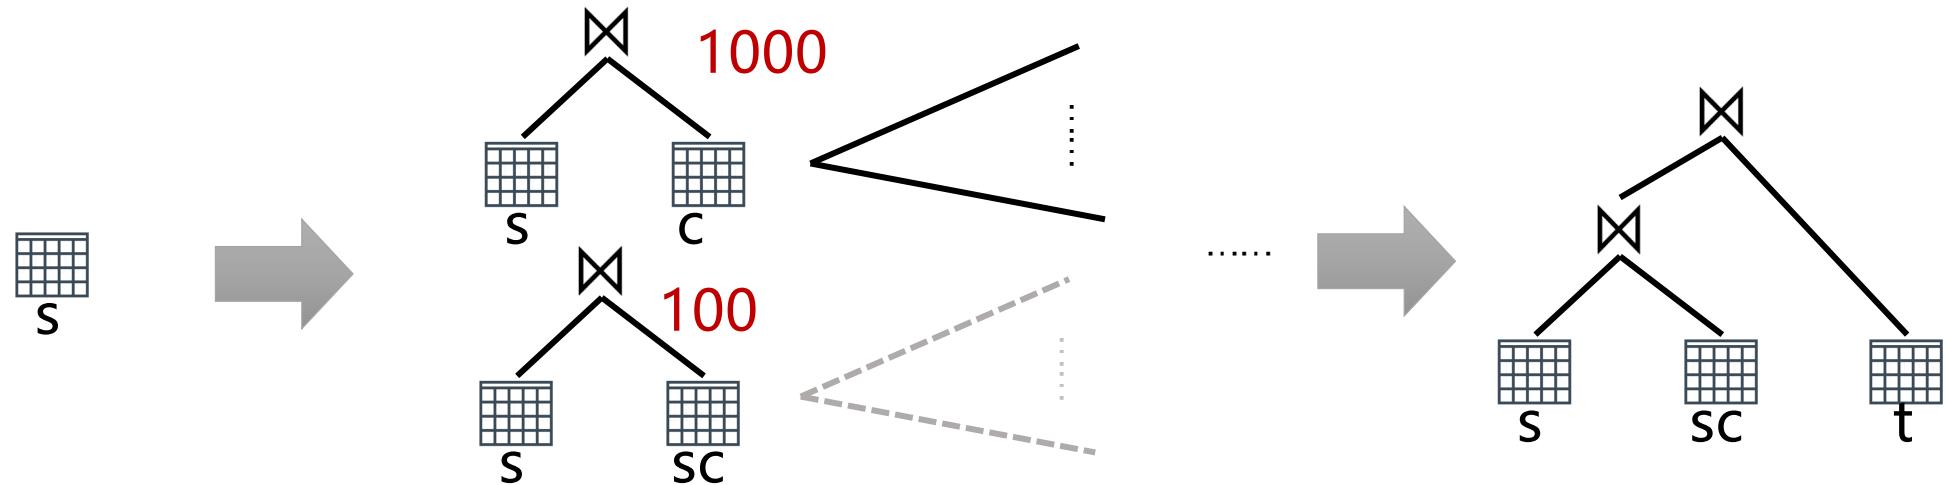
# Search by Value Networks

- Search plan space from scratch with value networks
  - Balsa: beam search



# Search by Value Networks

- Search plan space from scratch with value networks
  - LOGER:  $\epsilon$ -beam search



Expand the a random number (decided by  $k$  and  $\epsilon$ ) of branches in the search process

# Search by Value Networks

- Model training
  - Exact label is difficult to obtain
  - Using the latency of the best plan we can obtain as training label
  - Update model with collected workload in periodical
- Shortcomings:
  - Cold-start
    - Balsa/BASE: Bootstrap with existing cost model
  - Large search space
    - Not effective to obtain truly good plans

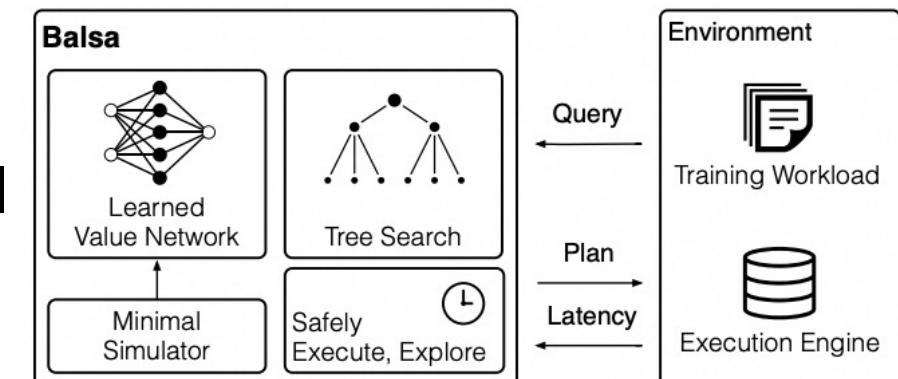
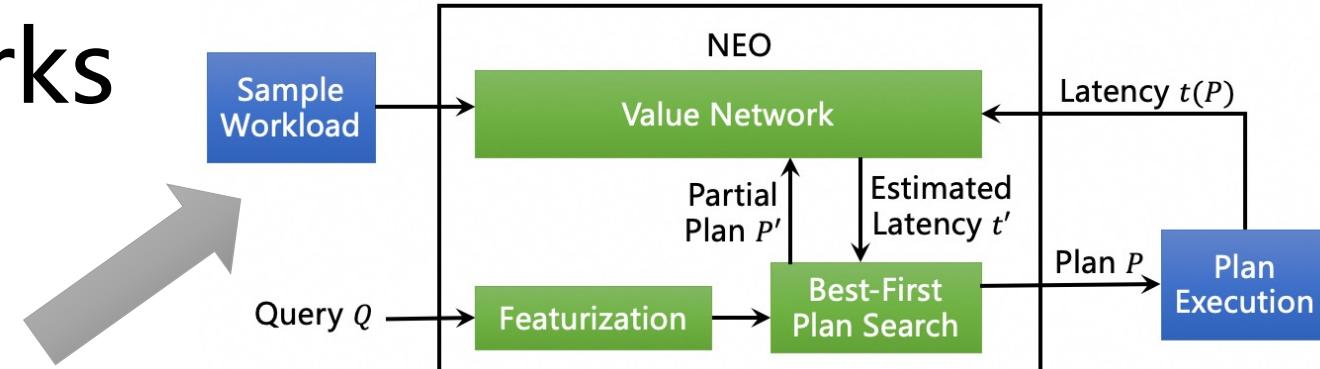


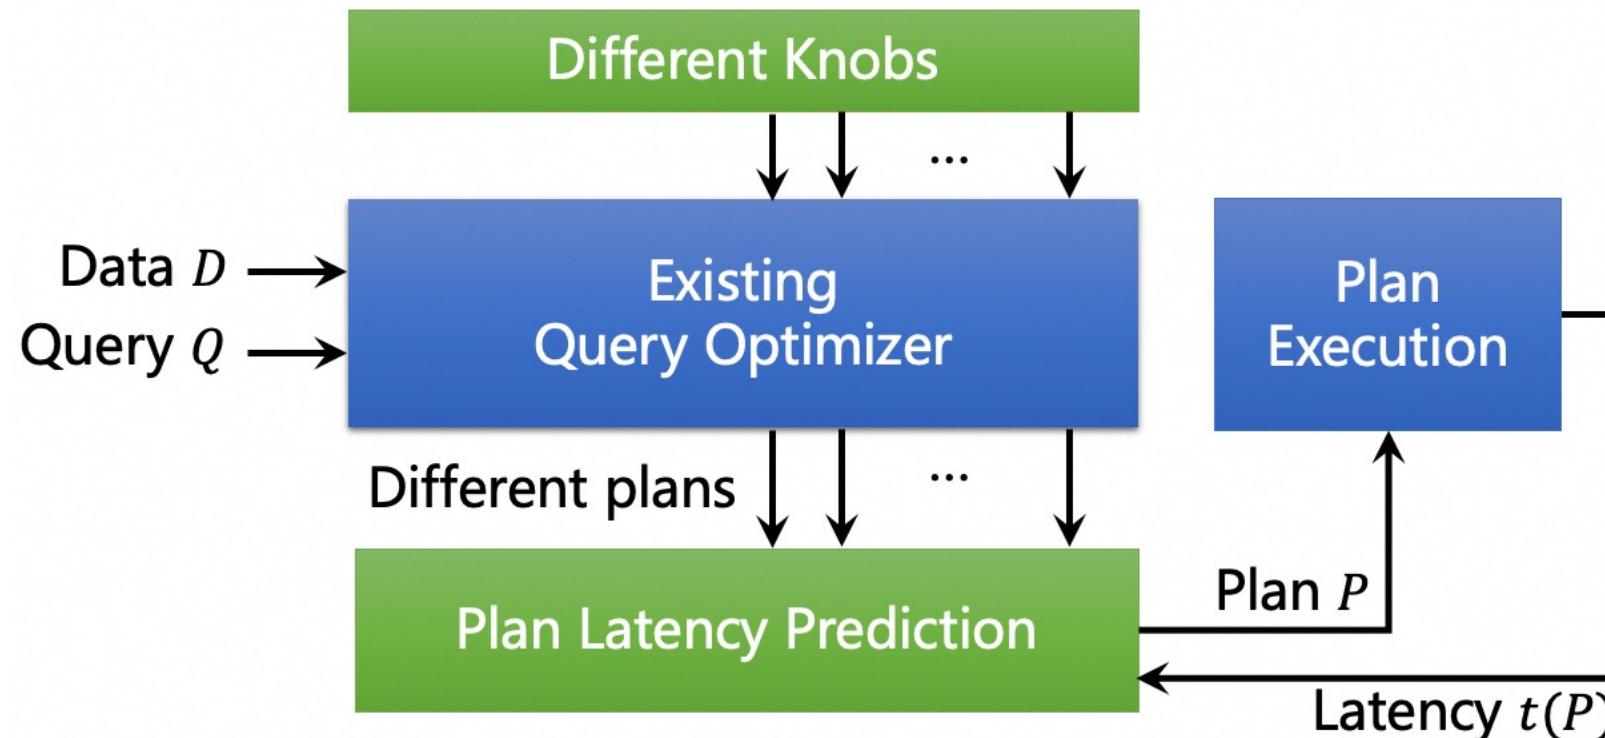
Figure from [Balsa: Learning a Query Optimizer Without Expert Demonstrations, SIGMOD 2022]

# Part 3: Learn to Explore Plan Space

- Search by Value Networks
- Search by Steering Existing Query Optimizers
- Search and Rank Plans
- Quality Control of Learned Query Optimizers

# Search by Steering Existing Query Optimizer

- Tune the existing query optimizer to search different paths in the plan space
  - Reduce the search space
  - Easy for deployment



# Different Tuning Knobs

- Bao: tune the hint set
  - Enable/Disable certain types of operations
    - Disable the nested loop join
- HyperQO: tune the leading hints on join order



Figure from [Bao: Making Learned Query Optimization Practical, SIGMOD 2021]

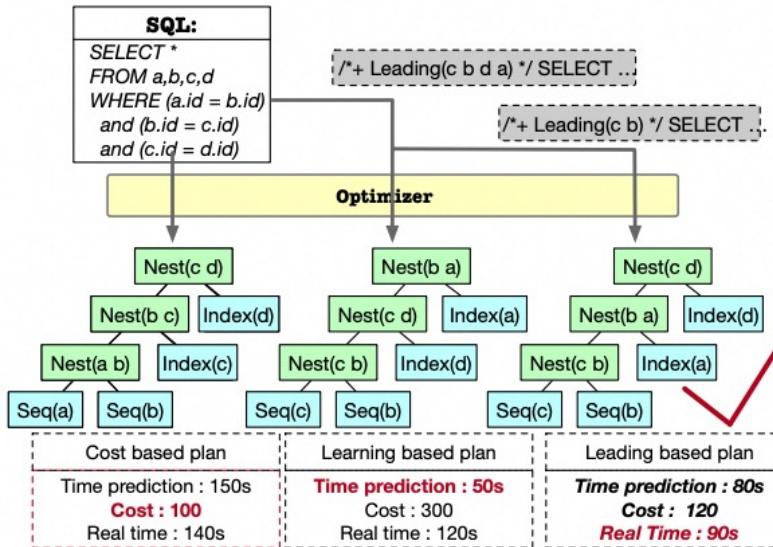
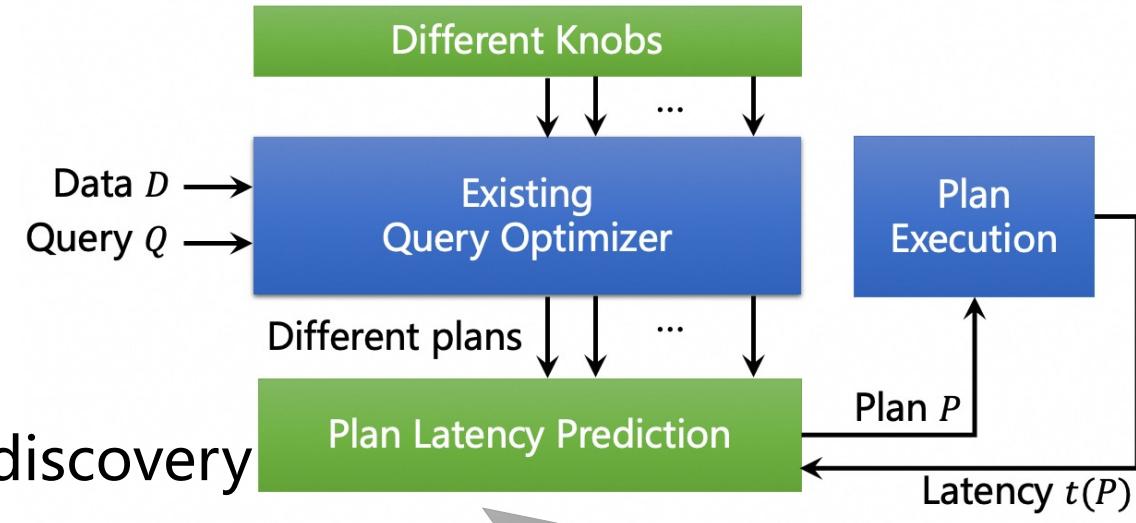


Figure from [Cost-based or Learning-based? A Hybrid Query Optimizer for Query Plan Selection, VLDB 2022]

# Disadvantages

- Large space of knobs
  - $h$  hints  $\rightarrow 2^h$  possible hint sets
  - AutoSteer: Automated suitable hint-set discovery
- Latency prediction is a very difficult problem
  - Large learning space: data, query, system and other factors ...
  - Select in-optimal plans
  - Slow model updating
  - ...

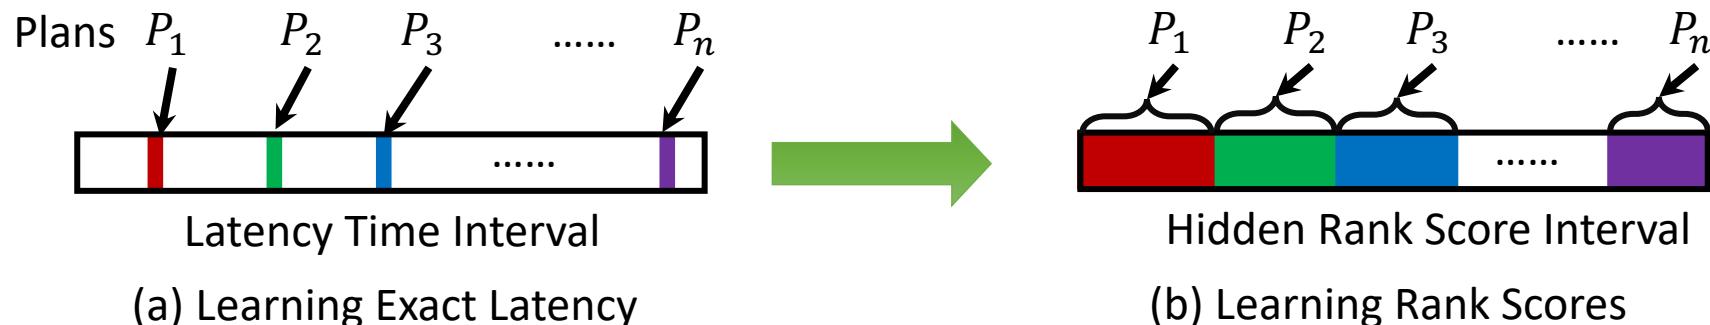


# Part 3: Learn to Explore Plan Space

- Search by Value Networks
- Search by Steering Existing Query Optimizers
- Search and Rank Plans
- Performance Quality Control

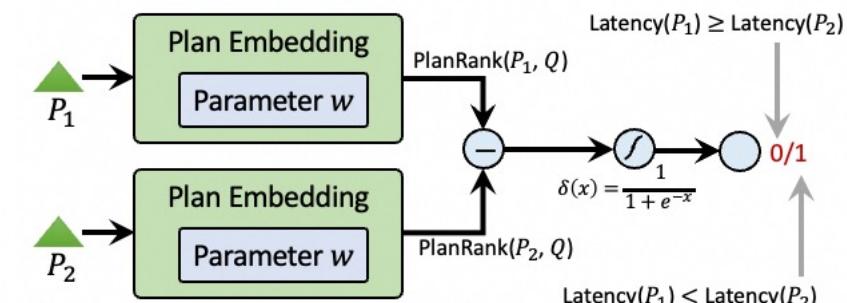
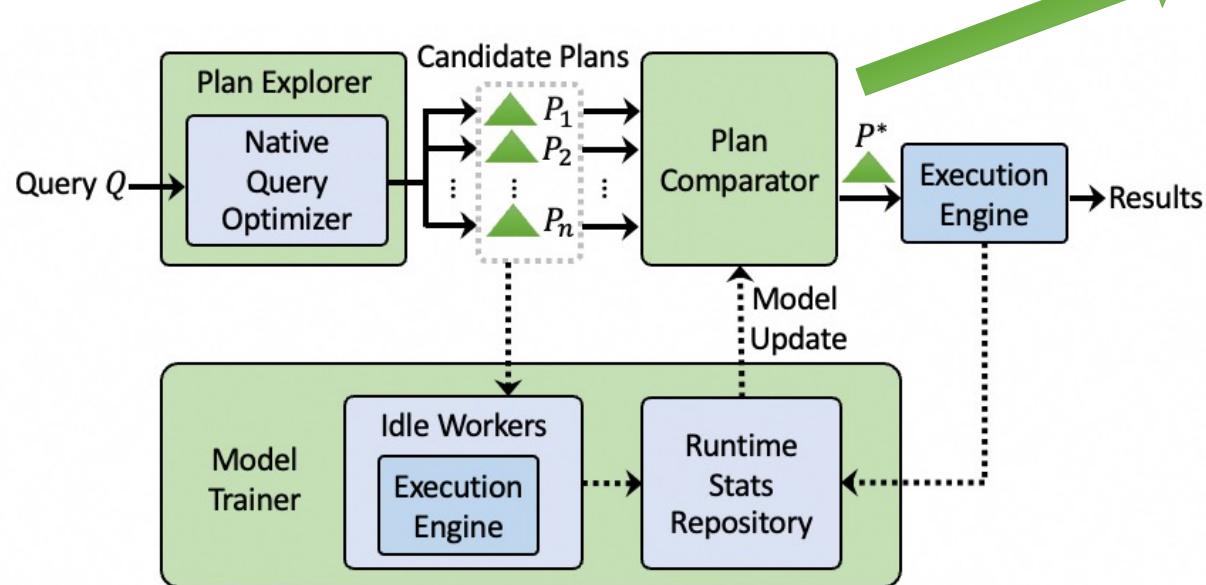
# Query Optimizer: A Better Solution

- Is it really necessary to know the exact time of each plan?
- No, we just need to know which plan is better!
- From learning-to-predict to learning-to-rank
  - Low learning difficulty with more benefits
  - Accurate model, fast learning and updating...



# Query Optimizer: Learn to Rank Plans

- Lero: learning-to-rank based query optimizer
  - Pairwise learning-to-rank model



Plan comparator  $\text{CmpPlan}(P_1, P_2)$

**Commutativity**

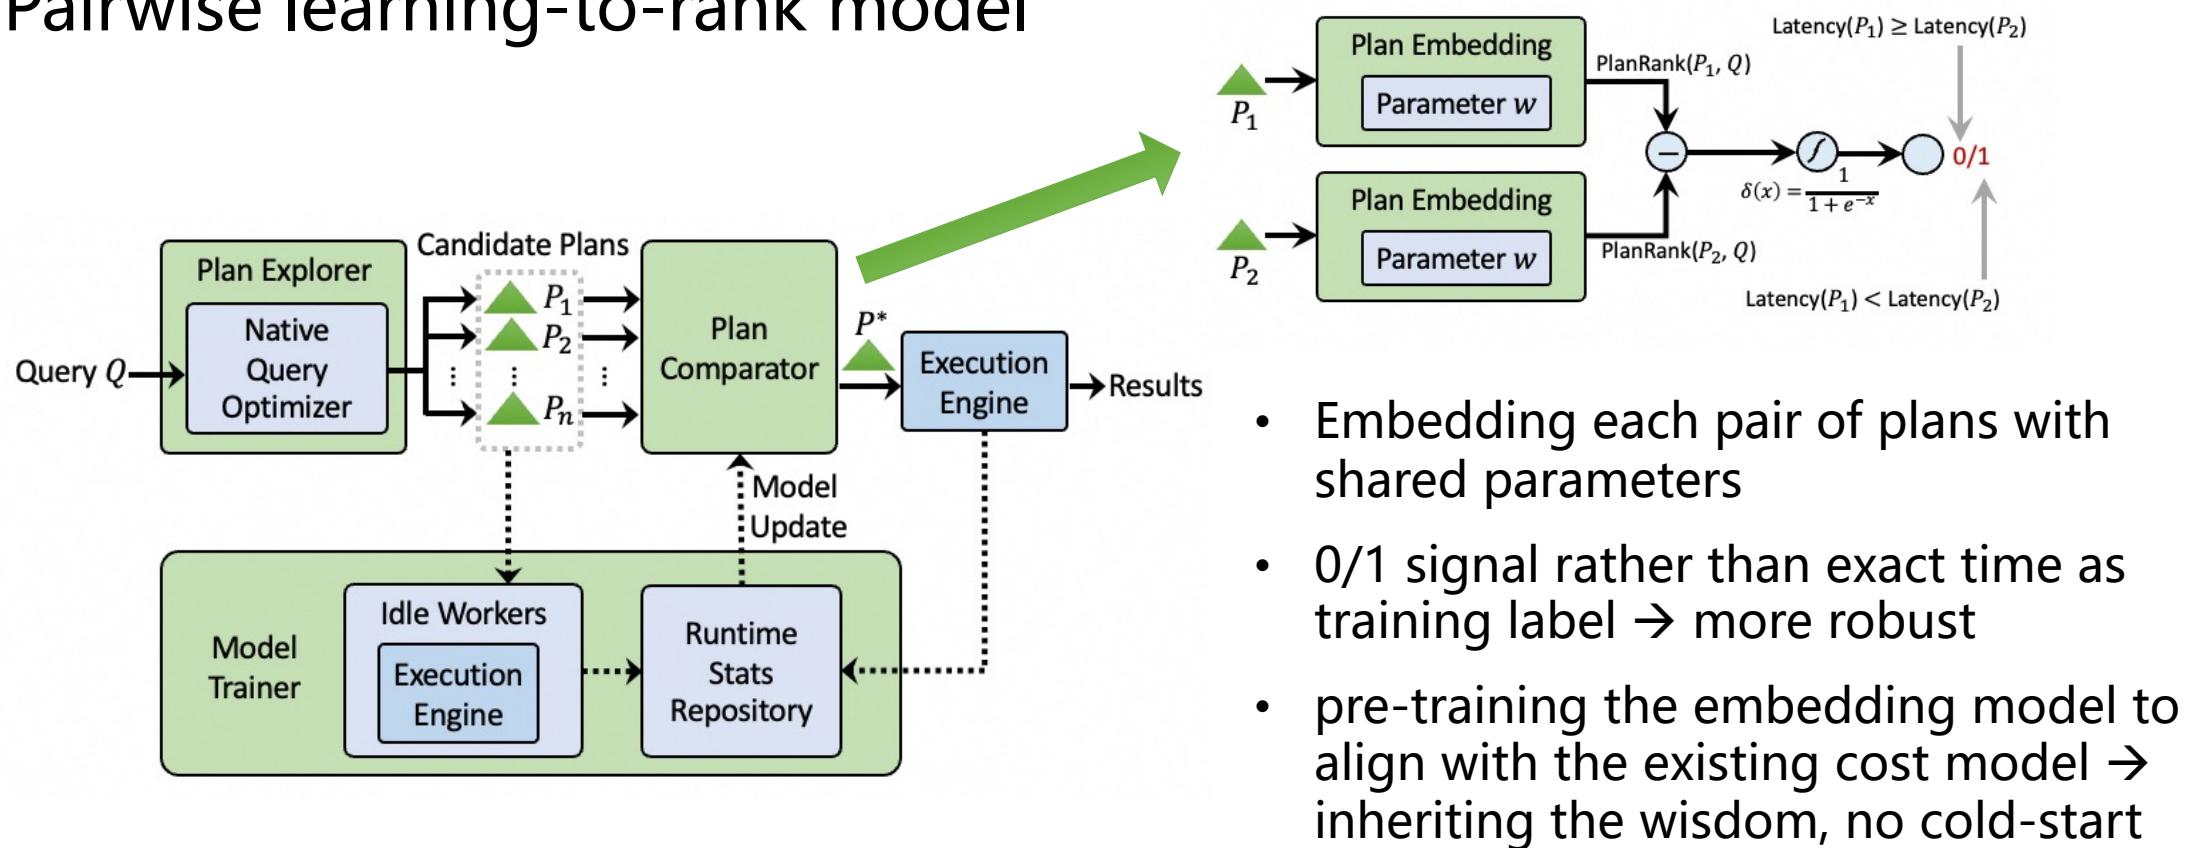
$$\bullet \quad \text{CmpPlan}(P_1, P_2) = 1 - \text{CmpPlan}(P_2, P_1)$$

**Transitivity**

$$\bullet \quad P_1 > P_2, P_2 > P_3 \Rightarrow P_1 > P_3$$

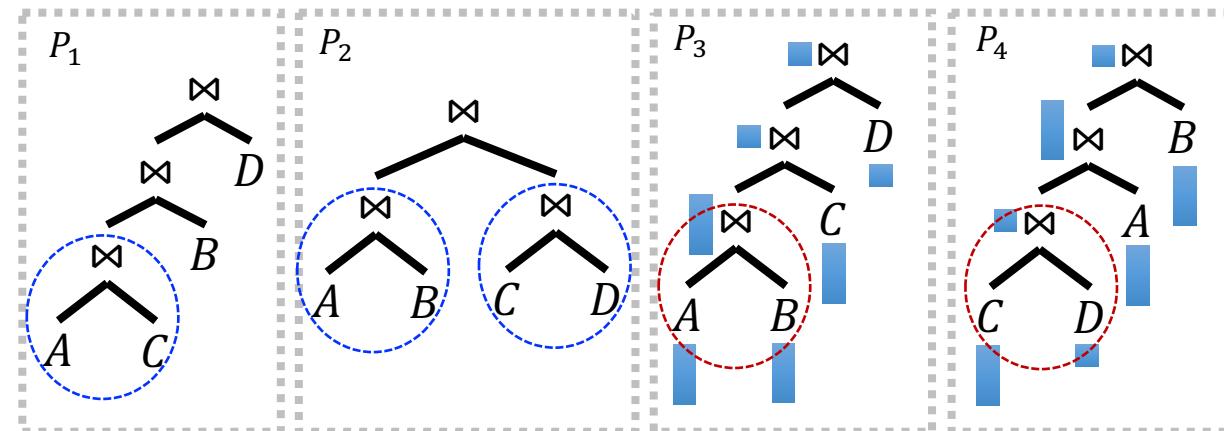
# Query Optimizer: Learn to Rank Plans

- Lero: learning-to-rank based query optimizer
  - Pairwise learning-to-rank model



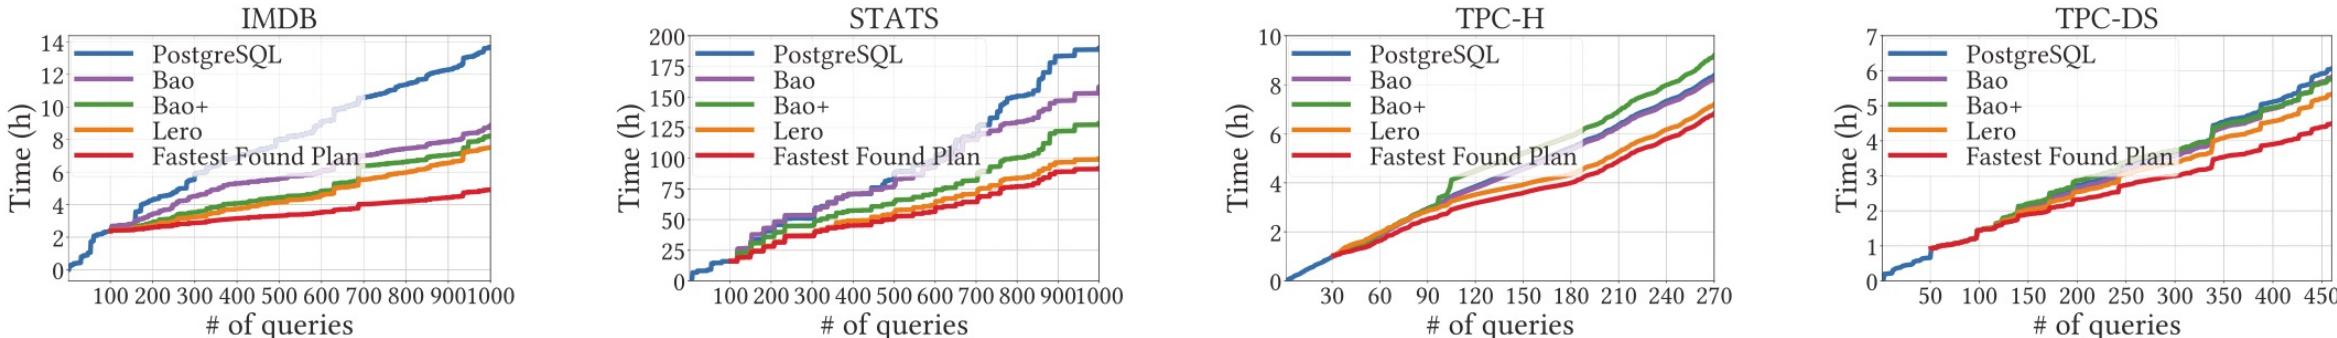
# Query Optimizer: Learn to Rank Plans

- Plan exploration in Lero: tune cardinality of all sub-queries
  - Group sub-queries together
  - Scale up/down their cardinality → could find better one with theoretical analysis
- Plan diversity
  - Different plan shapes
  - Different join orders

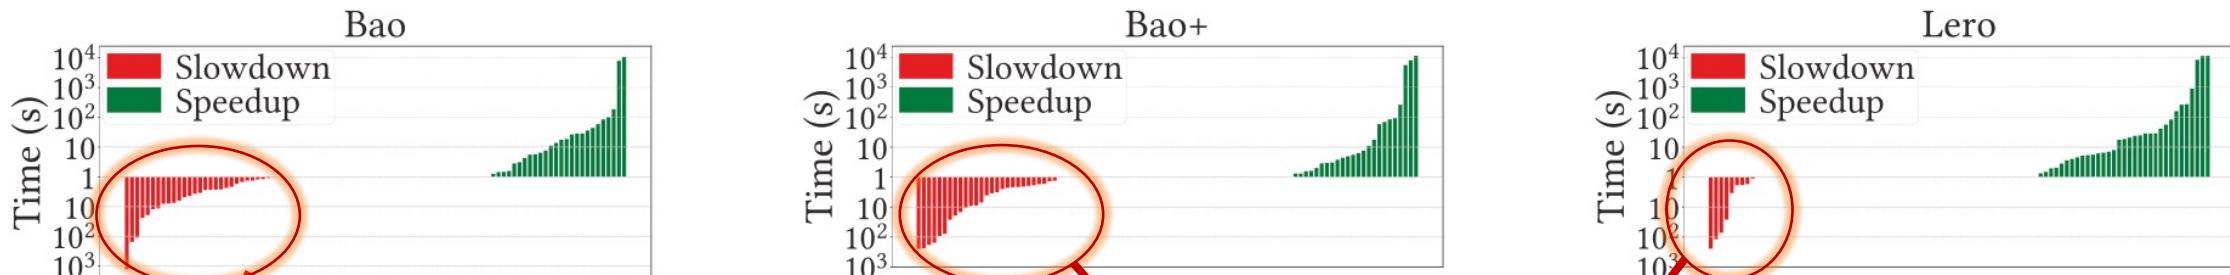


# Benchmark Evaluation

- Lero performs better than others
  - Static and dynamic settings



**Figure 6: Performance curve of different query optimizers since deployment on benchmarks.**



**Figure 5: Per-query execution time of different query optimizers in comparison with PostgreSQL on the STATS benchmark.**  
Could we ultimately eliminate such performance regressions without affecting the benefits?

**Table 1: Overall performance of different query optimizers.**

Query Optimizer	Execution Time (in hour)			
	STATS	IMDB	TPC-H	TPC-DS
PostgreSQL	20.19	1.15	0.94	1.68
Bao	15.32	0.47	1.17	1.55
Bao+	13.85	0.41	0.89	1.57
Lero	<b>11.32</b>	<b>0.35</b>	<b>0.74</b>	<b>1.47</b>
<b>Fastest Found Plan</b>	<b>10.73</b>	<b>0.19</b>	<b>0.72</b>	<b>1.39</b>

# Part 3: Learn to Explore Plan Space

- Search by Value Networks
- Search by Steering Existing Query Optimizers
- Search and Rank Plans
- Quality Control of Learned Query Optimizers

# Learned Query Optimizer: Quality Control

- HyperQO: an ensemble model to learn multiple prediction results of each plan
  - Plans having large variance are filtered before selection
  - Small variance ≠ accurate estimation
    - Bad plans may be reserved if all models make the same mistake

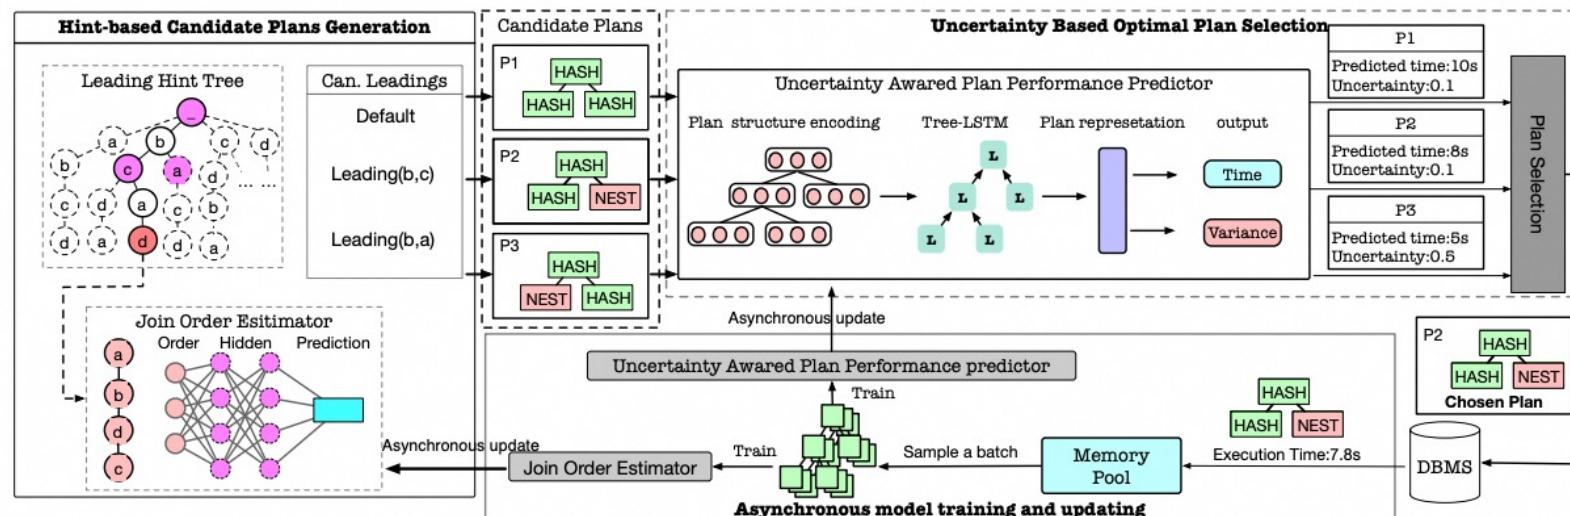
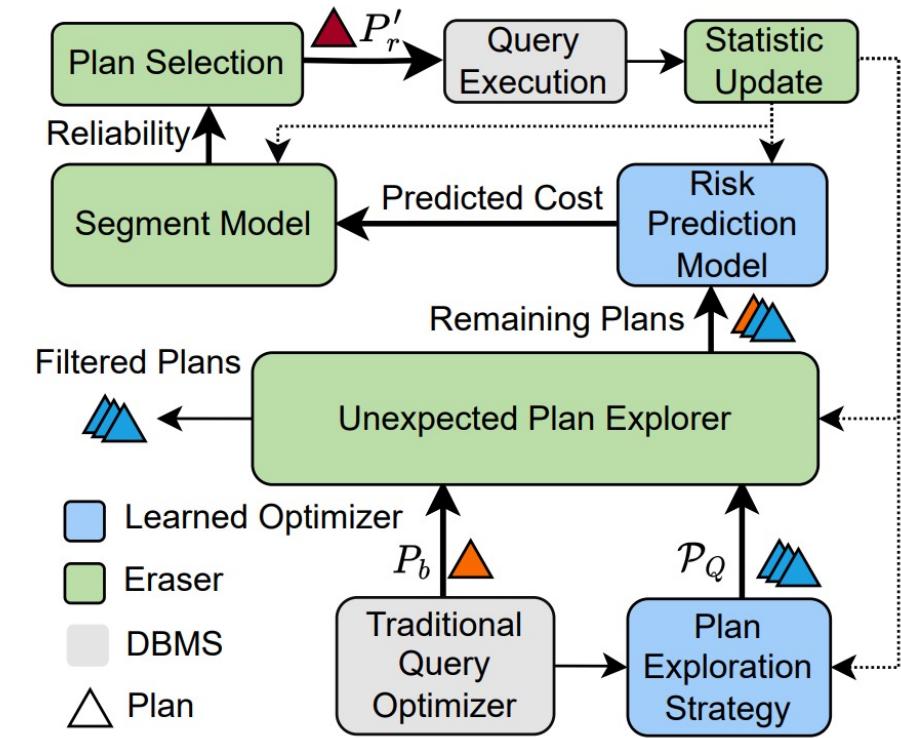


Figure from [Cost-based or Learning-based? A Hybrid Query Optimizer for Query Plan Selection, VLDB 2022]

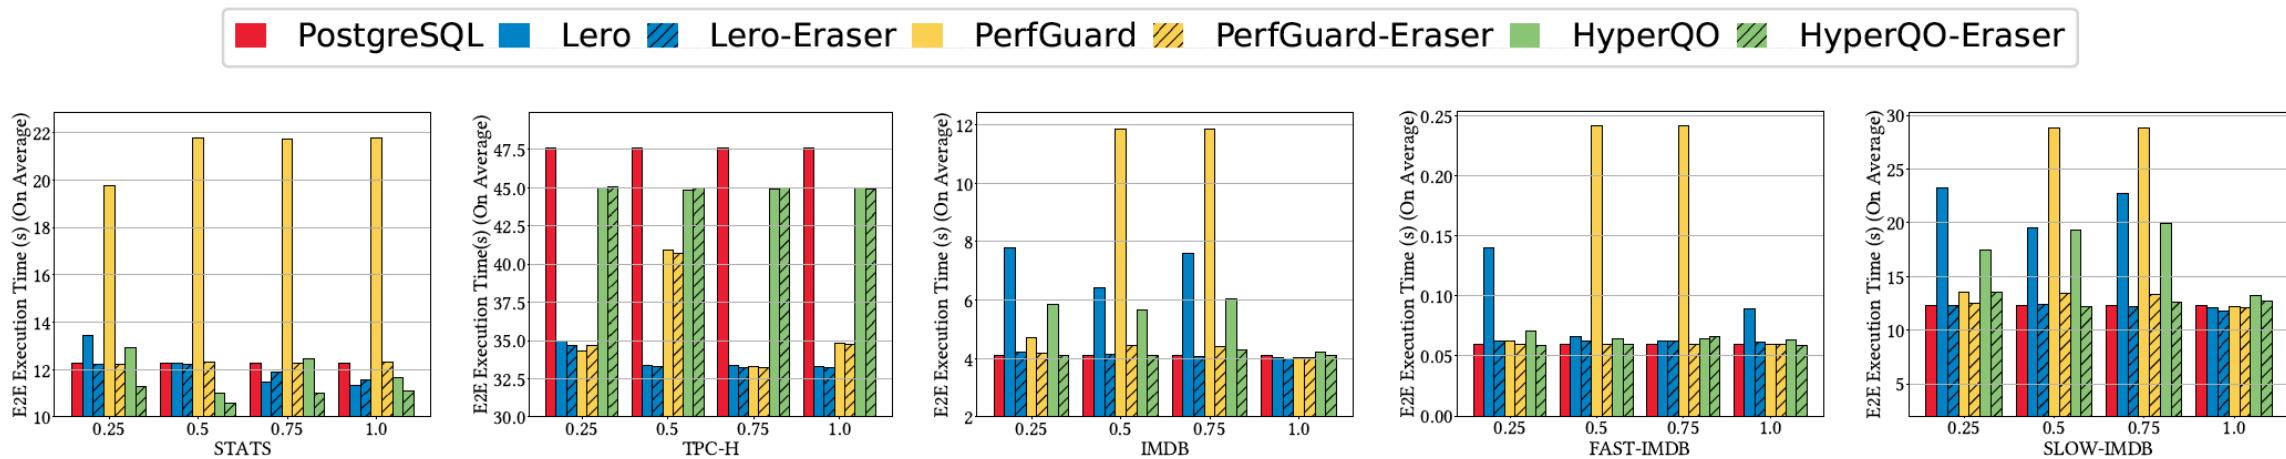
# Learned Query Optimizer: Quality Control

- Eraser: eliminate regressions while preserving benefits
- Unexpected plan explorer
  - A coarse-grained filter removing all highly risky plans with unseen feature values
- Segment model
  - A fine-grained model to evaluate the quality of each plan for plan selection



# Learned Query Optimizer: Quality Control

- Evaluation results of Eraser
  - Learned query optimizer performs worse than traditional one
    - Eraser could eliminate the performance regressions
  - Learned query optimizer performs better than traditional one
    - Eraser has very little impact on the performance improvements



# Part 4: Applications and Deployment

- Prototype Applications
  - Learned cost model for SCOPE in Microsoft
  - Learn to steer SCOPE in Microsoft
  - OpenGauss
- PilotScope: A Middleware System for Deployment

# Part 4: Applications and Deployment

- Prototype Applications
  - Learned cost model for SCOPE in Microsoft
  - Learn to steer SCOPE in Microsoft
  - OpenGauss
- PilotScope: A Middleware System for Deployment

# Learned Cost Model in SCOPE

- SCOPE: the big data system for data analytics in Microsoft
- Learned cost models could improve plan quality of production jobs

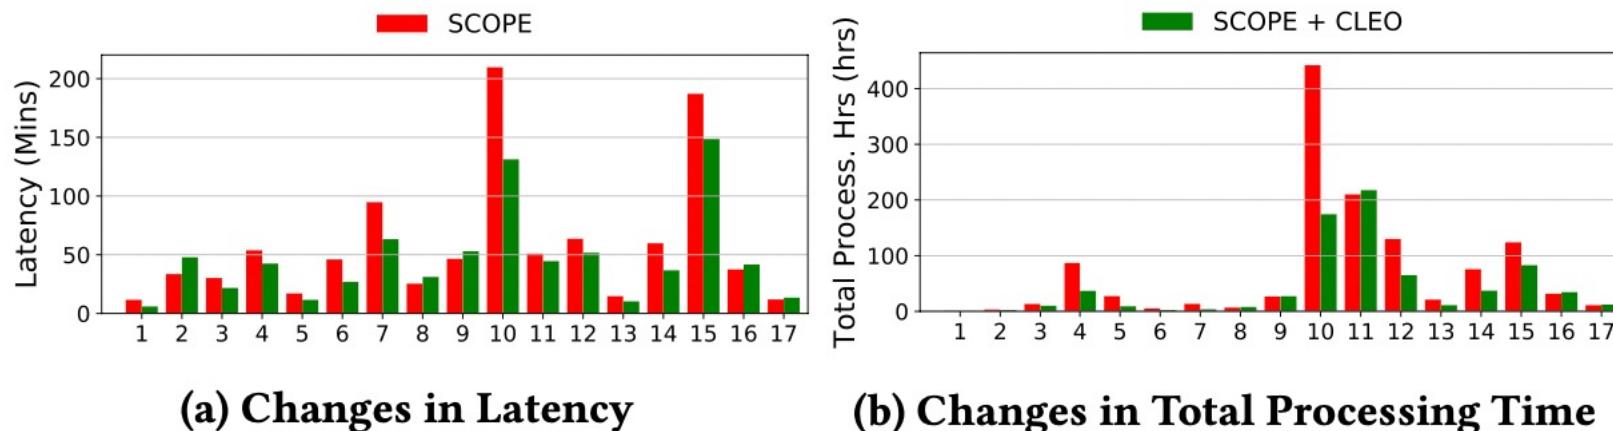


Figure from [Cost Models for Big Data Query Processing: Learning, Retrofitting, and Our Findings, SIGMOD 2020]

# SCOPE and Its Hints

- SCOPE: tuning different hints
  - 256 hints in total, 100-150 hints are used frequently
  - Typically 10 – 20 hints are used in a single job
  - Analyze to discover suitable hint set to steer SCOPE

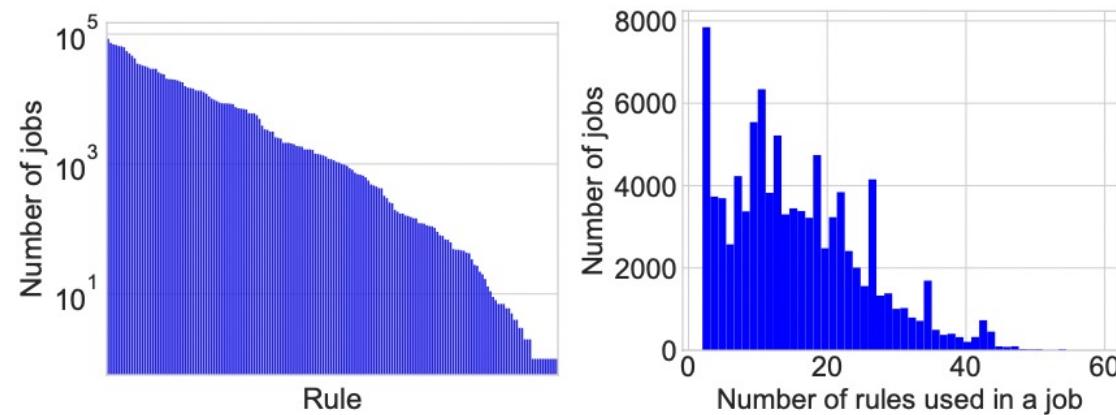


Figure from [Steering Query Optimizers: A Practical Take on Big Data Workloads, SIGMOD 2021]

# Learn to Steer SCOPE

- Online deployment: large improvements, small regressions

	1			2			3		
	Mean	90P	99P	Mean	90P	99P	Mean	90P	99P
Best	5458	14K	14.8K	19.8K	26K	27K	2966	13.8K	15.3K
Default	6461	16.3K	18.3K	20.7K	26.9K	28.9K	3304	14.7K	16.8K
Learned	<b>5724</b>	<b>14.7K</b>	<b>15.4K</b>	<b>20.2K</b>	<b>26.2K</b>	<b>27K</b>	<b>3252</b>	<b>14.6K</b>	<b>16.8K</b>

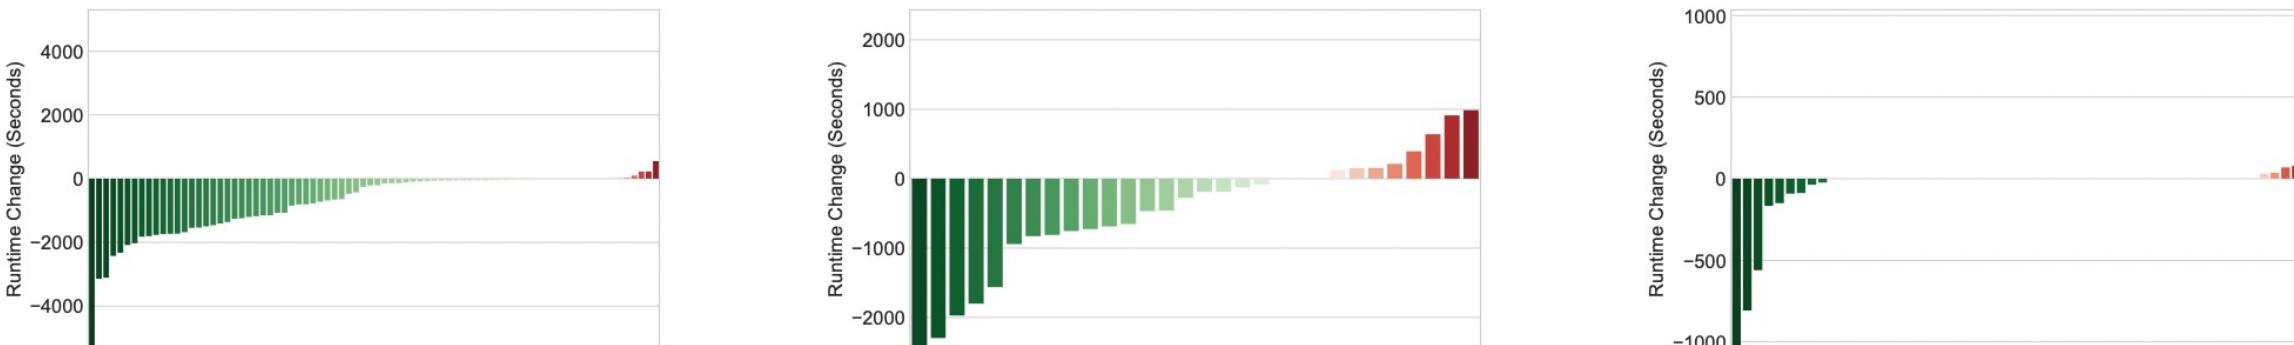


Figure and Table from [Steering Query Optimizers: A Practical Take on Big Data Workloads, SIGMOD 2021]

# OpenGauss: Automatic DB

- Multiple learned components are deployed into GaussDB
  - Query rewriter
  - Cardinality/Cost Estimator
  - Plan Enumerator
  - ...
- Traditional Methods as backups
- Difficult to extend to other AI4DB methods and DB systems

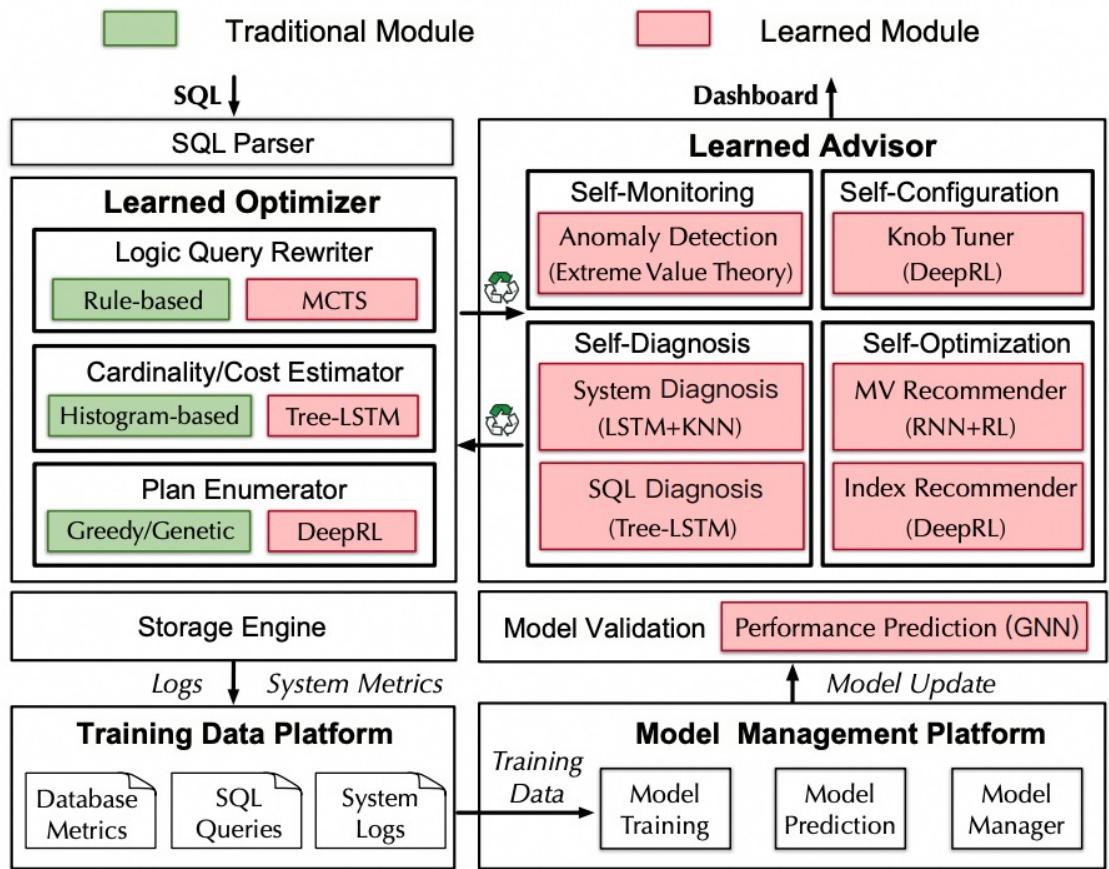


Figure from [openGauss: An Autonomous Database System, VLDB 2021]

# Part 4: Applications and Deployment

- Prototype Applications
  - Learned cost model for SCOPE in Microsoft
  - Learn to steer SCOPE
  - OpenGauss
- PilotScope: A Middleware System for Deployment

# PilotScope: Motivations

- AI4DB algorithms are developing very fast .....

- A fundamental problem is still not well studied

- Could they actually contribute to real world DBMS?

- If so, to what extent?
  - If not, why not?

- This could only be answered by actual deployment (**gold standard**)

- Examine in real systems, on real data, report real performance
  - The pros & cons of each method

.....

→ guide future research towards the right direction



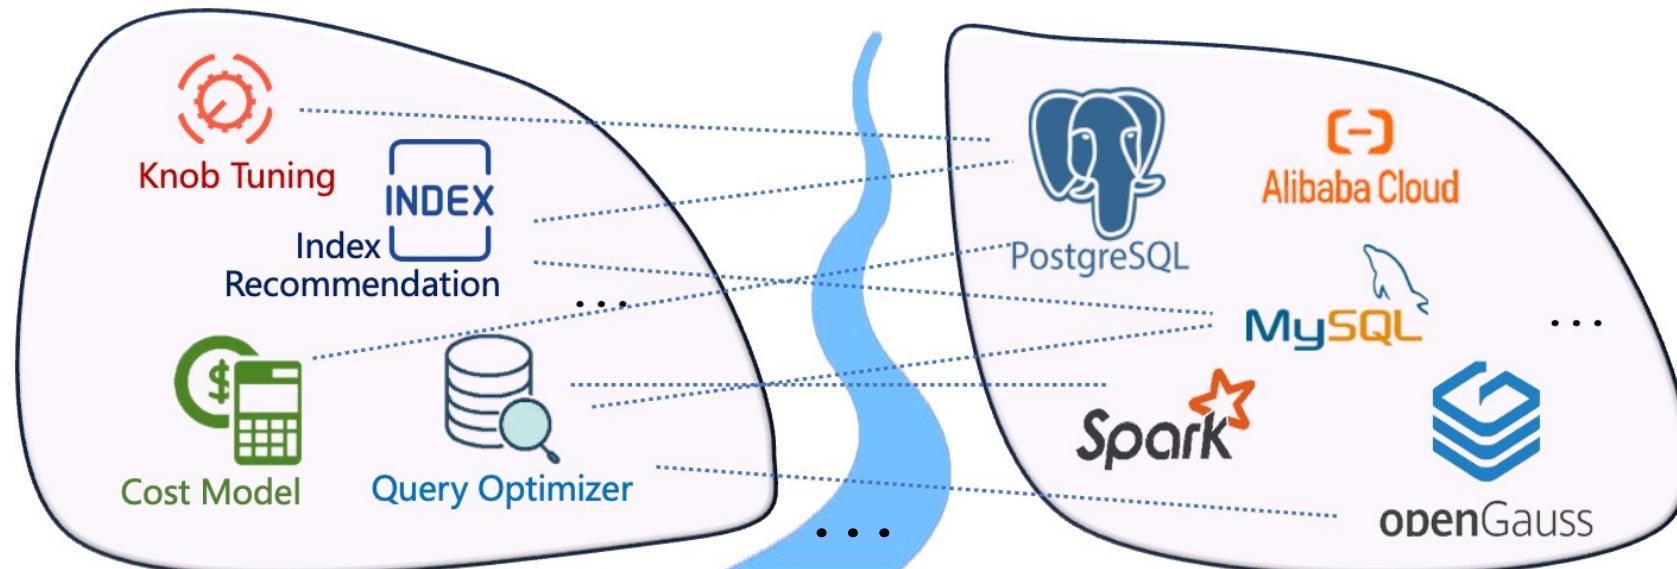
...



# PilotScope: Existing Solutions

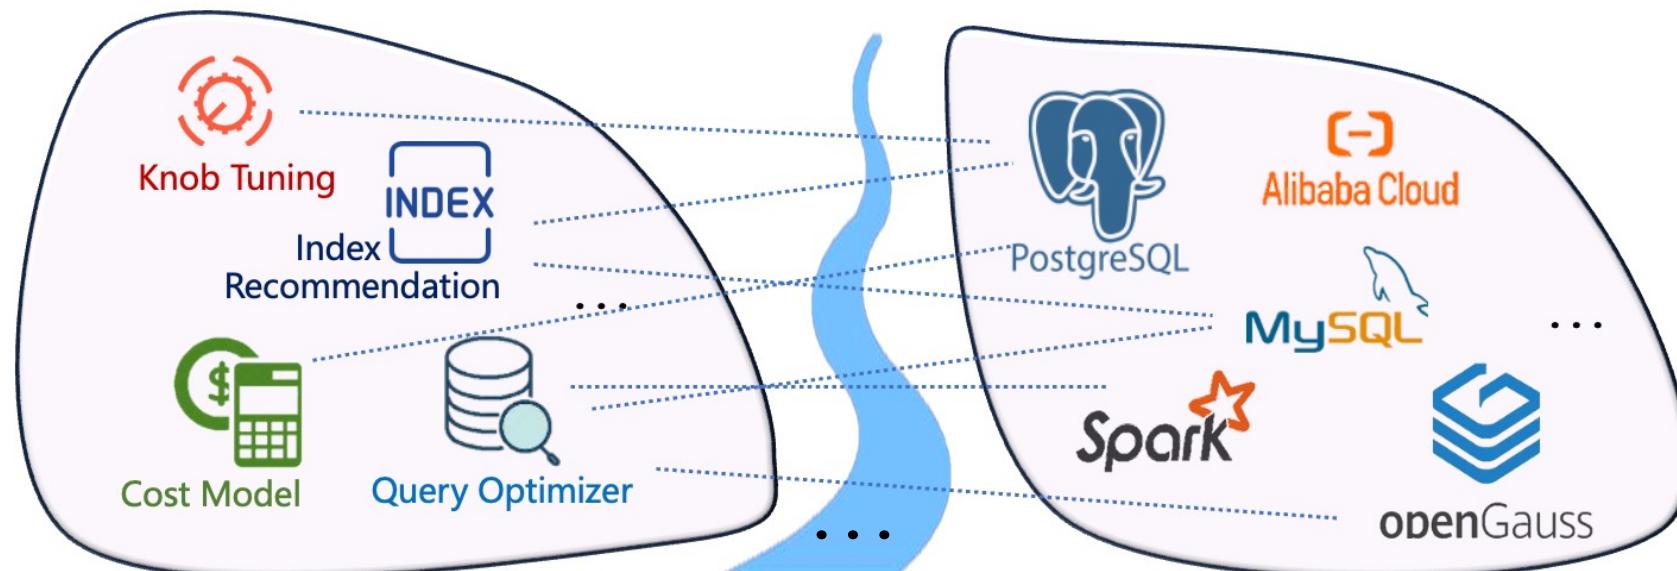
- Deploying specific algorithm into specific DB
  - Learned query optimizer/cost model → SCOPE
  - Knob tuning/query optimizer → openGauss
  - Workload forecasting/behavior modeling → NoisePage

.....



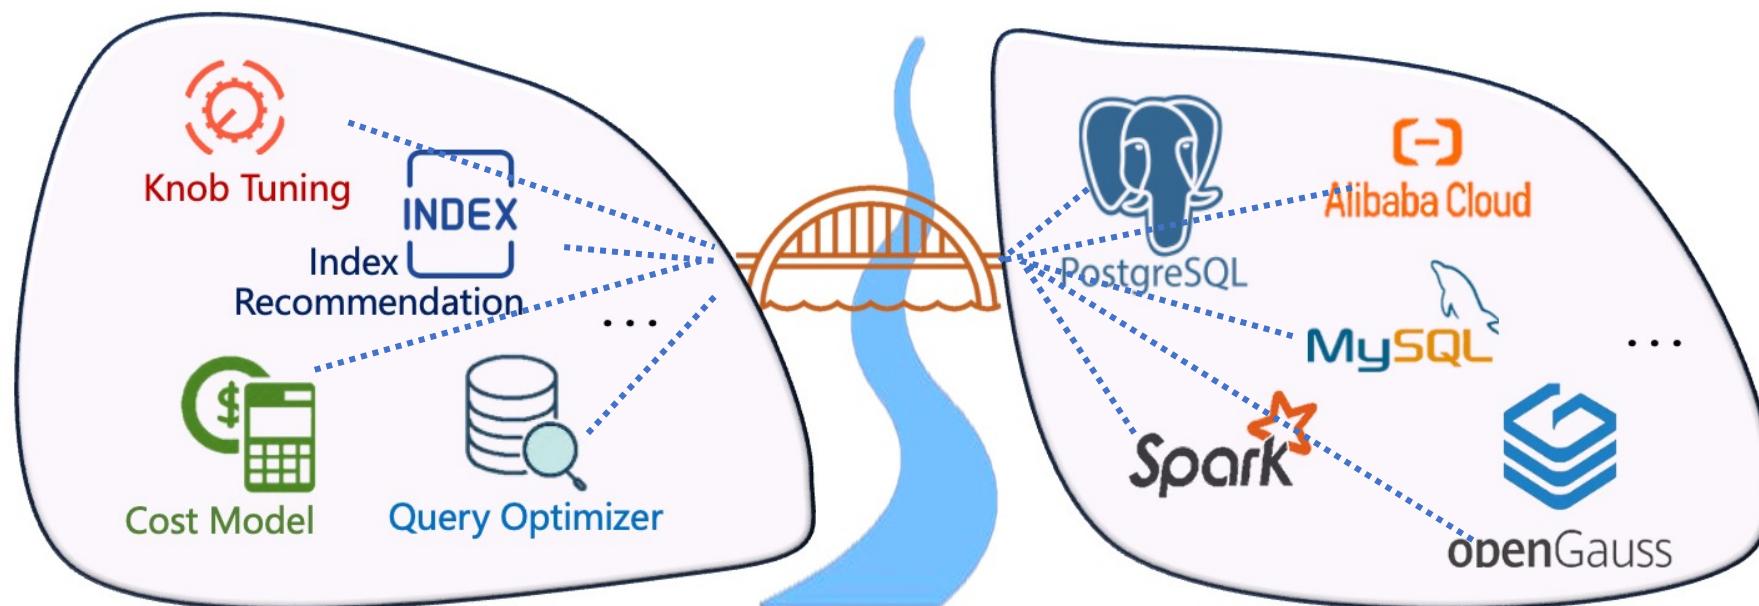
# PilotScope: Existing Solutions

- Limitations of existing solutions
  - Not general to other AI4DB algorithms & databases
  - Very expensive:  $O(m * n)$  tailored solutions for  $m$  algorithms &  $n$  databases
  - Not friendly to developers
    - Require close cooperation and learning between ML and DB developers



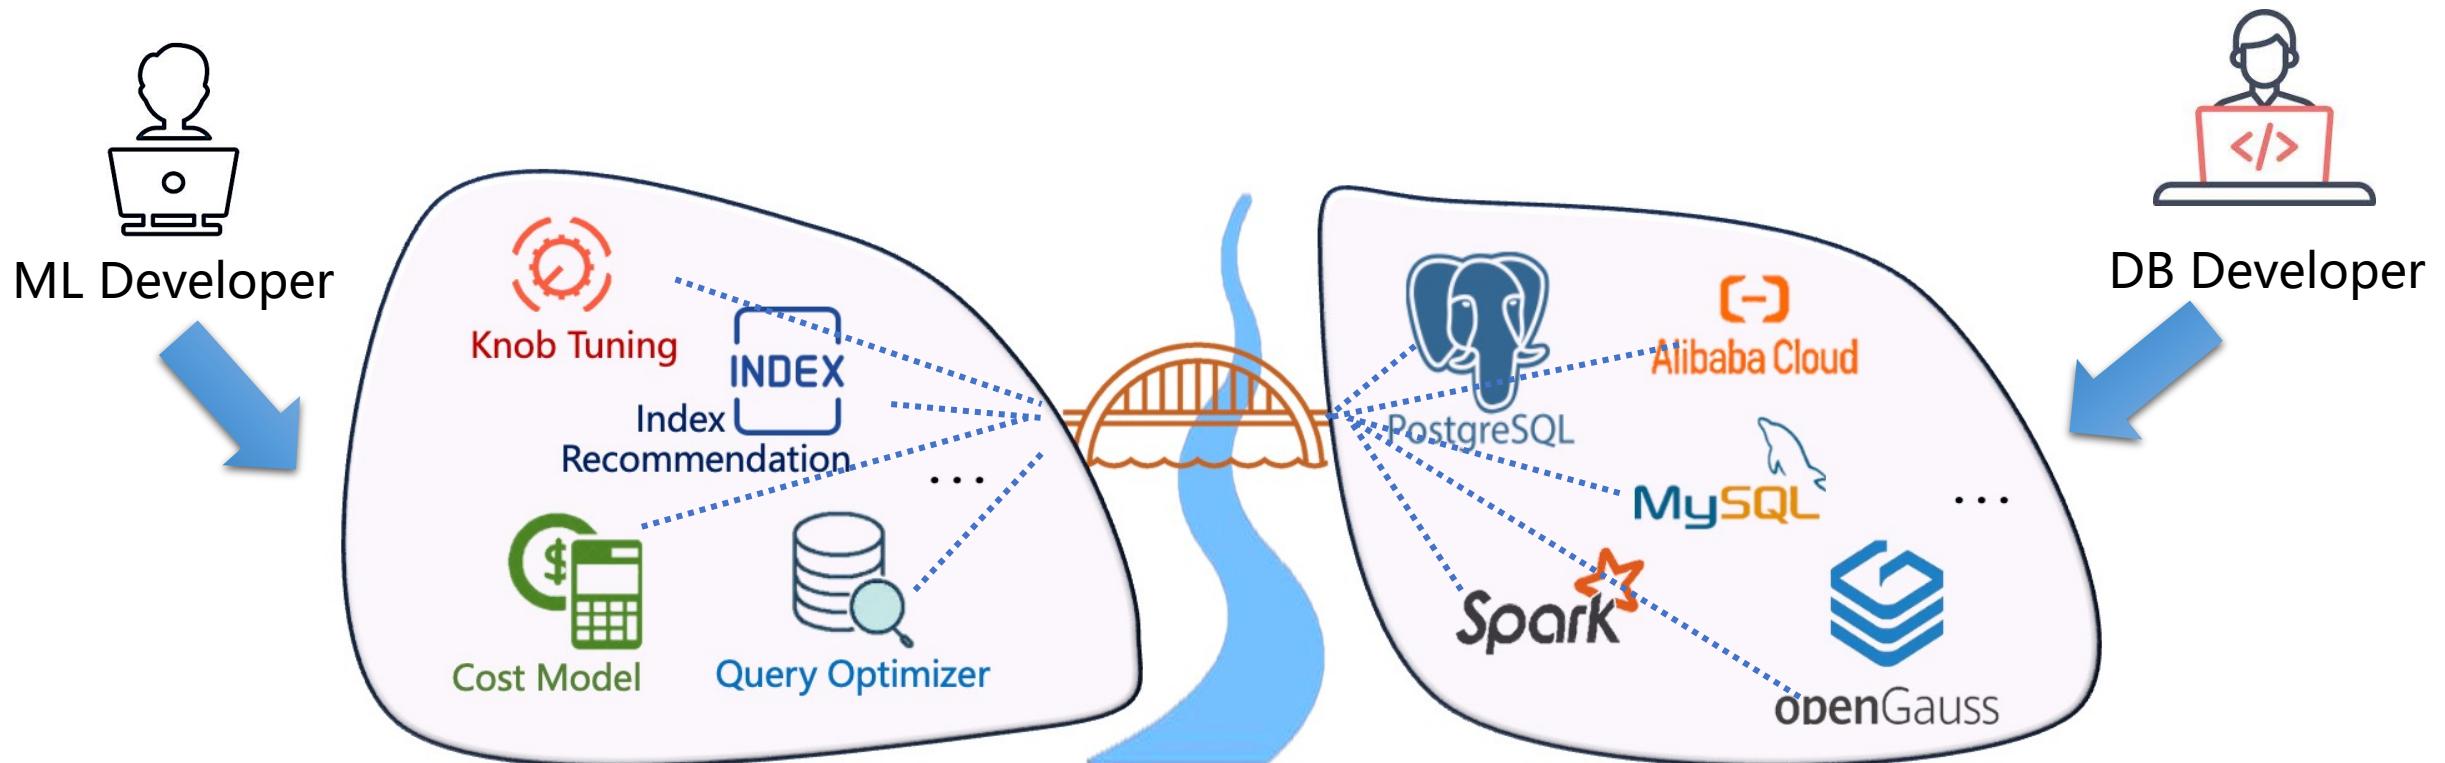
# PilotScope: Our Principle

- What we desire of the AI4DB deployment?
  - One system generally applicable for multiple algorithms & databases
    - General to algorithm: support different AI4DB tasks
    - General to system: the same ML program could run on different systems
    - Easy to extensible to other AI4DB algorithms/database systems



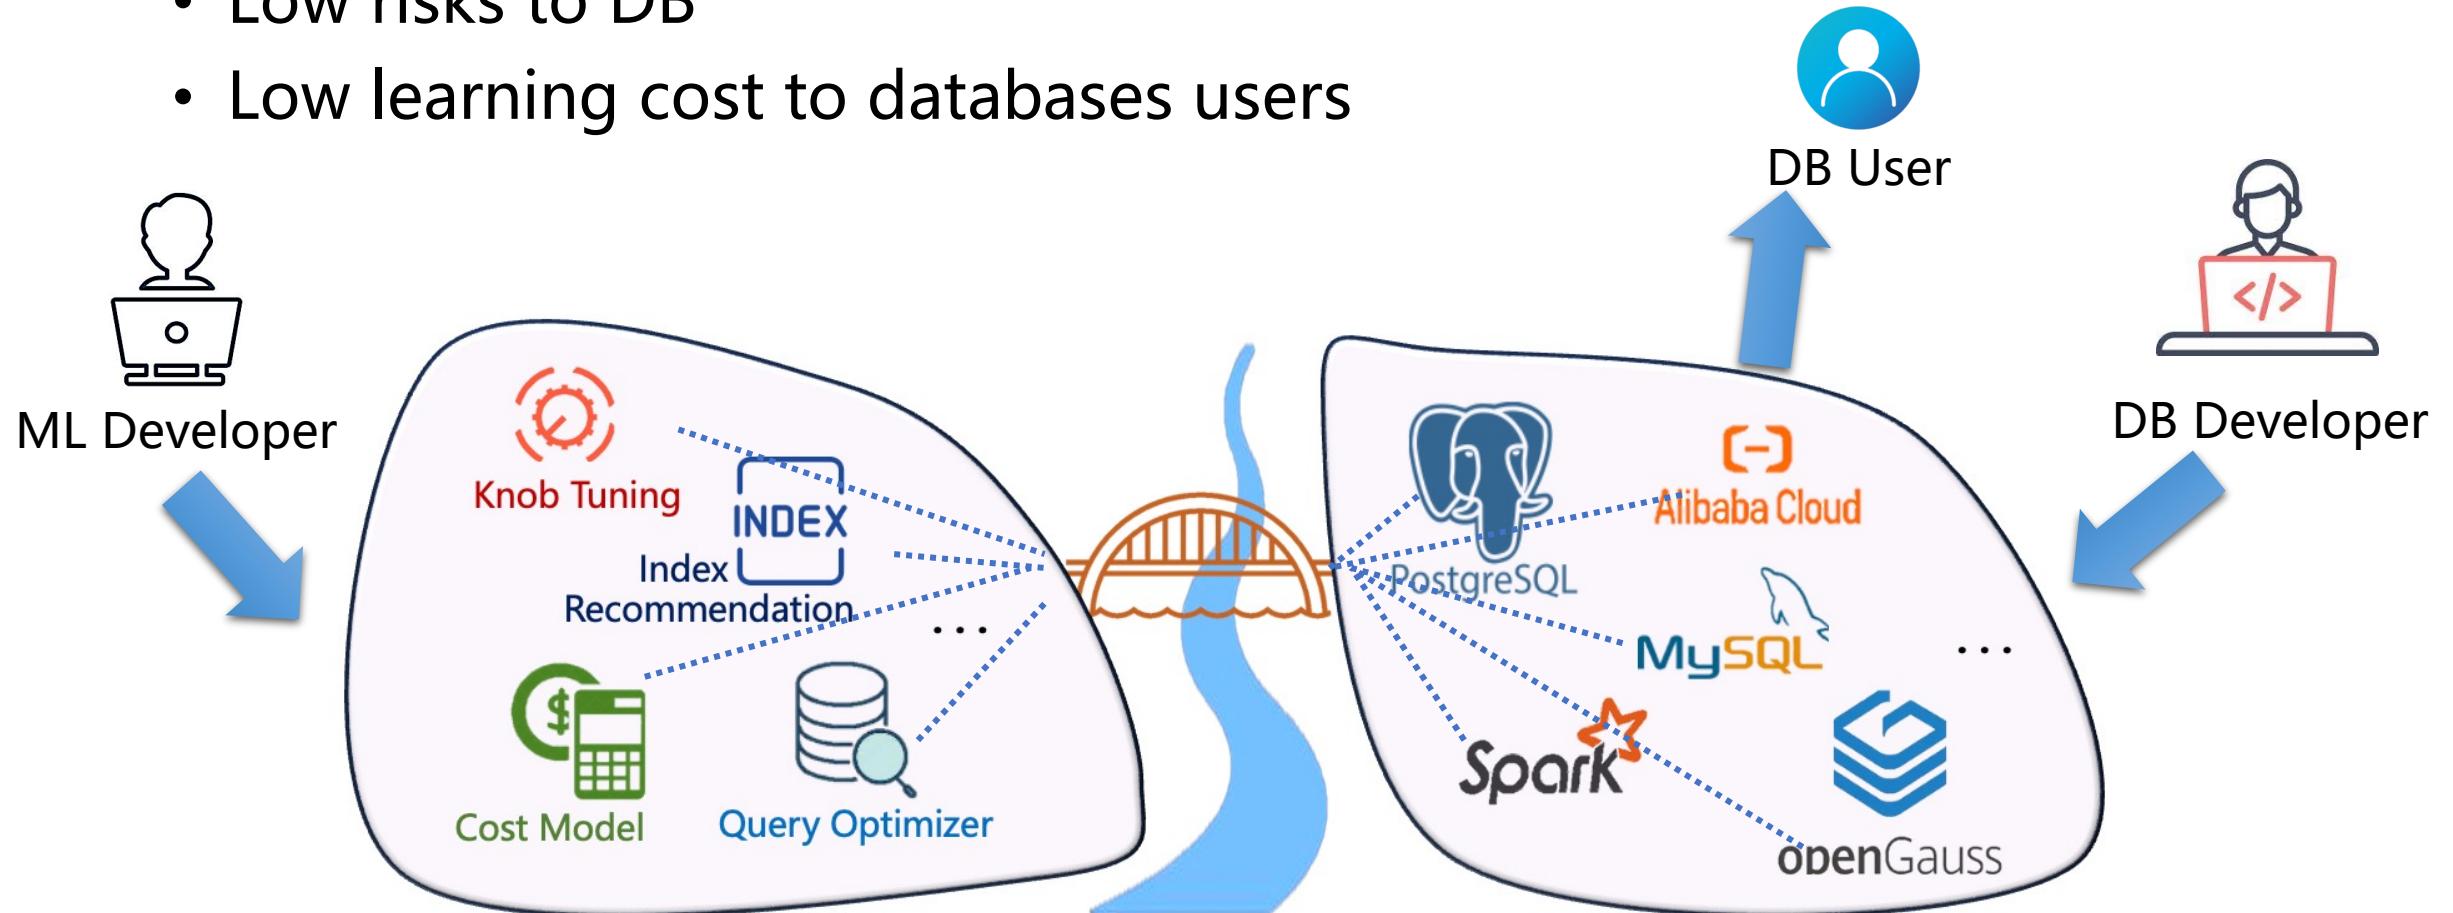
# PilotScope: Our Principle

- What we desire of the AI4DB deployment?
  - ML & DB developers could work independently
    - Developers could focus on writing their own programs
    - Not need to know details in other side



# PilotScope: Our Principle

- What we desire of the AI4DB deployment?
  - The minimal impact to the native databases
    - Low risks to DB
    - Low learning cost to databases users

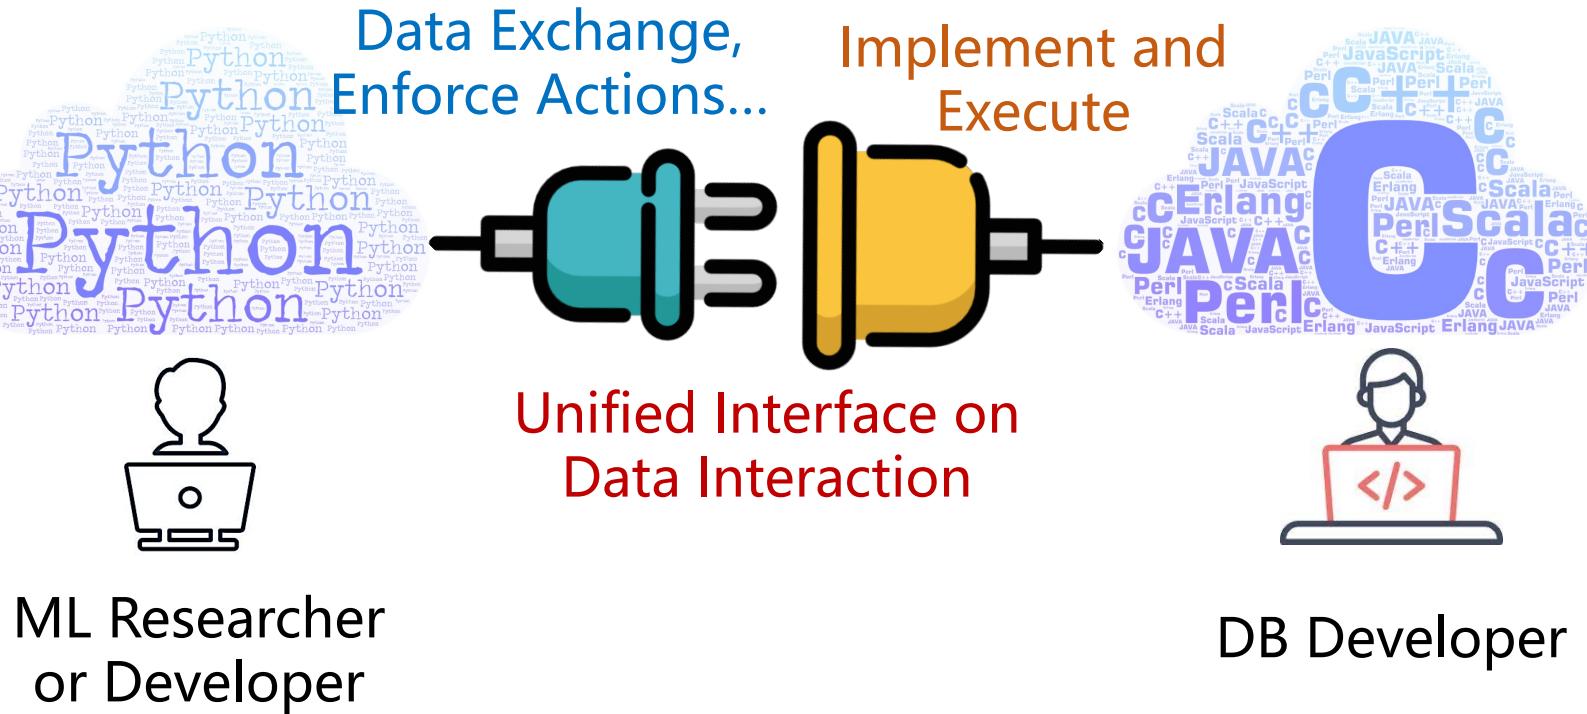


# PilotScope: Challenges

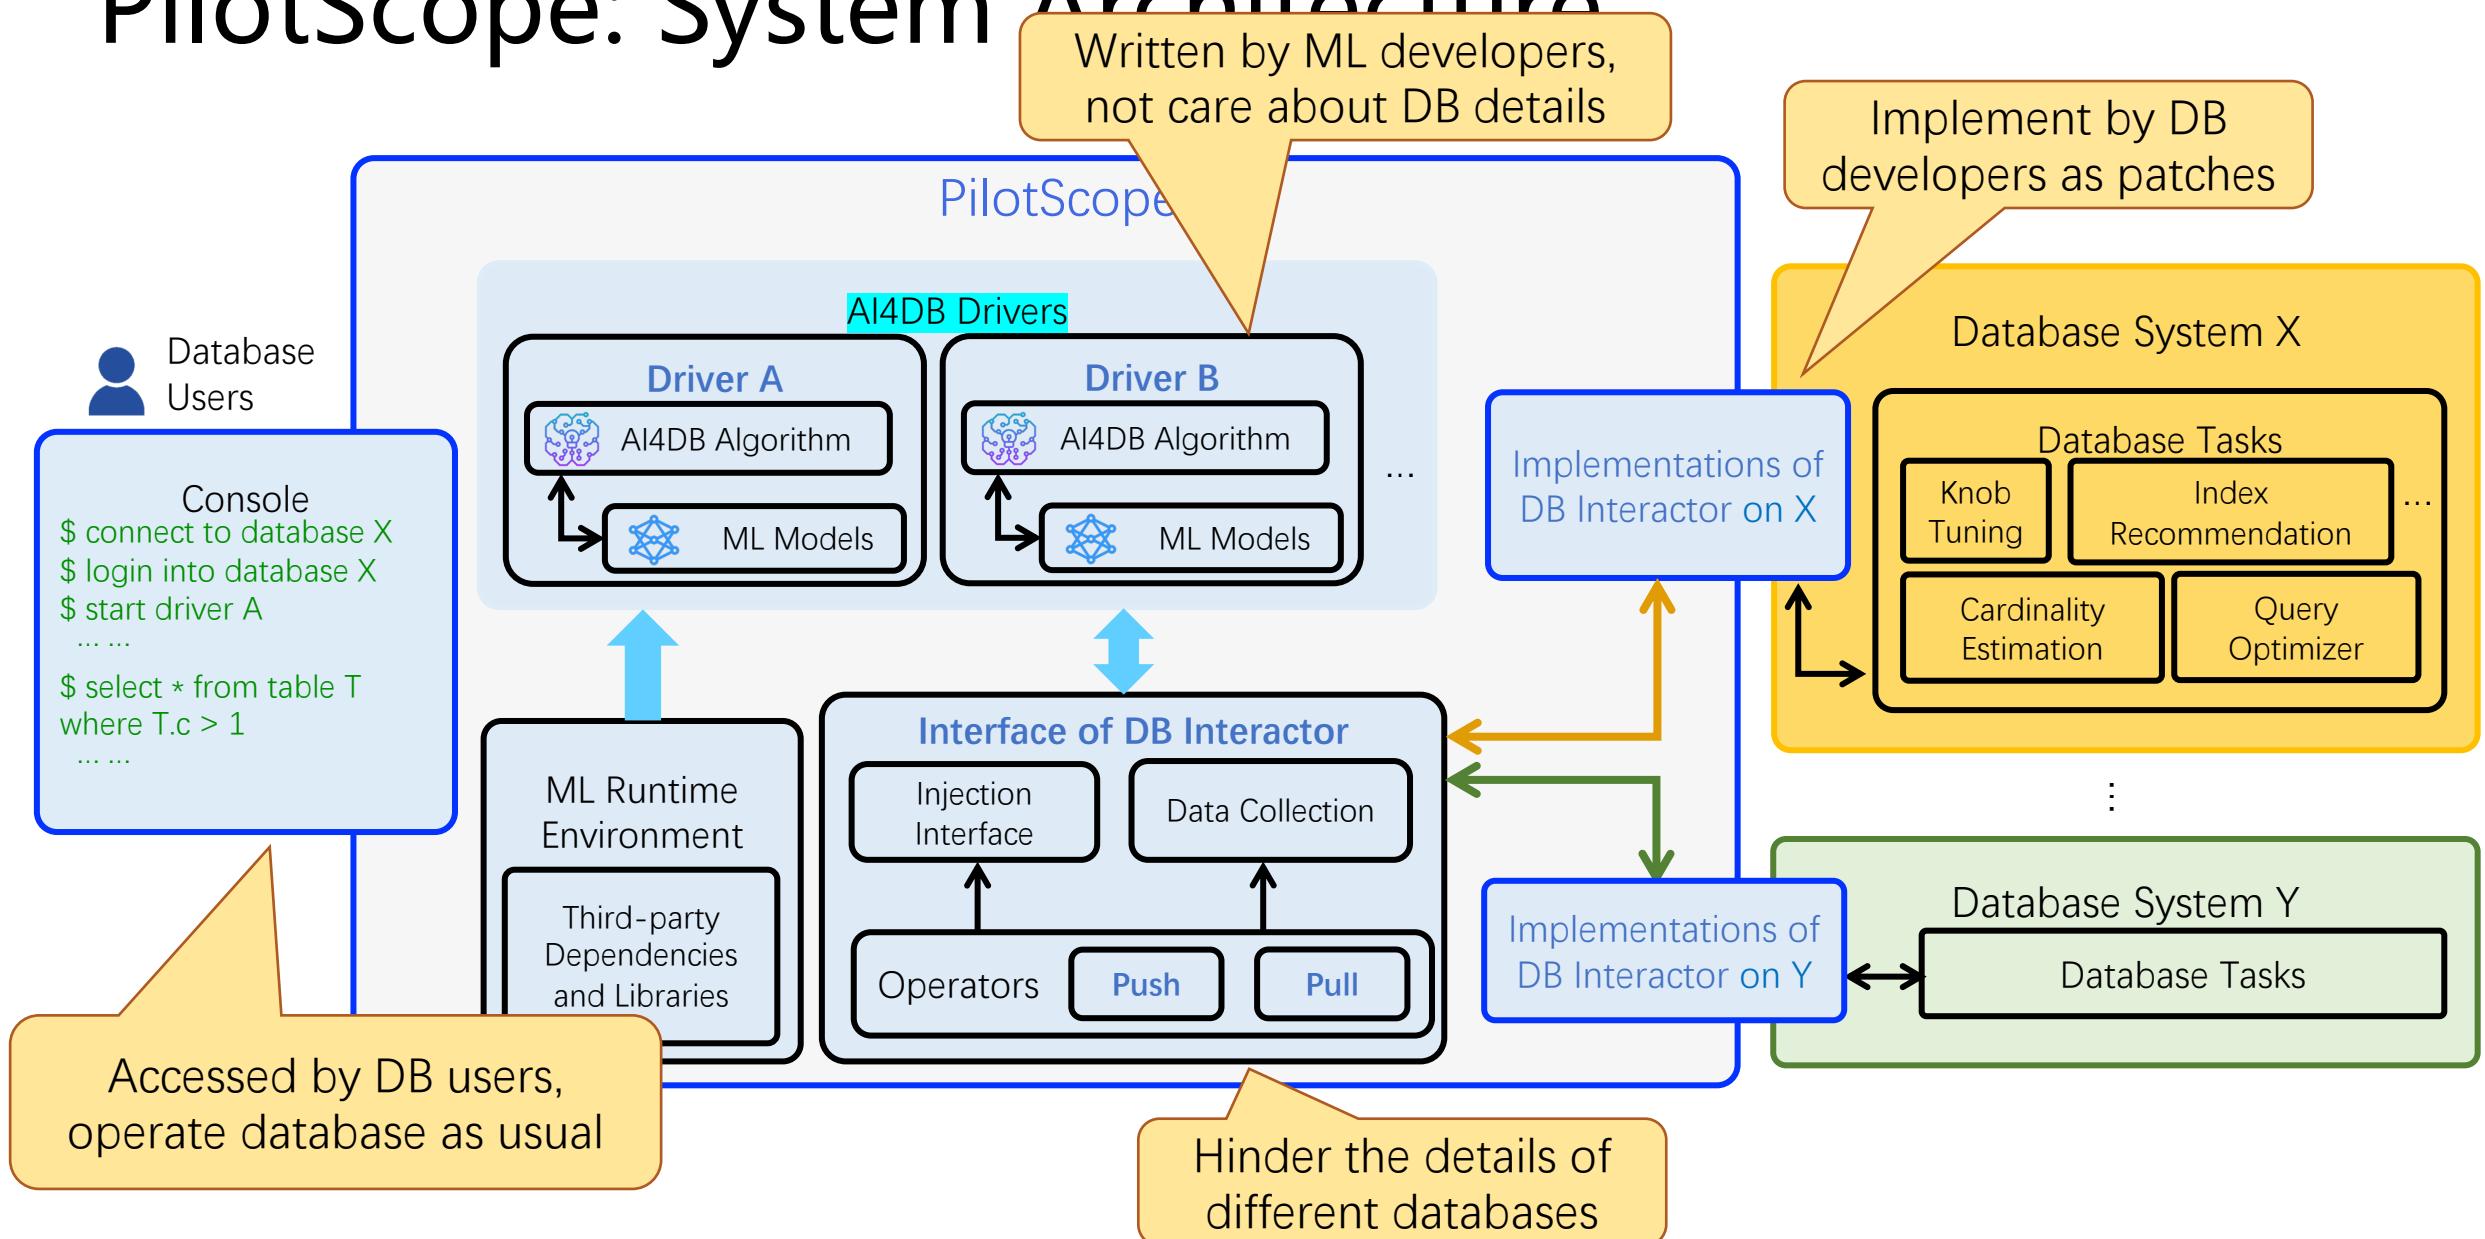
- A desirable deployment system is very challenging
  - ML and DB have very different paradigms: language, dependency, ...
  - AI4DB algorithms are very diversified
    - An AI4DB task usually involves a long pipeline on data interaction
  - Database systems are very complex
    - Difficult and risky to modify the codebase
- .....

# PilotScope: Framework

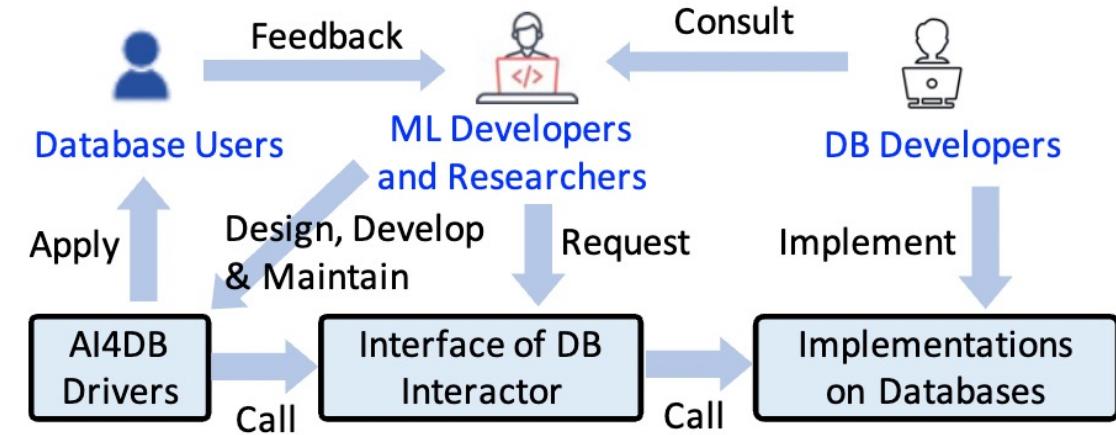
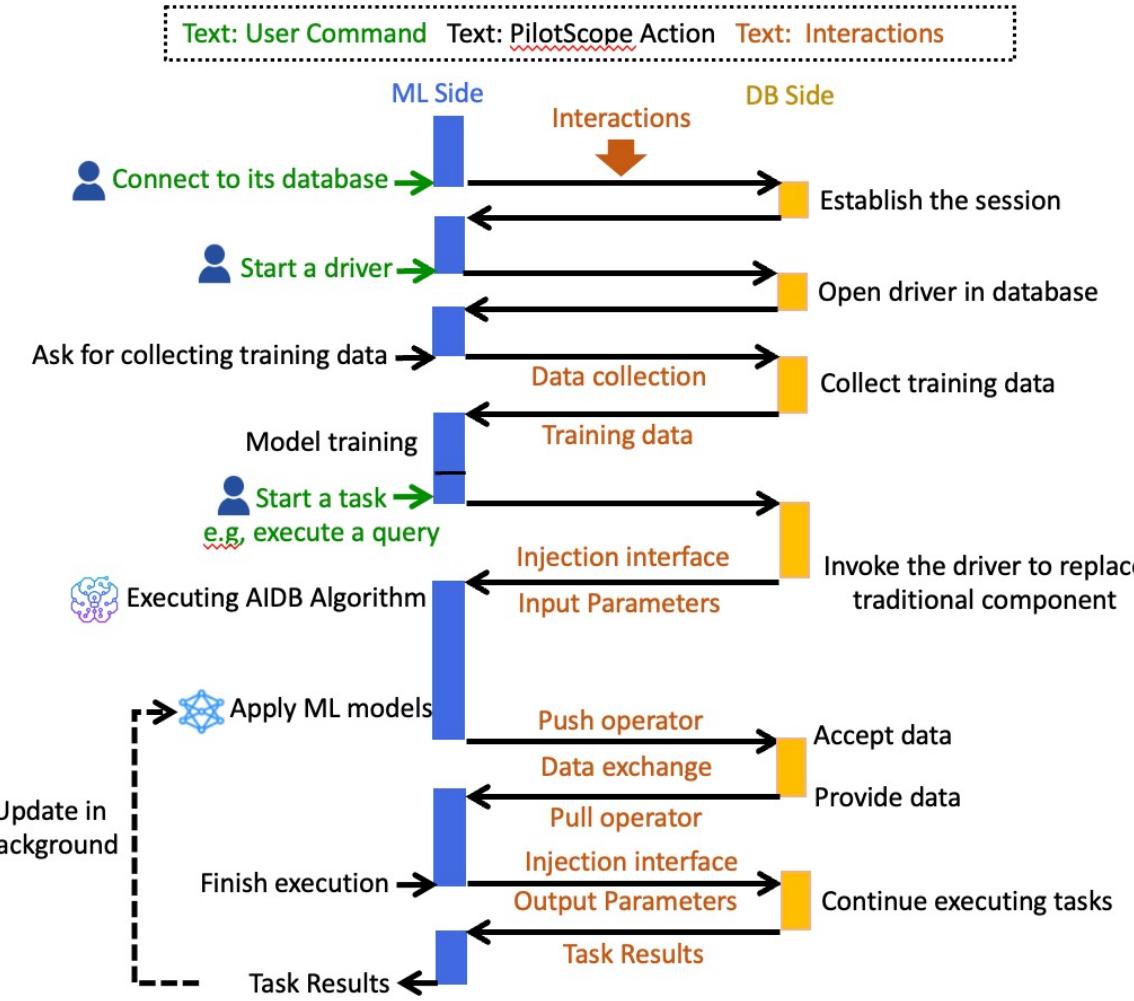
- How to bridge the gap between AI and DB?



# PilotScope: System Architecture

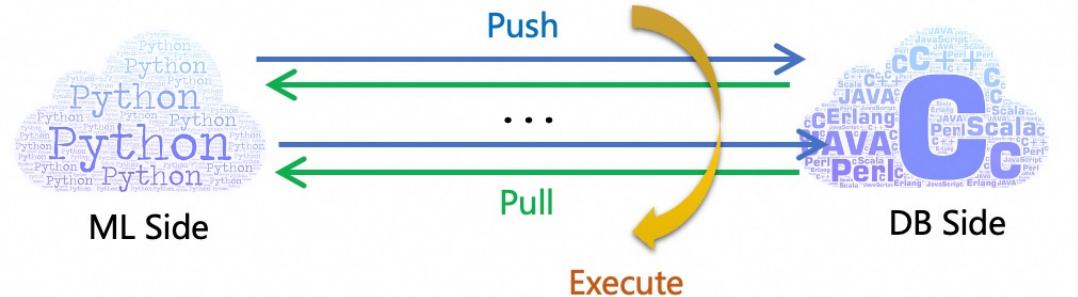


# PilotScope: Workflow



# PilotScope: DB Interactors

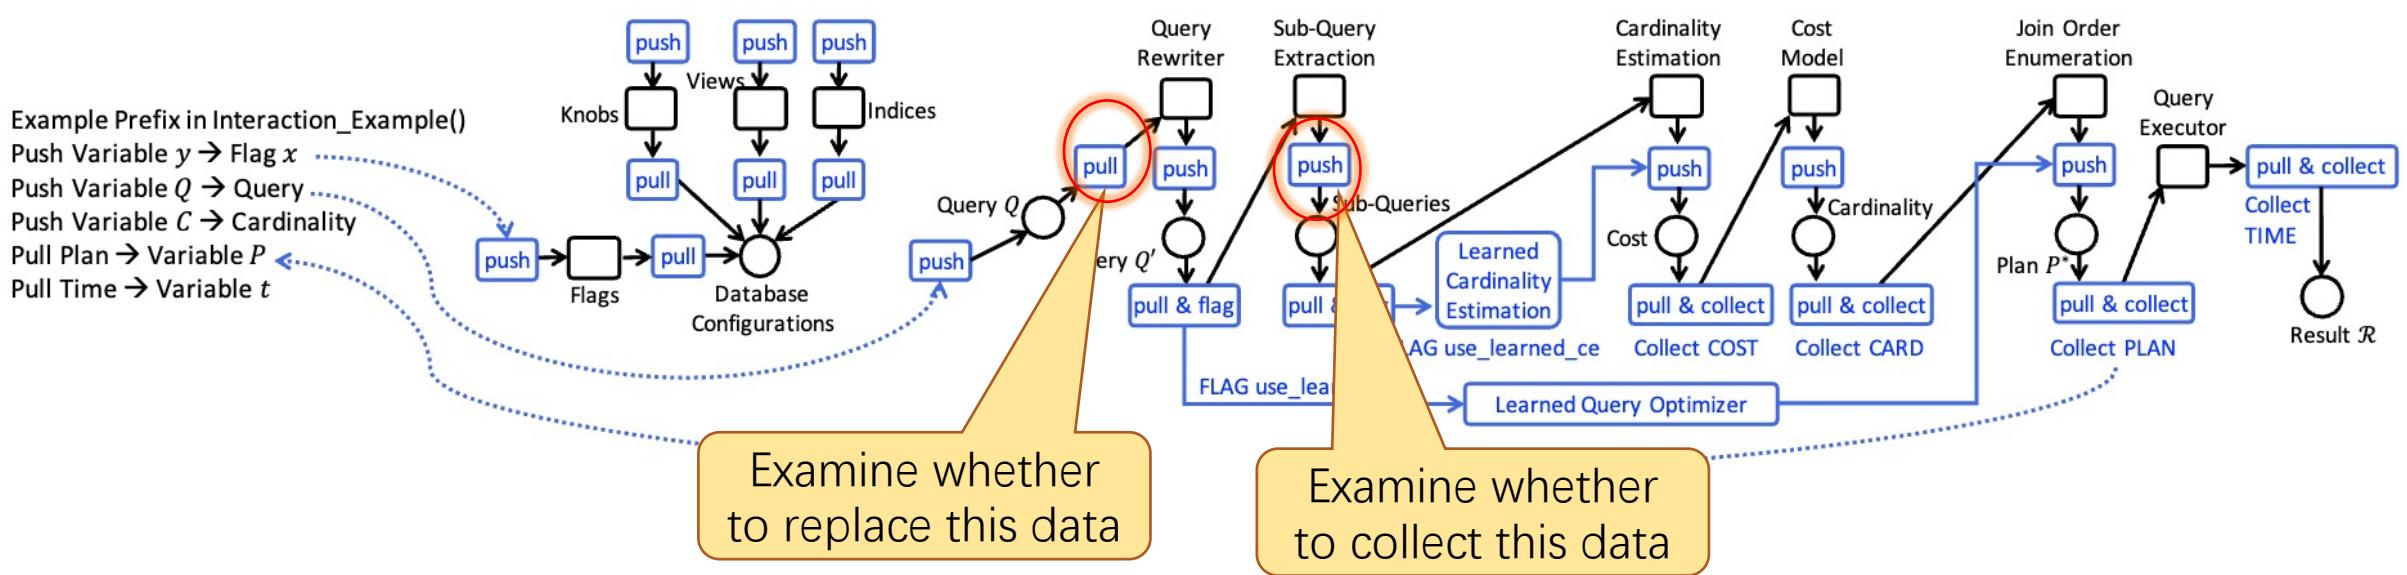
- Flexible APIs
  - Push: send data to DB
  - Pull: fetch data from DB
  - Any combination in a session
  - Execute: trigger execution
- Supported data types
  - Knob/hint/index
  - Sub-query/cardinality/cost
  - Plan/Latency
  - ...



```
def test_all_data_interactor(self):  
    knob = {"max_parallel_workers_per_gather": 0}  
    hint = {"enable_hashjoin": "off", "enable_mergejoin": "on"}  
    query = ("select count(*) from posts as p, postlinks as pl, posthistory as ph where p.id = pl.postid and pl.postid = ph.postid and " +  
             "p.creationdate>=1279570117 and ph.creationdate>=1279585800 " +  
             "and p.score < 50;")  
    data_interactor = PilotDataInteractor(self.config)  
    data_interactor.push_hint(hint)  
    data_interactor.push_knob(knob)  
    data_interactor.pull_physical_plan()  
    data_interactor.pull_subquery_card()  
    data_interactor.pull_estimated_cost()  
    data_interactor.pull_record()  
    data_interactor.execute(query)
```

# PilotScope: Implementation

- APIs in the interface: patches before/after each component
  - Built-in tools: commands/hooks, ...



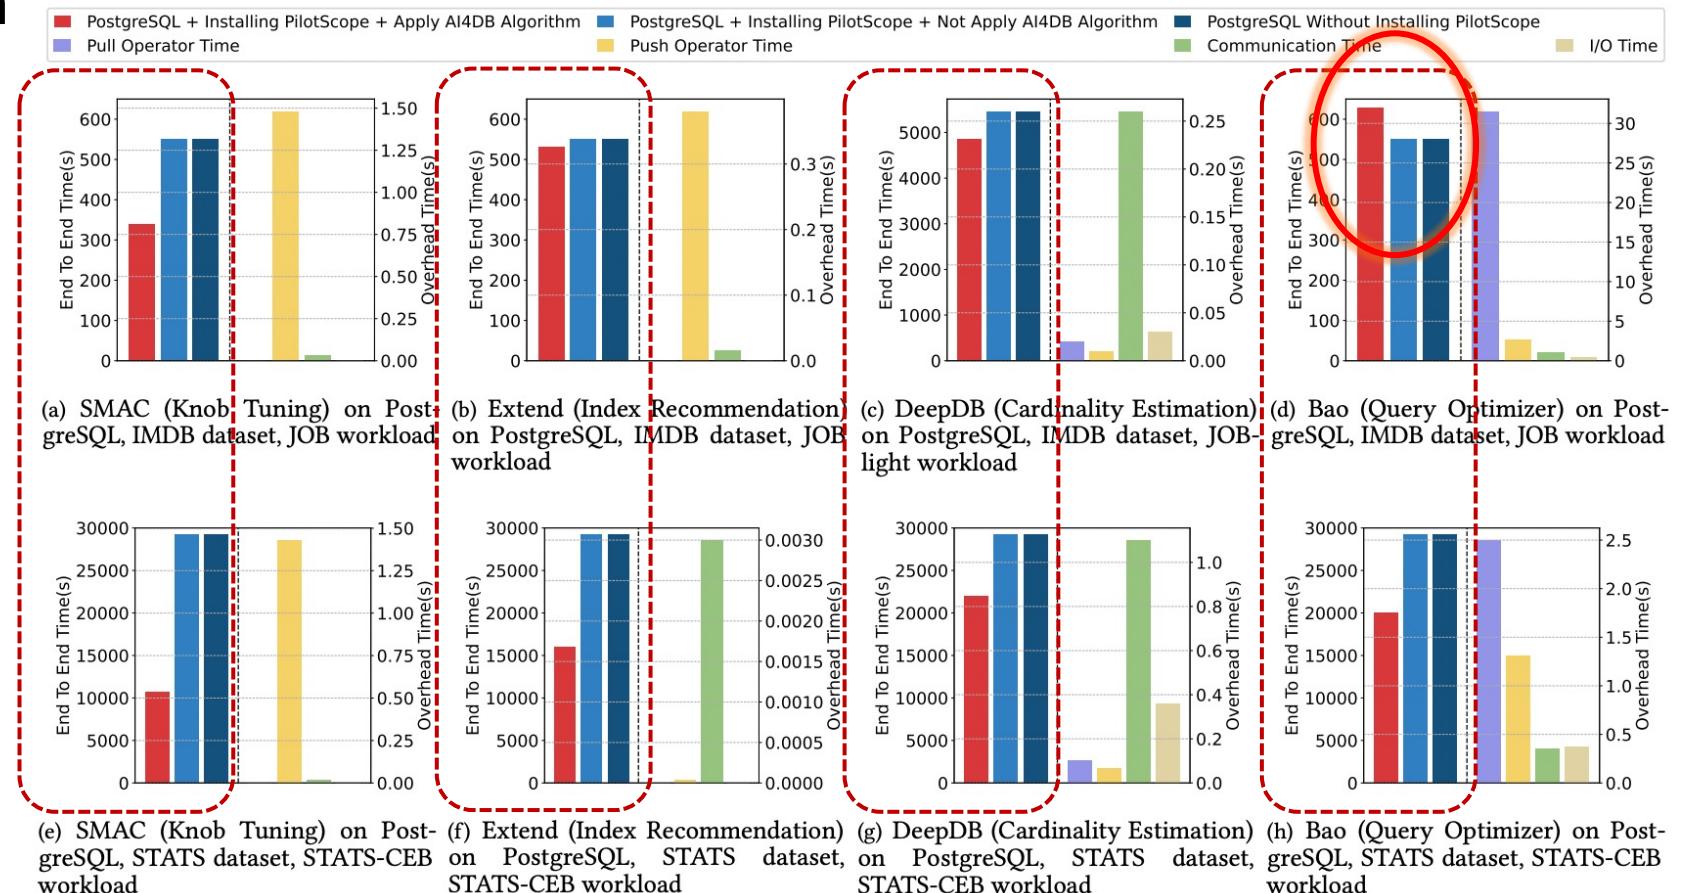
# PilotScope: Applications

- 4 AI4DB tasks, 15+ algorithms on PostgreSQL and Spark
- Easy to extend to similar tasks: view advisor, learned cost model...

Task	Algorithm	Description
Knob Tuning	SMAC [26]	Bayesian Optimization
	GP-BO [44]	
Index Recommendation	DDPG [56]	Reinforcement Learning
	AutoAdmin [17]	Enlarging Small Set
	DB2Advisor [49]	
	Extend [45]	
Cardinality Estimation	DTA [1]	Reducing Large Set
	Drop [51]	
	Relaxation [14]	Reducing Large Set
E2E Query Optimizer	Cophy [20]	Linear Programming
	MSCN [30]	Query-Driven Model
Cardinality Estimation	NeuroCard [53]	Data-Driven Model
	DeepDB [25]	
E2E Query Optimizer	Bao [40]	Tuning Hint
	Lero [59]	Scaling Cardinality

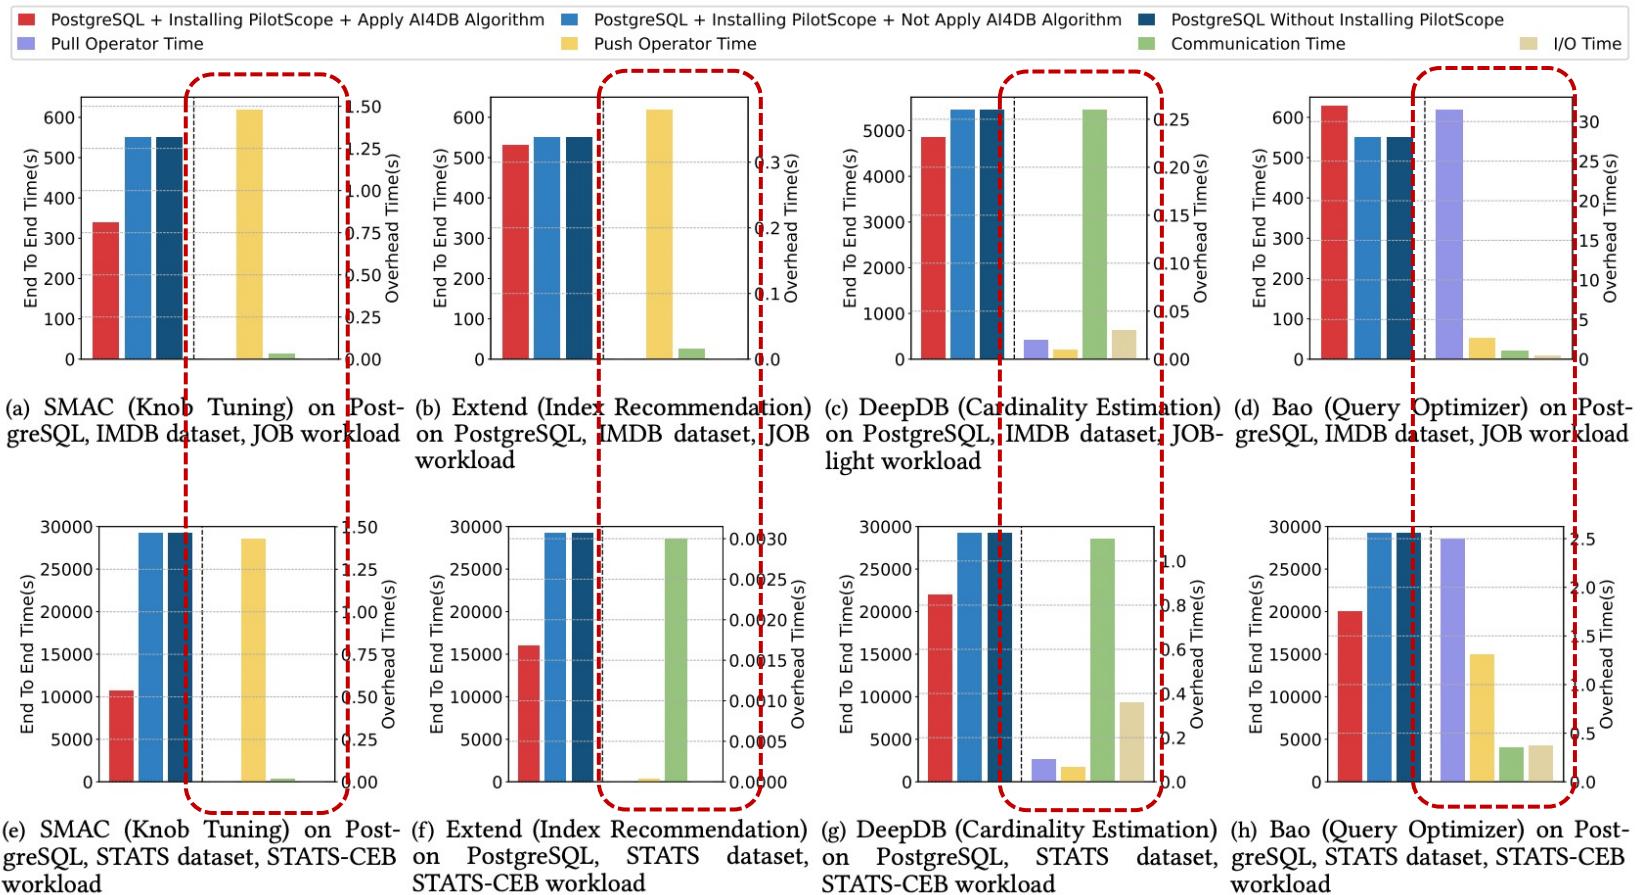
# PilotScope: Evaluations

- AI4DB algorithms may not always work well
  - By PilotScope, we know how/when it could perform well/bad → guidelines for future research



# PilotScope: Evaluations

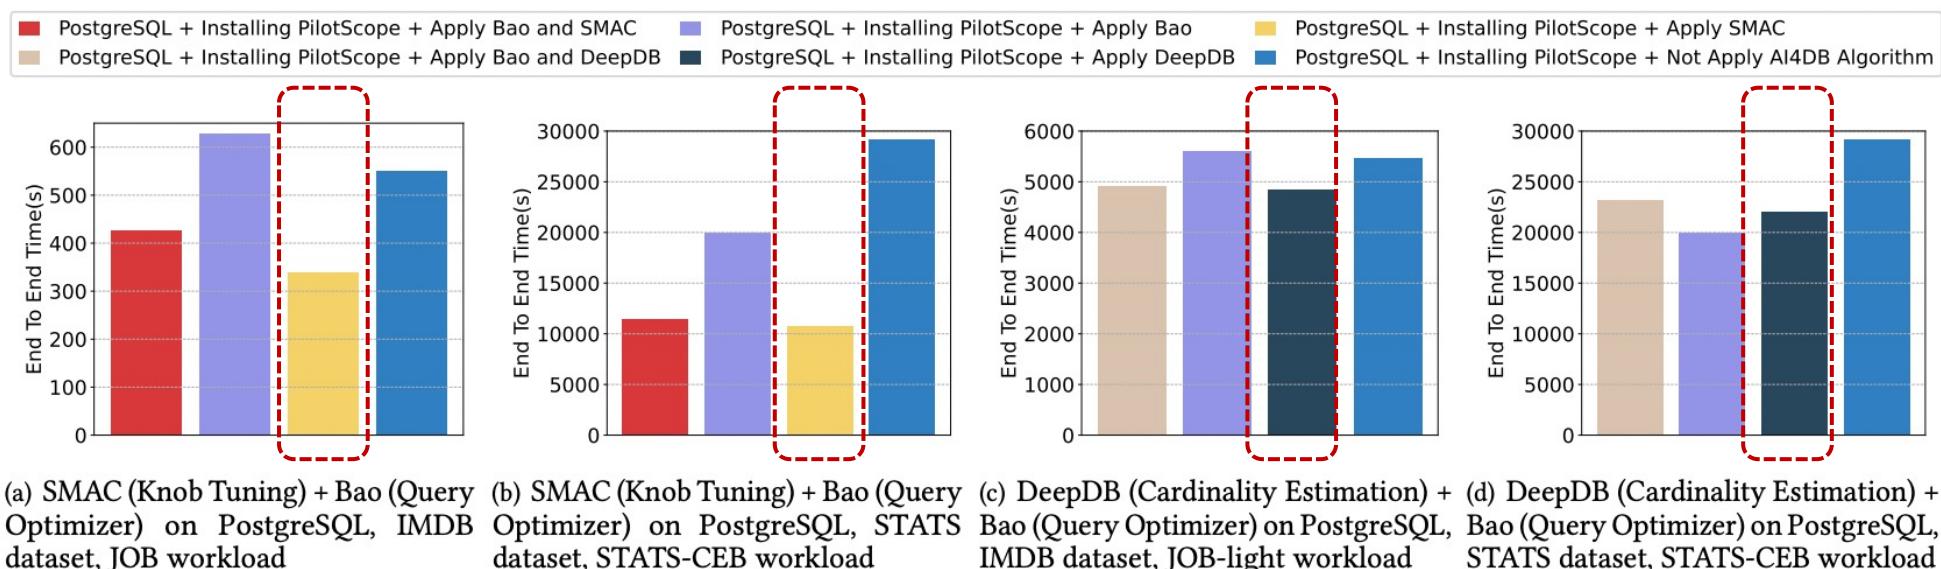
- Overhead of PilotScope is negligible



Experimental data from [PilotScope: steering databases with ML drivers , VLDB 2024]

# PilotScope: Evaluations

- PilotScope could support to deploy and run multiple AI4DB algorithms at the same time
- $1 + 1 > 2?$  → not always hold, correlations are complex
  - A new and interesting topic in AI4DB



Experimental data from [PilotScope: steering databases with ML drivers , VLDB 2024]

# PilotScope: More to Come

- Open-Source: <https://github.com/alibaba/pilotscope>
- Welcome for the contributions  
to build this community together



language Python language C language Scala license Apache 2.0 contributions Welcome  
docs Usage Guideline docs Develop Guideline docs API Reference  
AI4DB driver Knob Tuning AI4DB driver Index Recommendation AI4DB driver Cardinality Estimation  
AI4DB driver E2E Query Optimizer  
database PostgreSQL 13.1 database Spark 3.1

Figure generated by Tongyi Wanxiang (<https://tongyi.aliyun.com/wanxiang/>)

# PilotScope: Demo

Demo Time!



# Part 5: Summary and Future Opportunities

- Summary
- Future Opportunities

# Summary

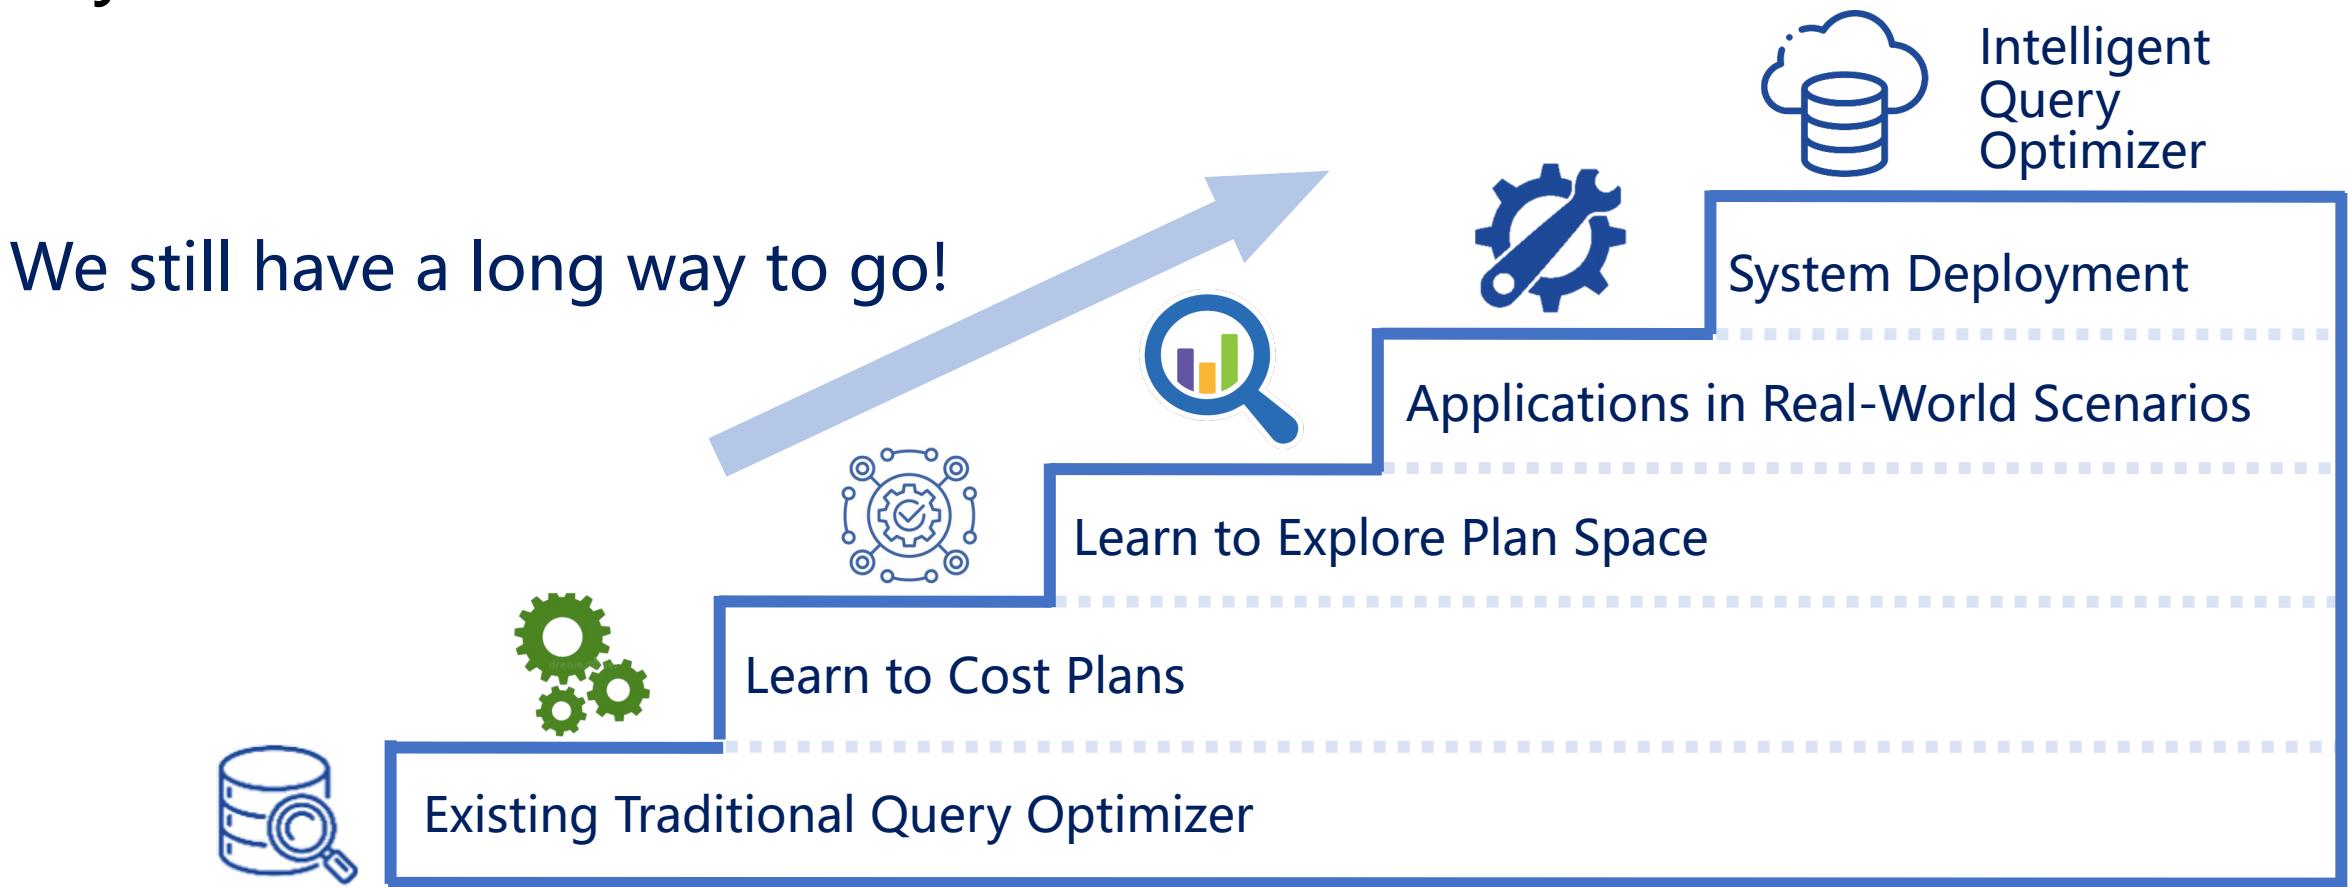
- Query optimizer is a suitable experimental plot for AI4DB → the prosperity of learned methods are developed in recent
- Learn to cost plans
  - Cardinality plays the most important role
  - Cardinality estimation: data-driven, query-driven and hybrid methods
  - Cost model: models for single/concurrent queries
  - Exhibit advantages in benchmark evaluations, but still not ready for actual deployment with some shortcomings

# Summary

- Learn to explore plan space
  - Learn the value network to explore plans from scratch
  - Learn to steer existing query optimizers with different knobs
  - Lero: from learning-to-predict to learning-to-rank
  - Eraser: eliminate performance regressions
- Applications and deployment
  - Some attempts on customized systems and tasks
  - PilotScope: a middleware for ease of deployment, generally applicable to multiple algorithms and systems

# Future Work

- Final goals: practical and intelligent query optimizer and beyond



# Future Work

- Opportunity I: cardinality estimation
  - Fusion of data-driven and query-driven methods
  - Adaptive and general methods: different data, workloads, OLAP/OLTP, ...
- Pretrained cross-database model is possible routine
  - One model generally applicable to different databases/settings
  - Very friendly to deploy and apply

# Future Work

- Opportunity II: optimize cardinality (cost) estimation with others tasks together
  - Complex relations of knob/index/... to query optimizer
  - 1 + 1 > 2, towards a more powerful query optimizer

# Future Work

- Opportunity III: LLMs and learned query optimizer
  - LLMs could help
    - Text2SQL
    - Query rewriter
    - Data cleaning

# Examples that AI/LLMs Help a Lot

- Text-to-SQL (DataWorks Copilot)

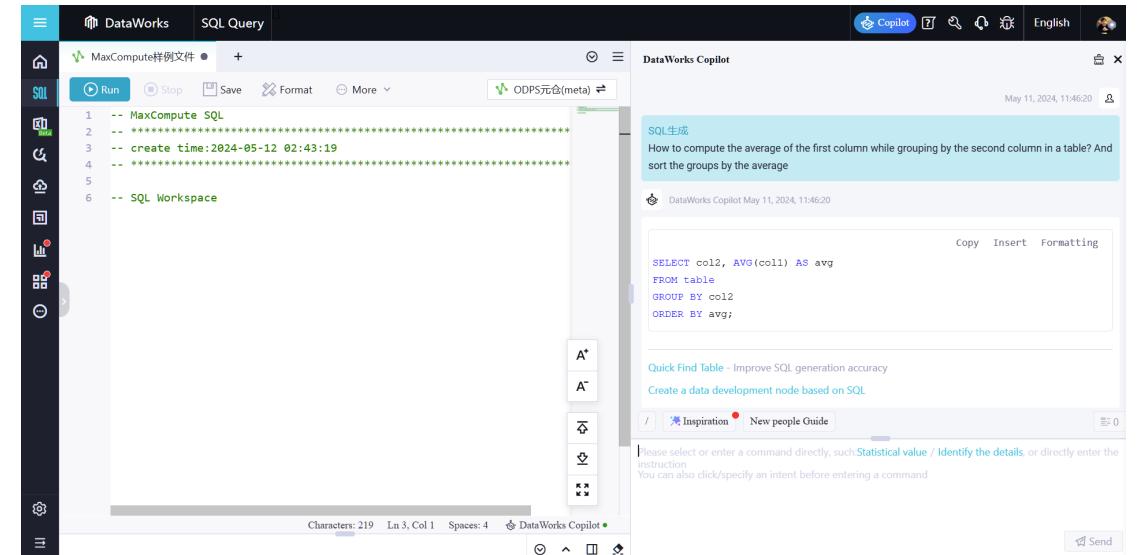
## Prompt for LLM

- (1) Meta-data (table names, info, key-foreign key)
- (2) SQL examples to join the tables
- (3) (instruction, SQL) example pairs
- (4) Text-to-SQL instruction (user's input)

(1)-(3) are *hidden from the end user*

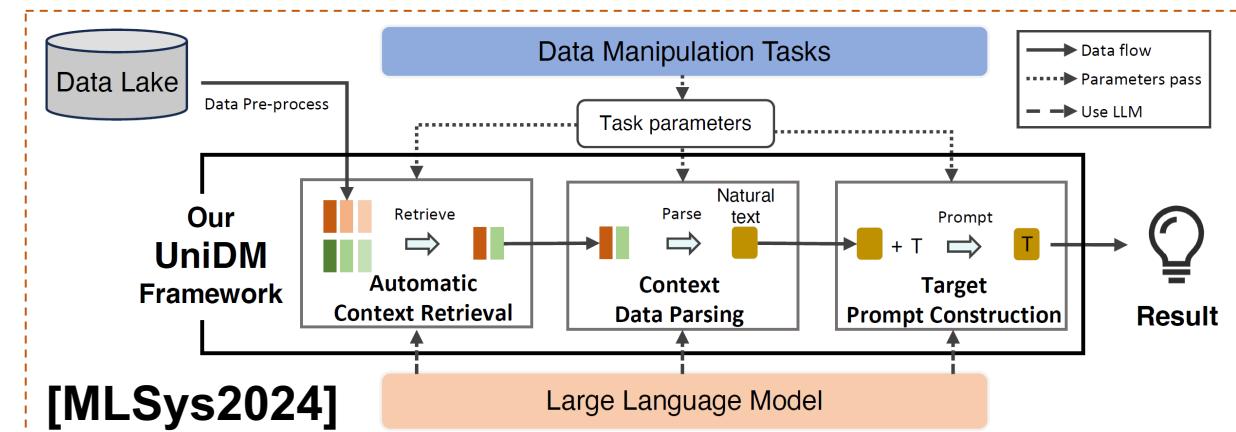
[4] *is the user's input*

[VLDB2024]



- AI-aided data manipulation (UniDM)

- Handling missing data
- Data transformation
- Error detection



# Future Work

- Opportunity III: LLMs and learned query optimizer
  - Can LLM help query plan generation?
  - How to fuse LLM with the information of data distributions?
    - Interplay between large and small models
      - Combine LLM reasoning and data knowledge representations together
      - Efficiency and effectiveness balance

# Future Work

- Opportunity IV: extend the techniques to more AI4DB or even DB4AI tasks beyond query optimizer
  - Index recommendation
  - View advisors
  - System diagnosis
  - ...
- From intelligent query optimizer to intelligent database

# Learned Query Optimizer: What is New and What is Next

## Q & A



red.zr@alibaba-inc.com



lianggui.wlg@alibaba-inc.com

bolin.ding@alibaba-inc.com

jingren.zhou@alibaba-inc.com